

Semantic Matching in Medieval Latin: A Benchmark, Layerwise Analysis, and Geometric Repair

Anonymous ACL submission

Abstract

Medieval canon law survives as families of manuscript copies, shaped by spelling variation, fragmentary copying, and editorial change. Scholars who consolidate duplicate passages and link related witnesses need strict semantic matching. We introduce a Medieval Latin benchmark of 1,705 canon-law fragments grouped into 840 meaning-equivalent families, and we define two tasks on it: pairwise duplicate detection and open-set directory routing. A layerwise analysis of six pre-trained models from three architecture families separates two effects that are easy to conflate. Anisotropy is universal: every model has a layer where the mean cosine between arbitrary passages reaches 0.23 to 0.95, and no baseline layer in any model separates equivalent from non-equivalent passages by more than 0.29 in mean cosine. Catastrophic mid-depth collapse is not universal: the Latin and multilingual T5 encoders fall to AUROC 0.497, 0.539, and 0.654 in their middle layers, while LaBSE, Qwen3-0.6B, and KaLM-mini improve close to monotonically with depth. All-but-the-Top post-processing, fit only on training embeddings, repairs all six models. Best-layer AUROC moves into a 0.971 to 0.984 band, the cosine gap rises to between 0.525 and 0.610, and five-seed top-1 routing accuracy tightens from a 36.6-point spread across models to a 1.0-point spread. SIF pooling gives a weaker frequency-based correction, and it helps only the collapsed T5 encoders. We then ask how local attributions for the retrieval score change under the repair, using integrated gradients and retrieval-adapted MaRC. The answer is narrow: ABTT improves leave-one-out rank faithfulness in every LaTa, PhilTa, and mT5-base model-method cell, while ERASER-style rationale metrics show mixed effects.

1 Introduction

Medieval canon law preserves early records of legal reasoning in Western Europe. The corpus does




Figure 1: Cosine similarity distributions for PhilTa at the last layer and at the diagnostic mid layer, before and after ABTT. Equivalent pairs share a directory; non-equivalent pairs do not. At the collapsed layer, ABTT separates the two distributions more clearly.

not survive as one stable text. Manuscript copies differ in spelling, wording, fragment boundaries, and editorial history. Scholars often need to decide whether two passages express the same legal content despite those surface changes. That decision drives everyday editorial work: consolidating duplicates, linking witnesses, and tracing how a text was transmitted.

We study this problem as strict semantic matching. The benchmark contains 1,705 Medieval Latin canon-law fragments arranged into 840 families of meaning-equivalent variants. We evaluate two tasks. Task A asks whether two passages express the same meaning. Task B turns the same embeddings into an open-set workflow: a system either routes a new passage to an existing directory or marks the passage as new.

A simple retrieval system embeds each passage and ranks candidates by cosine similarity. A final-layer vector, though, says nothing about where inside the network the retrieval signal appears, weakens, or disappears. We therefore work layer by layer, because the failure mode is geometric.

Our layerwise view of six models separates two effects that are often reported together. The first,

anisotropy, is universal. Every model we test has a layer where pooled passages crowd into a narrow cone, with mean cosine between arbitrary passages reaching 0.23 to 0.95. No baseline layer of any model separates equivalent from non-equivalent passages by more than 0.29 in mean cosine. The second, a catastrophic mid-depth collapse of retrieval, is not. LaTa, PhilTa, and mT5-base lose almost all pairwise separation in their middle layers, while LaBSE, Qwen3-0.6B, and KaLM-mini improve close to monotonically with depth. The collapse is therefore a symptom of the T5 encoder family on this corpus, not a property of pre-trained encoders in general.

Anisotropy motivates our use of All-but-the-Top (ABTT) post-processing (Mu and Viswanath, 2018). We apply it as a repair after pooling, with no update to model parameters: fit the dominant principal components on training embeddings, remove those directions from train and test embeddings, then score passages with cosine similarity. The repair, unlike the collapse, is universal. ABTT lifts every model into a narrow high-accuracy band on both tasks, and at each T5 encoder’s worst baseline layer it recovers pairwise AUROC to within 1.0 point of that same model’s best ABTT layer. We also report SIF weighting (Arora et al., 2017), which reweights tokens by frequency before pooling. SIF corrects part of the same dominant direction, but it only helps the models whose retrieval has collapsed, and it never matches ABTT.

Geometry diagnoses collapse and recovery, but it does not select the best retrieval layer. For operational retrieval and attribution we use a train-only rule, described in Section 4.2, which locks LaTa layer 7, PhilTa layer 1, and mT5-base layer 1 before attribution metrics are computed.

Finally, we compare local explanations of the score before and after the repair. Integrated gradients (Sundararajan et al., 2017) gives a gradient view of token importance, and retrieval-adapted MaRC (Brinner and Zarrieß, 2023) gives a learned-mask view after replacing MaRC’s classification target with pairwise retrieval cosine. Both explain a fixed query-candidate score, and both stay local diagnostics rather than global accounts of model behavior.

This paper makes four contributions: (1) we introduce a benchmark for strict semantic matching in Medieval Latin canon law, with a pairwise task and an open-set routing task; (2) we separate universal anisotropy from the T5-specific mid-depth re-

trieval collapse across six models, measuring both with top-PC dominance, effective rank, and cosine concentration; (3) we show that ABTT repairs all six models without updating parameters, compressing a 36.6-point spread in five-seed top-1 routing accuracy to 1.0 point, while SIF gives a weaker frequency-based correction that helps only the T5 encoders; and (4) we adapt MaRC-style mask optimization to bi-encoder retrieval and compare it with integrated gradients at predeclared retrieval-selected layers, finding leave-one-out rank faithfulness gains and mixed ERASER-style rationale metrics.

2 Related Work

Anisotropy, layer geometry, and repair. Pre-trained spaces occupy narrow cones (Gao et al., 2019; Ethayarajh, 2019) where a few coordinates dominate cosine (Kovaleva et al., 2021; Timkey and van Schijndel, 2021), and corrections remove or rescale those directions (Mu and Viswanath, 2018; Su et al., 2021) or downweight frequent tokens (Arora et al., 2017). Information spreads across depth (Tenney et al., 2019; Liu et al., 2019), intermediate layers can beat the last (Skean et al., 2025), and multilingual work is rarer (Rajaei and Pilehvar, 2022). Layerwise profiles put the mid-depth anisotropy peak in decoder-only models (Razhigaev et al., 2024); T5 enters that work with encoder and decoder states pooled (Godey et al., 2024), and its weak mean-pooled similarity is a final-layer result (Ni et al., 2022). Anisotropy is neither universal (Machina and Mercer, 2024) nor always harmful (Ait-Saada and Nadif, 2023), so we report geometry and retrieval separately and reuse ABTT with a train-only layer rule.

Latin and classical-language NLP. Latin has dedicated encoders and encoder-decoders (Bamman and Burns, 2020; Riemenschneider and Frank, 2023a) is covered by the multilingual encoder-decoder mT5 (Xue et al., 2021), and appears in multilingual and decoder-only embedders (Feng et al., 2022; Zhang et al., 2025; Hu et al., 2025). Shared tasks target tagging (Sprugnoli et al., 2020), retrieval work allusion (Riemenschneider and Frank, 2023b), and charter work semantic change (Liu et al., 2024); no general benchmark (Thakur et al., 2021; Muennighoff et al., 2023) covers medieval Latin retrieval.

Text reuse in digital humanities. Text-reuse systems match n-grams (Coffee et al., 2012; Smith et al., 2013) and add distributional semantics for allusive reuse (Manjavacas et al., 2019; Burns et al., 2021), ranking candidates in a closed collection. Our routing task is open-set: a fragment may belong to no family.

Faithfulness of token attributions. Integrated gradients (Sundararajan et al., 2017) and learned masks (Brinner and Zarri  , 2023) score tokens, and their faithfulness is debated (Jain and Wallace, 2019; Jacovi and Goldberg, 2020). Erasure protocols (DeYoung et al., 2020; Atanasova et al., 2020) target classification and retrieval explanation targets ranking (Anand et al., 2022); we apply both to a bi-encoder cosine score.

3 Task and Dataset

3.1 Task Definition

The benchmark defines two tasks over the same fragments and the same pooled embeddings. They differ in what the system has to decide.

Task A: pairwise duplicate detection. Let p_1 and p_2 be two passages and let $\equiv$ denote semantic equivalence. The label is

$$y = \begin{cases} 1 & \text{if } p_1 \equiv p_2, \\ 0 & \text{otherwise.} \end{cases}$$

The model produces $\cos(e_{p_1}, e_{p_2})$ from pooled embeddings. Task A uses no directory structure and no threshold at scoring time. Evaluation uses the test $n \times n$ cosine matrix.

Task B: directory routing. Task B uses the same embeddings in a scholar-facing retrieval workflow. Given a query passage q and the existing labelled directories, the system routes q to a directory or marks q as new. We evaluate two routing variants. In autonomous routing, each query receives the maximum cosine score to any other test file. A threshold τ , fit on training same-directory and different-directory pairs, decides whether the query enters its nearest neighbour’s directory or the new class. In top- k ranking, each directory receives the best document match score, and a pseudo-option __NEW__ receives score τ . The simulated expert succeeds when the labelled directory appears in the shortlist.

3.2 Dataset Construction

The corpus contains 1,705 Medieval Latin canon-law fragments from the Carolingian Canon Law project (CCL), an ongoing digital edition founded in 2009 (Firey, 2009; Eichbauer, 2014). Its manuscripts are mostly ninth-century, and most have never been edited to modern standards. Transcription is manual, mostly from manuscript images and for some texts from eighteenth-century printings. Before 2019 transcribers worked from microfilm, and CCL contributors converted the files to TEI-P5. Since 2019 they work in an editor beside a IIIF image of the folio, and each transcription then passes a proofreading desk, where an approved proofreader verifies it against the image to the letter and adds TEI markup. The interface credits each unit’s transcriber and proofreader. We hold the fragments as plain text with the labels as supplied, and we are confirming with CCL how they were extracted from the TEI.

CCL labels every unit twice, and our directories reuse the second label. A siglum names the manuscript witness; a key names the original source, such as a fourth-century council or a papal decretal. Most of our 840 directories are keys, a minority name the source by a modern edition citation or a biblical reference, and a handful group passages whose source is unidentified, all of them singletons that contribute no positive pairs. The files in a directory are surviving hand copies of one source, each free to vary from the others, so the benchmark asks whether an unidentified witness belongs to a known source. Directory sizes range from one to ten files. Variants inside a directory count as positive matches and files from different directories count as negatives.

3.3 Train/Test Split

We split the corpus into 847 training files and 858 test files with a leak-free protocol (Table 1). The 545 singleton directories are distributed at the file level. Doubleton directories are assigned as whole folders, 54 to train and 54 to test, so both files from a train doubleton contribute a positive training pair. For directories with at least three files, the split varies the train allocation inside each size class while keeping each class near a balanced file split. This avoids a rule where every size-three directory contributes no positive training pairs. The final split contains 565 positive training pairs and 595 positive test pairs.

Category	Dirs	Train	Test
Singletons (1 file)	545	272	273
Doubletons (2 files)	108	108	108
Multi-file (≥ 3)	187	467	477
Total	840	847	858

Table 1: Train and test split statistics under the v2 varied-allocation protocol. Doubleton directories are assigned as whole folders. The test set contains 535 existing files, with a same-directory partner in test, and 323 new files.

All fitting uses training data only. This includes ABTT principal components, routing thresholds, ABTT dimensionality choices, and the operational layer rule used for attribution.

4 Methodology

Our pipeline has three stages: pool token states from one layer into a passage vector, correct the geometry, then score pairs by cosine.

4.1 Models

We evaluate six pre-trained models from three architecture families:

- **LaTa**: a T5-based seq2seq model pre-trained on Latin (Riemenschneider and Frank, 2023a), with 12 encoder layers.
- **PhilTa**: a philologically oriented Latin T5 variant (Riemenschneider and Frank, 2023a), with 12 encoder layers.
- **mT5-base**: a multilingual T5 encoder-decoder (Xue et al., 2021), with 12 encoder layers.
- **LaBSE**: a language-agnostic BERT sentence encoder covering 109 languages (Feng et al., 2022), with 12 encoder layers.
- **Qwen3-0.6B**: a multilingual decoder-only embedding model with 28 layers (Zhang et al., 2025).
- **KaLM-mini**: a multilingual decoder-only embedding model with 24 layers (Hu et al., 2025).

4.2 Embedding Extraction and Layer Selection

We extract hidden-state representations from every transformer block, starting at layer 1 and skipping

layer 0, the static embedding table. Token representations at each layer are pooled into one passage vector by plain mean pooling or SIF weighting (Arora et al., 2017). For decoder-only models, last-token pooling is another common option; Appendix C compares it with mean pooling under ABTT.

The paper-facing pooling baseline removes only tokenizer artifacts with no lexical content, such as whitespace, padding, and tokenizer-prefix-only tokens. Punctuation and numerals remain in the pooled representations. Attribution figures display original decoded tokens, even when a token contributes little.

Layer geometry and layer selection serve different roles. We use label-free geometry diagnostics, including top-PC dominance, effective rank, and cosine concentration, to identify collapsed layers and test whether ABTT repairs them. For operational retrieval and attribution, we then apply a pre-declared train-only retrieval rule. The rule chooses the earliest layer within 0.5 percentage points of the best training-set directory accuracy at rank 1 under ABTT. It selects LaTa layer 7, PhilTa layer 1, and mT5-base layer 1 for attribution, with $D = 10$ ABTT components. The strongest diagnostic collapse layers are LaTa layer 8, PhilTa layer 6, and mT5-base layer 5, and we keep those as mechanism checks rather than attribution layers.

4.3 Post-Processing

ABTT. Our main post-processing method is All-but-the-Top (ABTT) (Mu and Viswanath, 2018). For each model and layer, we center the pooled training embeddings, compute the top D principal components by SVD, and remove those components from both train and test embeddings. We search $D \in \{1, 2, 3, 5, 7, 10\}$ and choose D by training-set Task B directory accuracy at rank 1. We apply it as a retrieval-specific repair for anisotropic Latin manuscript representations.

SIF. The main comparison also includes SIF weighting (Arora et al., 2017), which divides each token weight by its estimated corpus frequency and so down-weights high-frequency tokens before the passage vector is formed. SIF and ABTT act on the same problem at different stages. SIF reweights tokens before pooling; ABTT removes dominant directions after pooling. Token probabilities for SIF come from training files only. We report four settings throughout: baseline mean pooling, SIF

alone, ABTT alone, and SIF followed by ABTT.

PCA whitening (excluded). PCA whitening (Su et al., 2021) rescales principal directions rather than only removing the top components. It leaves the different-directory cosine mean at or below 0.014, so the learned threshold degenerates: across all 100 whitening cells every file is routed to an existing directory, existing accuracy is 1.000, new accuracy never exceeds one file in 323, and assignment accuracy sits at the 62.4 percent class prior. That prior beats baseline in 60 of 100 cells while directory accuracy at rank 1 never exceeds 0.415, so the apparent gain is a thresholding artifact and we exclude whitening from the main tables.

4.4 Local Attribution Methods

We evaluate local explanations for a fixed query-candidate retrieval score. Integrated gradients (Sundararajan et al., 2017) gives a gradient-based attribution of pairwise cosine similarity to input tokens. We compute per-sequence token attributions for each side of the pair. The derived heatmap in Figure 5 combines token-token cosine similarity with the marginal integrated-gradient scores from the two sequences.

Our learned-mask attribution is adapted from MaRC (Brinner and Zarrieß, 2023). MaRC was proposed for classification, where a soft input mask preserves a target class score while sparsity and smoothness regularizers encourage concise rationales. We keep the per-instance mask-optimization idea and replace the classification target with pairwise retrieval cosine. Given a query-candidate pair, the mask is optimized on one side while the partner representation stays fixed. Swapping the two sides gives candidate-side attribution.

Let $\mathbf{E}(\mathbf{x}_q) \in \mathbb{R}^{T_q \times d_{\text{emb}}}$ denote the input-embedding lookup, $E_{>0}$ the encoder with the lookup removed, and $\mathbf{z}_c^*$ the candidate representation held fixed during optimization. We parameterize $\mathbf{m} = \sigma(\boldsymbol{\zeta})$ with free logits $\boldsymbol{\zeta} \in \mathbb{R}^{T_q}$ and write $\mathcal{I}$ for content-token positions. For the ABTT variant, define $s_{\text{ABTT}}(\mathbf{a}, \mathbf{b}) = \cos(\text{ABTT}(\mathbf{a}), \mathbf{b})$. The retrieval-adapted MaRC loss is

$$\mathcal{L}(\mathbf{m}) = (s_{\text{ABTT}}(\tilde{\mathbf{z}}_q(\mathbf{m}), \mathbf{z}_c^*) - s_{\text{ABTT}}(\mathbf{z}_q, \mathbf{z}_c^*))^2 + \lambda \sum_{i \in \mathcal{I}} \frac{m_i}{|\mathcal{I}|} + \gamma \sum_{i, i+1 \in \mathcal{I}} \frac{|m_{i+1} - m_i|}{|\mathcal{I}| - 1}, \quad (1)$$

where $\tilde{\mathbf{z}}_q(\mathbf{m}) = \text{pool}(E_{>0}(\text{diag}(\mathbf{m}) \mathbf{E}(\mathbf{x}_q)), \mathbf{a}_q)$. The baseline variant uses raw cosine in place of

the ABTT-corrected cosine in Equation 1. We run 200 Adam steps per pair, initialize $\zeta_0 = \log 9$, and freeze logits outside $\mathcal{I}$ at $-\infty$.

4.5 Evaluation Metrics

Task A. Both Task A metrics come from the test $n \times n$ pairwise cosine matrix. AUROC treats same-directory test pairs as positives and different-directory pairs as negatives. Cosine gap is the mean same-directory cosine minus the mean different-directory cosine.

Task B. For autonomous routing, we report overall assignment accuracy, existing accuracy, and new accuracy. Assignment accuracy scores only the existing-versus-new decision: a file with a partner in test is correct when its maximum cosine to another test file is at least τ , and a file without one is correct when that maximum falls below τ . It never checks which directory the file would enter, so a degenerate threshold that routes everything as existing already attains the 62.4 percent class prior. Read it next to directory accuracy at rank 1, which does require the correct directory. For top- k ranking, cumulative top- K accuracy is the fraction of test queries whose labelled directory appears among the top K options.

Attribution. For attribution, we adapt ERASER-style sufficiency and comprehensiveness metrics (DeYoung et al., 2020) to retrieval cosine. Sufficiency keeps the top-ranked query tokens and measures how much of the full pair cosine remains. Comprehensiveness removes those tokens and measures the score drop. MinFrac reports the smallest retained query-token fraction needed to recover a fixed share of the full cosine. We also report ρ_{LOO} , the Spearman correlation between attribution magnitude and the leave-one-out drop in the same retrieval score.

5 Results

The results move from diagnosis to use: Task A for every model and correction, the layerwise view separating universal anisotropy from the T5-specific collapse, Task B routing under the train-only layer rule, then attribution at the locked layers.

5.1 Task A: Pairwise Duplicate Detection

Table 2 gives the best-layer Task A result for every model under the four post-processing settings. Two patterns matter. First, baseline quality is uneven:

Model	Task A AUROC				Task A cosine gap			
	Base	SIF	ABTT	SIF+ABTT	Base	SIF	ABTT	SIF+ABTT
LaTa	0.938 ₁₂	0.944 ₁	0.971 ₁₂	0.965 ₇	0.238 ₁₂	0.365 ₁	0.525 ₁₂	0.579 ₇
PhilTa	0.939 ₁	0.953 ₁	0.976 ₁	0.981 ₈	0.231 ₁	0.319 ₁	0.573 ₁	0.573 ₈
mT5-base	0.839 ₁₂	0.846 ₁	0.979 ₁	0.971 ₁	0.082 ₁₂	0.036 ₁	0.550 ₁	0.519 ₁
LaBSE	0.957 ₁₂	0.956 ₁₂	0.984 ₁₁	0.974 ₁₀	0.183 ₁₂	0.184 ₁₂	0.610 ₁₁	0.579 ₁₀
Qwen3-0.6B	0.958 ₂₈	0.941 ₂₂	0.972 ₂	0.972 ₂	0.098 ₂₈	0.053 ₂₂	0.553 ₂	0.540 ₂
KaLM-mini	0.970 ₂₃	0.963 ₂₃	0.980 ₁	0.974 ₄	0.056 ₂₃	0.056 ₂₃	0.568 ₁	0.535 ₄

Table 2: Task A pairwise duplicate detection for all six models under four post-processing settings. Each cell is a test-set score at the layer chosen by training-set AUROC, given as the subscript; ABTT uses D tuned on train. Baseline AUROC spans 0.839 to 0.970, and ABTT lifts every model into a 0.971 to 0.984 band. Cosine gap is the mean same-directory minus the mean different-directory cosine. Per-layer grids: Appendix Tables 7, 11, and 14.

AUROC ranges from 0.839 for mT5-base to 0.970 for KaLM-mini. Second, ABTT removes most of that unevenness. Every model lands between 0.971 and 0.984, and the cosine gap rises from a baseline range of 0.056 to 0.238 up to a range of 0.525 to 0.610. The gain is largest where the baseline is weakest: mT5-base gains 0.140 AUROC, LaTa 0.032, and KaLM-mini 0.011.

The cosine gap column explains why a strong AUROC is not enough. KaLM-mini has the best baseline AUROC at 0.970 but the smallest gap at 0.056: it ranks candidates correctly, yet leaves almost no margin for a threshold. ABTT raises every model’s gap by at least 0.28, and that margin is what makes a single learned threshold usable in Task B.

SIF as a frequency-based complement. SIF and ABTT act on the same dominant direction from different sides, and Tables 2 and 3 show how far the frequency side gets. On Task B assignment accuracy, SIF alone adds 9.4 points for LaTa, 8.2 for PhilTa, and 7.3 for mT5-base, the three models whose retrieval collapses. It adds 0.1 point for LaBSE and subtracts 1.9 for KaLM-mini and 6.3 for Qwen3-0.6B. This fits the mechanism. A common component built mostly from frequent tokens can be attenuated by reweighting those tokens before pooling, but a component that survives that reweighting needs the geometric step after pooling. SIF alone never matches ABTT on either task for any model, and adding SIF before ABTT moves best-layer assignment accuracy by at most 1.8 points either way. We read SIF as a partial account of the anisotropy rather than a competing repair, and keep the full sweeps in Appendix E.

5.2 Layerwise Geometry Diagnosis

Figures 2 and 3 show the layerwise retrieval profile, and they separate the two effects. The retrieval col-

lapse belongs to the T5 family. Baseline AUROC bottoms out at 0.497 for LaTa at layer 6, which is chance, at 0.539 for PhilTa at layer 10, and at 0.654 for mT5-base at layer 5. ABTT at those same layers gives 0.963, 0.982, and 0.977. The other three behave differently: baseline AUROC rises with depth, from 0.807 to 0.957 for LaBSE, from 0.859 to 0.958 for Qwen3-0.6B, and from 0.888 to 0.954 for KaLM-mini, and their layer minima are 0.807, 0.859, and 0.862.

Anisotropy, in contrast, is present in all six. At each model’s most collapsed layer, meaning the layer with the largest top-PC variance share, the mean pairwise cosine on the test split is 0.230 for LaTa, 0.584 for PhilTa, 0.314 for mT5-base, 0.878 for LaBSE, 0.953 for Qwen3-0.6B, and 0.887 for KaLM-mini. ABTT with $D = 10$ drives all six to within 0.002 of zero and raises effective rank from as low as 1.00 to between 147.50 and 176.36 (Appendix A). The two families arrive there by different routes. The T5 encoders concentrate variance in one direction, with top-PC shares of 0.956, 0.862, and 1.000, so the cone is low-rank and retrieval dies with it. The non-T5 models hold top-PC shares of 0.506, 0.167, and 0.396 but still carry a large common component, so passages sit close together without losing their relative order. ABTT removes that common component in both cases, which is why the repair generalizes even though the symptom does not.

Figure 1 gives the same diagnosis at the score-distribution level: at PhilTa’s collapsed layer the positive and negative cosine scores overlap heavily before ABTT and separate more clearly after it.

Two-dimensional projections give a visual check. Under baseline mean pooling, points from different directories overlap; under ABTT-only with $D = 10$, same-directory points form tighter groups, with winnable-fragment silhouette scores rising by 0.09 to 0.75 on the three T5 encoders. Appendix H




Figure 2: Task A AUROC against normalized layer depth for all six models. The three T5 encoders in the top row lose almost all pairwise separation in their middle layers, with LaTa reaching chance. The three non-T5 models in the bottom row improve close to monotonically with depth and never dip below 0.80. ABTT holds every model near its ceiling at every depth, so the repair does not depend on the collapse being present.




Figure 3: Cosine gap against normalized layer depth for all six models. The baseline gap stays under 0.29 at every layer of every model, and under 0.11 at every layer of Qwen3-0.6B, so raw cosine carries almost no decision margin anywhere. ABTT lifts all six models to a gap near 0.5 across the whole depth profile. SIF raises the gap for LaTa and PhilTa and leaves the remaining four models close to baseline.

shows t-SNE (van der Maaten and Hinton, 2008) and UMAP (McInnes et al., 2018) projections for all six models.

5.3 Task B: Directory Routing

Task B connects the layerwise analysis to the retrieval workflow through the two variants of Section 3.1: autonomous routing and top- k ranking. The full per-layer routing and ranking tables for the T5 encoders are in Appendix B; Appendix D reports the non-T5 models.

Table 3 reports best-layer autonomous routing. Directory accuracy at rank 1, the column that requires the correct directory, spans 39.3 points under the baseline, from 46.6 for mT5-base to 85.9

for KaLM-mini, so the model choice dominates. ABTT compresses that to 3.3 points, from 86.1 to 89.4. Assignment accuracy tracks it, moving from a 40.3-point spread to the range 88.5 to 91.7, with every ABTT cell above every baseline cell.

Table 4 repeats the top-1 comparison under five query and reference reseeds, so the compression is not an artifact of one draw. The baseline spread of 36.6 points, from 50.6 for mT5-base to 87.3 for KaLM-mini, falls to 1.0 point under the best per-model SIF plus ABTT configuration, from 89.4 to 90.4, with standard deviations of at most 1.1 points.

The same training signal also picks the attribution layers. We take the earliest layer within 0.5 percentage points of the best training-set directory

Model	Assignment accuracy				Directory accuracy @1			
	Base	SIF	ABTT	SIF+ABTT	Base	SIF	ABTT	SIF+ABTT
LaTa	73.8 ₁	83.2 ₁	88.5 ₈	90.2 ₁₂	72.1 ₁	81.8 ₁	86.1 ₈	88.2 ₁₂
PhilTa	70.9 ₁	79.0 ₁	91.1 ₁	91.4 ₁	69.3 ₁	77.7 ₁	88.3 ₁	88.9 ₁
mT5-base	47.2 ₁₂	54.5 ₁	90.3 ₁	90.3 ₂	46.6 ₁₂	54.1 ₁	88.5 ₁	87.5 ₂
LaBSE	83.3 ₁₁	83.4 ₁₂	90.8 ₁₁	89.6 ₁₁	81.5 ₁₁	81.9 ₁₂	88.5 ₁₁	86.9 ₁₁
Qwen3-0.6B	82.3 ₂₈	76.0 ₂₇	91.5 ₅	90.8 ₅	80.3 ₂₈	74.2 ₂₇	89.4 ₅	88.6 ₅
KaLM-mini	87.5 ₂₃	85.7 ₂₂	91.7 ₃	90.8 ₁	85.9 ₂₃	84.0 ₂₂	89.4 ₃	88.0 ₁

Table 3: Task B autonomous routing for all six models under four post-processing settings, in percent. Each cell is a test-set score at the layer chosen by training-set directory accuracy at rank 1, given as the subscript. Assignment accuracy scores the existing-versus-new decision alone, comparing each file’s maximum cosine against the train-fit threshold τ ; a degenerate threshold that routes everything as existing already attains the 62.4 percent class prior. Directory accuracy at rank 1 requires the correct directory, so read the two columns together. Per-layer grids: Appendix Tables 8, 12, and 15.

accuracy at rank 1 under ABTT. The rule lands on LaTa layer 7, PhilTa layer 1, and mT5-base layer 1. We lock those choices before computing any attribution metric, so the attribution results cannot pick their own layer.

The cumulative view also describes the scholar-facing shortlist. By Top-2, every model exceeds 95%. Gains after Top-3 are small, so two or three options cover nearly all successful cases in this split.

5.4 Attribution at Retrieval-Selected Layers

Table 5 reports the main attribution comparison: 200 positive query-candidate pairs per model for LaTa, PhilTa, and mT5-base, with baseline and ABTT variants of both integrated gradients and retrieval-adapted MaRC. Figure 4 in Appendix F plots the same rank-faithfulness column. ABTT repairs the scores being explained at these layers: LaTa layer 7 moves from AUROC 0.499 to 0.961 and directory accuracy 0.302 to 0.868, PhilTa layer 1 from 0.939 to 0.976 and 0.693 to 0.883, and mT5-base layer 1 from 0.823 to 0.979 and 0.451 to 0.885.

The clearest quantitative result is rank faithfulness. ABTT improves ρ_{LOO} in all six model-method cells: LaTa IG $-0.001 \rightarrow 0.396$, LaTa MaRC $0.042 \rightarrow 0.401$, PhilTa IG $0.182 \rightarrow 0.607$, PhilTa MaRC $0.190 \rightarrow 0.331$, mT5-base IG $0.164 \rightarrow 0.597$, and mT5-base MaRC $0.250 \rightarrow 0.392$. Across these cells, tokens with higher attribution align more closely with leave-one-out drops in the same retrieval score.

The ERASER-style metrics qualify this result. At the global headline thresholds, ABTT improves only 6 of 18 Sufficiency, Comprehensiveness, and MinFrac comparisons. MinFrac@0.80 worsens in all six cells. We therefore read the attribution

evidence as support for the repair through rank faithfulness, not as a uniform gain across rationale metrics.

Appendix G shows example-level views for one PhilTa retrieval case, including the integrated-gradient pair matrix before and after ABTT.

6 Discussion

Six models across three families narrow what we can claim. Anisotropy is universal on this corpus and so is the repair, but the dramatic mid-depth collapse is not. Every model has a layer where passages crowd into a cone, and no model produces a cosine margin above 0.29 at any layer without correction. ABTT then moves all six into a 0.971 to 0.984 AUROC band and a 3.3-point routing band. Only the T5 encoders in our set turn their anisotropy into a retrieval collapse. Published layerwise profiles place the mid-depth anisotropy peak in decoder-only models and do not separate encoder-decoder encoders (Razzhigaev et al., 2024; Godey et al., 2024), so we read this as a first report for T5 encoders rather than a property of the architecture. A study of one architecture family could easily report that collapse as the general finding; three families show that the general finding is the geometry and the correction, while the collapse is a family-level symptom.

The practical reading depends on the family. With a T5 encoder the middle layers are unusable without correction, so layer choice matters. With a non-T5 model the final layers are already usable, and the correction still adds 3.5 to 9.1 points of directory accuracy at rank 1, at a cost of one SVD per layer fit on 847 training files.

Geometry diagnoses the failure mode but does not select layers. Use it to find where representations collapse and recover, then pick retrieval layers

Model	Top-1	Top-2	Top-3	Top-4	Top-5
LaTa	90.1 $\pm$ 1.1	96.0 $\pm$ 0.4	97.1 $\pm$ 0.1	97.6 $\pm$ 0.4	97.7 $\pm$ 0.4
PhilTa	90.0 $\pm$ 0.9	95.7 $\pm$ 0.4	96.9 $\pm$ 0.4	97.4 $\pm$ 0.3	97.7 $\pm$ 0.2
mT5-base	89.8 $\pm$ 0.4	95.6 $\pm$ 0.5	96.6 $\pm$ 0.3	96.9 $\pm$ 0.2	97.2 $\pm$ 0.2
LaBSE	90.4 $\pm$ 0.8	95.5 $\pm$ 0.4	96.4 $\pm$ 0.3	97.1 $\pm$ 0.5	97.7 $\pm$ 0.4
Qwen3-0.6B	90.3 $\pm$ 0.6	95.5 $\pm$ 0.5	96.8 $\pm$ 0.3	97.2 $\pm$ 0.5	97.4 $\pm$ 0.4
KaLM-mini	89.4 $\pm$ 0.5	95.1 $\pm$ 0.4	96.4 $\pm$ 0.2	96.9 $\pm$ 0.3	97.3 $\pm$ 0.4

Table 4: Cumulative top- K accuracy for Task B, as mean $\pm$ standard deviation over five query/reference reseeds at a fixed v2 split, using the best per-model SIF+ABTT configuration. Each cell is the percentage of test queries whose labelled directory appears within the top K options. Per-layer top- k rankings: Appendix Tables 9, 13, and 16.

Model	Method	ρ_{LOO} base	ρ_{LOO} ABTT	$\Delta\rho_{\text{LOO}}$	Suff@25% base $\rightarrow$ ABTT	Comp@25% base $\rightarrow$ ABTT	MinFrac@0.80 base $\rightarrow$ ABTT	ERASER wins
LaTa	IG	-0.001	0.396	+0.397	0.976 $\rightarrow$ 0.681	0.719 $\rightarrow$ 0.581	0.041 $\rightarrow$ 0.379	0/3
	MaRC	0.042	0.401	+0.359	0.534 $\rightarrow$ 0.645	0.361 $\rightarrow$ 0.539	0.383 $\rightarrow$ 0.459	2/3
PhilTa	IG	0.182	0.607	+0.425	0.959 $\rightarrow$ 0.963	0.969 $\rightarrow$ 0.747	0.079 $\rightarrow$ 0.141	1/3
	MaRC	0.190	0.331	+0.141	0.911 $\rightarrow$ 0.754	0.963 $\rightarrow$ 0.521	0.104 $\rightarrow$ 0.358	0/3
mT5-base	IG	0.164	0.597	+0.433	0.892 $\rightarrow$ 1.027	0.110 $\rightarrow$ 0.717	0.116 $\rightarrow$ 0.214	2/3
	MaRC	0.250	0.392	+0.142	0.913 $\rightarrow$ 0.827	0.042 $\rightarrow$ 0.592	0.066 $\rightarrow$ 0.346	1/3

Table 5: **Attribution metrics at the predeclared operational layers for LaTa, PhilTa, and mT5-base.** 200 positive query-candidate pairs per model, for integrated gradients (IG) and retrieval-adapted MaRC. The primary faithfulness signal is ρ_{LOO} , which improves under ABTT in 6/6 model-method cells. ERASER-style cells give baseline $\rightarrow$ ABTT at one global threshold choice: Sufficiency and Comprehensiveness at 25% of query tokens, MinFrac at recovery threshold 0.80 (lower is better for MinFrac, higher otherwise). Boldface marks the better variant, so disagreements stay visible: ABTT wins 6/18 ERASER-style comparisons even though ρ_{LOO} improves consistently. Cross-variant comparisons are descriptive because ABTT changes both the attribution scores and the cosine function being explained. Full sweeps are in Appendix F.

with a train-only task signal. In our attribution experiment the two answers differ: the selected layers are not the strongest diagnostic collapse layers.

The attribution results need a narrow reading. Because ABTT changes the cosine function being explained, cross-variant comparisons are descriptive: the stable result is improved leave-one-out rank faithfulness in all six cells, while ERASER-style metrics show threshold-dependent tradeoffs.

Limitations

The study covers one domain, medieval canon law, and one language, Latin. The 1,705-fragment corpus is large for this field, but small by current NLP standards. We test six models; others may behave differently. The mid-depth retrieval collapse we document holds for the three T5 encoders in our set, and we cannot separate architecture, pre-training objective, and checkpoint as its cause. ABTT also needs enough text to estimate principal components, so smaller archives need separate study.

Task B simulates a human-in-the-loop workflow by counting a query as successful when its labelled directory appears in the shortlist. That measures ranking quality, not the full judgment load a working scholar carries. Our attribution analysis is also

local to fixed query-candidate scores. It does not explain model behavior globally, and it does not replace expert reading of retrieved passages. The consistent attribution effect is improved leave-one-out rank faithfulness; other rationale metrics remain threshold-dependent and should not be reported as a broad attribution-quality gain.

References

- Mira Ait-Saada and Mohamed Nadif. 2023. [Is anisotropy truly harmful? A case study on text clustering](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 1194–1203. Association for Computational Linguistics.
- Avishek Anand, Lijun Lyu, Maximilian Idahl, Yumeng Wang, Jonas Wallat, and Zijian Zhang. 2022. [Explainable information retrieval: A survey](#). *arXiv preprint arXiv:2211.02405*.
- Sanjeev Arora, Yingyu Liang, and Tengyu Ma. 2017. [A simple but tough-to-beat baseline for sentence embeddings](#). In *Proceedings of the 5th International Conference on Learning Representations (ICLR)*.
- Pepa Atanasova, Jakob Grue Simonsen, Christina Lioma, and Isabelle Augenstein. 2020. [A diagnostic study of explainability techniques for text classification](#). In *Proceedings of the 2020 Conference on*

672	<i>Empirical Methods in Natural Language Processing (EMNLP)</i> , pages 3256–3274. Association for Computational Linguistics.	728
673		729
674		
675	David Bamman and Patrick J. Burns. 2020. Latin BERT: A contextual language model for classical philology . <i>arXiv preprint arXiv:2009.10053</i> .	730
676		731
677		732
678	Marc Brinner and Sina Zarri��. 2023. Model interpretability and rationale extraction by input mask optimization . In <i>Findings of the Association for Computational Linguistics: ACL 2023</i> , pages 13722–13744. Association for Computational Linguistics.	733
679		734
680		735
681		
682		
683	Patrick J. Burns, James A. Brofos, Kyle Li, Pramit Chaudhuri, and Joseph P. Dexter. 2021. Profiling of intertextuality in Latin literature using word embeddings . In <i>Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies</i> , pages 4900–4907. Association for Computational Linguistics.	736
684		737
685		738
686		739
687		740
688		741
689		
690		
691	Neil Coffee, Jean-Pierre Koenig, Shakthi Poornima, Roelant Ossewaarde, Christopher Forstall, and Sarah Jacobson. 2012. Intertextuality in the digital age . <i>Transactions of the American Philological Association</i> , 142(2):383–422.	742
692		743
693		744
694		745
695		746
696	Jay DeYoung, Sarthak Jain, Nazneen Fatema Rajani, Eric Lehman, Caiming Xiong, Richard Socher, and Byron C. Wallace. 2020. ERASER: A benchmark to evaluate rationalized NLP models . In <i>Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics</i> , pages 4443–4458, Online. Association for Computational Linguistics.	747
697		
698		
699		
700		
701		
702		
703	Melodie H. Eichbauer. 2014. Carolingian Canon Law Project Dir. by Abigail Firey (review) . <i>Digital Philology: A Journal of Medieval Cultures</i> , 3(1):165–168.	748
704		749
705		750
706	Kawin Ethayarajh. 2019. How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings . In <i>Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)</i> , pages 55–65. Association for Computational Linguistics.	751
707		752
708		753
709		754
710		
711		
712		
713		
714	Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. 2022. Language-agnostic BERT sentence embedding . In <i>Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 878–891. Association for Computational Linguistics.	755
715		756
716		757
717		758
718		759
719		760
720	Abigail Firey. 2009. The Carolingian Canon Law Project . Collaboratory for Research in Computing for the Humanities, University of Kentucky. Ongoing digital edition; transcription protocols and canon index published on the project site.	761
721		762
722		763
723		764
724		765
725	Jun Gao, Di He, Xu Tan, Tao Qin, Liwei Wang, and Tie-Yan Liu. 2019. Representation degeneration problem in training natural language generation models . In <i>Proceedings of the 7th International Conference on Learning Representations (ICLR)</i> .	766
726		767
727		768
		769
		770
		771
		772
		773
		774
		775
		776
		777
		778
		779
		780
		781
		782
		783
		784

785	<i>SIGHUM Workshop on Computational Linguistics for Cultural Heritage, Social Sciences, Humanities and Literature</i> , pages 104–114. Association for Computational Linguistics.	841
786		842
787		843
788		844
789	Leland McInnes, John Healy, and James Melville. 2018.	
790	UMAP: Uniform manifold approximation and projection for dimension reduction. <i>arXiv preprint arXiv:1802.03426</i> .	846
791		847
792		848
793	Jiaqi Mu and Pramod Viswanath. 2018. All-but-the-top: Simple and effective postprocessing for word representations . In <i>Proceedings of the 6th International Conference on Learning Representations (ICLR)</i> .	849
794		850
795		851
796		852
797	Niklas Muennighoff, Nouamane Tazi, Loic Magne, and Nils Reimers. 2023. MTEB: Massive text embedding benchmark . In <i>Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics</i> , pages 2014–2037. Association for Computational Linguistics.	853
798		854
799		855
800		856
801		857
802		858
803	Jianmo Ni, Gustavo Hernandez Abrego, Noah Constant, Ji Ma, Keith Hall, Daniel Cer, and Yinfei Yang. 2022. Sentence-T5: Scalable sentence encoders from pre-trained text-to-text models . In <i>Findings of the Association for Computational Linguistics: ACL 2022</i> , pages 1864–1874. Association for Computational Linguistics.	859
804		860
805		861
806		862
807		863
808		864
809		865
810	Sara Rajae and Mohammad Taher Pilehvar. 2022. An isotropy analysis in the multilingual BERT embedding space . In <i>Findings of the Association for Computational Linguistics: ACL 2022</i> , pages 1309–1316. Association for Computational Linguistics.	866
811		867
812		868
813		869
814		870
815	Anton Razzhigaev, Matvey Mikhalechuk, Elizaveta Goncharova, Ivan Oseledets, Denis Dimitrov, and Andrey Kuznetsov. 2024. The shape of learning: Anisotropy and intrinsic dimensions in transformer-based models . In <i>Findings of the Association for Computational Linguistics: EACL 2024</i> , pages 868–874. Association for Computational Linguistics.	871
816		872
817		873
818		874
819		875
820		876
821		877
822	Frederick Riemenschneider and Anette Frank. 2023a. Exploring large language models for classical philology . In <i>Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 15181–15199. Association for Computational Linguistics.	878
823		879
824		880
825		881
826		882
827		883
828	Frederick Riemenschneider and Anette Frank. 2023b. Graecia capta ferum victorem cepit. detecting Latin allusions to Ancient Greek literature . In <i>Proceedings of the Ancient Language Processing Workshop</i> , pages 30–38. INCOMA Ltd., Shoumen, Bulgaria.	884
829		885
830		886
831		887
832		888
833	Oscar Skea, Md Rifat Arefin, Dan Zhao, Niket Nikul Patel, Jalal Naghiyev, Yann LeCun, and Ravid Shwartz-Ziv. 2025. Layer by layer: Uncovering hidden representations in language models . In <i>Proceedings of the 42nd International Conference on Machine Learning (ICML)</i> , volume 267 of <i>Proceedings of Machine Learning Research</i> , pages 55854–55875. PMLR.	889
834		890
835		891
836		892
837		893
838		894
839		895
840		896
		897
	David A. Smith, Ryan Cordell, and Elizabeth Maddock Dillon. 2013. Infectious texts: Modeling text reuse in nineteenth-century newspapers . In <i>2013 IEEE International Conference on Big Data</i> , pages 86–94. IEEE.	
	Rachele Sprugnoli, Marco Passarotti, Flavio Massimiliano Cecchini, and Matteo Pellegrini. 2020. Overview of the EvaLatin 2020 evaluation campaign . In <i>Proceedings of LT4HALA 2020 - 1st Workshop on Language Technologies for Historical and Ancient Languages</i> , pages 105–110. European Language Resources Association (ELRA).	
	Jianlin Su, Jiarun Cao, Weijie Liu, and Yangyiwen Ou. 2021. Whitening sentence representations for better semantics and faster retrieval . In <i>Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)</i> . Association for Computational Linguistics.	
	Mukund Sundararajan, Ankur Taly, and Qiqi Yan. 2017. Axiomatic attribution for deep networks. In <i>Proceedings of the 34th International Conference on Machine Learning (ICML)</i> , pages 3319–3328. PMLR.	
	Ian Tenney, Dipanjan Das, and Ellie Pavlick. 2019. BERT rediscovers the classical NLP pipeline . In <i>Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics</i> , pages 4593–4601. Association for Computational Linguistics.	
	Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. 2021. BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models . In <i>Proceedings of the 35th Conference on Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track</i> .	
	William Timkey and Marten van Schijndel. 2021. All bark and no bite: Rogue dimensions in transformer language models obscure representational quality . In <i>Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing</i> , pages 4527–4546. Association for Computational Linguistics.	
	Laurens van der Maaten and Geoffrey Hinton. 2008. Visualizing data using t-SNE. <i>Journal of Machine Learning Research</i> , 9:2579–2605.	
	Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. 2021. mT5: A massively multilingual pre-trained text-to-text transformer . In <i>Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies</i> , pages 483–498. Association for Computational Linguistics.	
	Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, Huan Lin, Baosong Yang, Pengjun Xie, An Yang, Dayiheng Liu, Junyang Lin, Fei Huang, and Jingren	

Zhou. 2025. [Qwen3 embedding: Advancing text embedding and reranking through foundation models](#). *arXiv preprint arXiv:2506.05176*.

A Layer Diagnostics and Attribution Layer Selection

This appendix separates diagnostic layers from operational attribution layers. The diagnostic collapse layers identify where geometry is most anisotropic. The operational attribution layers are selected by training-set retrieval accuracy. Table 6 lists both for the three T5 encoders, together with the top-PC share and effective rank before and after ABTT.

Across the main-model layer grid, PC1 dominance and low effective rank are strongly associated with ABTT gains over baseline. These associations explain why ABTT helps, but they do not replace the train-only retrieval rule used for layer selection.

B Expanded Per-Layer Results for T5 Encoders

The main text compresses the headline retrieval results into Tables 2 and 3, which report one train-selected layer per model and method. This appendix provides the expanded per-layer Task A duplicate-detection, Task B autonomous-routing, and Task B top- k ranking tables for LaTa, PhilTa, and mT5-base.

Model	Operational layer	Diagnostic layer	PC1 before	PC1 after	Eff. rank before→after
LaTa	7	8	0.956	0.036	1.34→168.40
PhilTa	1	6	0.862	0.039	1.83→155.02
mT5-base	1	5	1.000	0.044	1.00→151.74

Table 6: Diagnostic collapse layers and operational attribution layers for the three T5 encoders. PC1 and effective-rank values compare baseline geometry with ABTT-D10 at the diagnostic collapse layer. Operational layers come from the train-only retrieval rule.

Task A: Pairwise Duplicate Detection (main)					
Model	Layer	AUROC (base)	AUROC (ABTT)	Cosine gap (base)	Cosine gap (ABTT)
LaTa	1	0.933	0.969	0.286	0.565
	2	0.565	0.969	0.020	0.553
	3	0.515	0.972	-0.046	0.555
	4	0.502	0.970	-0.065	0.564
	5	0.498	0.971	-0.068	0.572
	6	0.497	0.963	-0.067	0.569
	7	0.499	0.961	-0.070	0.564
	8	0.504	0.964	-0.073	0.558
	9	0.504	0.964	-0.071	0.557
	10	0.507	0.970	-0.070	0.551
	11	0.510	0.973	-0.065	0.550
	12	0.938	0.971	0.238	0.525
PhilTa	1	0.939	0.976	0.231	0.573
	2	0.911	0.971	0.161	0.551
	3	0.601	0.972	0.085	0.549
	4	0.551	0.976	0.064	0.551
	5	0.543	0.978	0.054	0.551
	6	0.541	0.980	0.048	0.551
	7	0.542	0.979	0.047	0.554
	8	0.543	0.983	0.047	0.555
	9	0.539	0.982	0.042	0.540
	10	0.539	0.982	0.040	0.530
	11	0.540	0.982	0.041	0.511
	12	0.911	0.981	0.181	0.487
mT5-base	1	0.823	0.979	0.027	0.550
	2	0.823	0.973	0.026	0.535
	3	0.805	0.968	0.025	0.517
	4	0.800	0.975	0.040	0.503
	5	0.654	0.977	0.218	0.482
	6	0.657	0.973	0.238	0.482
	7	0.657	0.979	0.243	0.482
	8	0.656	0.980	0.248	0.471
	9	0.660	0.981	0.235	0.475
	10	0.659	0.975	0.225	0.459
	11	0.660	0.971	0.213	0.449
	12	0.839	0.970	0.082	0.441

Table 7: Per-layer Task A pairwise duplicate-detection metrics for the three headline models (LaTa, PhilTa, mT5-base). baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. Both AUROC and cosine gap are computed over the test $n \times n$ cosine matrix with no threshold and no directory routing: AUROC treats same-directory test pairs as positives, and cosine gap is the difference between mean same-directory and mean different-directory cosine. Rows in bold mark the best layer per model, selected by ABTT AUROC.

Task B: Autonomous Routing (main, file-level, τ -thresholded)							
Model	Layer	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)	Overall (base)	Overall (ABTT)
LaTa	1	0.630	0.910	0.916	0.907	0.738	0.909
	2	0.153	0.905	0.913	0.892	0.439	0.900
	3	0.256	0.882	0.842	0.950	0.477	0.908
	4	0.389	0.893	0.718	0.920	0.513	0.903
	5	0.374	0.892	0.715	0.913	0.502	0.900
	6	0.378	0.884	0.712	0.913	0.503	0.895
	7	0.372	0.875	0.715	0.916	0.501	0.890
	8	0.379	0.884	0.706	0.885	0.502	0.885
	9	0.366	0.886	0.724	0.873	0.501	0.881
	10	0.344	0.873	0.752	0.882	0.498	0.876
	11	0.310	0.867	0.786	0.901	0.490	0.880
	12	0.576	0.890	0.932	0.833	0.710	0.868
PhilTa	1	0.591	0.918	0.904	0.901	0.709	0.911
	2	0.479	0.895	0.867	0.885	0.625	0.892
	3	0.193	0.897	0.889	0.879	0.455	0.890
	4	0.148	0.886	0.923	0.873	0.439	0.881
	5	0.252	0.879	0.796	0.873	0.457	0.876
	6	0.338	0.905	0.700	0.817	0.474	0.872
	7	0.301	0.901	0.746	0.830	0.469	0.874
	8	0.262	0.875	0.808	0.885	0.467	0.879
	9	0.303	0.890	0.762	0.870	0.476	0.882
	10	0.299	0.908	0.780	0.793	0.480	0.865
	11	0.256	0.869	0.824	0.836	0.470	0.857
	12	0.338	0.880	0.898	0.876	0.549	0.879
mT5-base	1	0.135	0.877	0.981	0.947	0.453	0.903
	2	0.110	0.890	0.975	0.910	0.436	0.897
	3	0.161	0.794	0.926	0.941	0.449	0.850
	4	0.269	0.787	0.858	0.882	0.491	0.823
	5	0.422	0.746	0.659	0.864	0.512	0.790
	6	0.636	0.755	0.489	0.836	0.580	0.786
	7	0.594	0.793	0.517	0.851	0.565	0.815
	8	0.557	0.809	0.560	0.842	0.558	0.822
	9	0.364	0.841	0.746	0.814	0.508	0.831
	10	0.314	0.794	0.783	0.876	0.491	0.825
	11	0.297	0.763	0.780	0.854	0.479	0.797
	12	0.198	0.834	0.926	0.799	0.472	0.821

Table 8: Per-layer Task B autonomous routing accuracy for the three headline models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. For every test file we compute $s_i = \max_{j \neq i} \cos(e_i, e_j)$ against the rest of the test set and compare to the learned threshold τ (fit on train via F_1 -optimal cut on same-directory vs. different-directory pairs). The decision is binary first (existing if $s_i \geq \tau$, else new); only files predicted existing are then routed to their top-1 neighbour’s directory. We report **existing accuracy** (restricted to files whose true partner is in the test set), **new accuracy** (files that should be flagged as novel), and **overall assignment accuracy**. Rows in bold mark the best layer per model, selected by ABTT overall assignment accuracy.

Task B: Top-K Ranking (main, single-seed v2)							
Model	Layer	Acc@1 (base)	Acc@1 (ABTT)	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)
LaTa	1	0.721	0.880	0.630	0.910	0.916	0.907
	2	0.383	0.869	0.153	0.905	0.913	0.892
	3	0.350	0.882	0.256	0.882	0.842	0.950
	4	0.303	0.881	0.389	0.893	0.718	0.920
	5	0.297	0.875	0.374	0.892	0.715	0.913
	6	0.297	0.874	0.378	0.884	0.712	0.913
	7	0.302	0.868	0.372	0.875	0.715	0.916
	8	0.302	0.861	0.379	0.884	0.706	0.885
	9	0.309	0.858	0.366	0.886	0.724	0.873
	10	0.317	0.852	0.344	0.873	0.752	0.882
	11	0.331	0.857	0.310	0.867	0.786	0.901
	12	0.689	0.841	0.576	0.890	0.932	0.833
PhilTa	1	0.693	0.883	0.591	0.918	0.904	0.901
	2	0.603	0.869	0.479	0.895	0.867	0.885
	3	0.389	0.865	0.193	0.897	0.889	0.879
	4	0.369	0.859	0.148	0.886	0.923	0.873
	5	0.334	0.854	0.252	0.879	0.796	0.873
	6	0.310	0.845	0.338	0.905	0.700	0.817
	7	0.325	0.847	0.301	0.901	0.746	0.830
	8	0.350	0.857	0.262	0.875	0.808	0.885
	9	0.337	0.857	0.303	0.890	0.762	0.870
	10	0.344	0.833	0.299	0.908	0.780	0.793
	11	0.355	0.829	0.256	0.869	0.824	0.836
	12	0.528	0.854	0.338	0.880	0.898	0.876
mT5-base	1	0.451	0.885	0.135	0.877	0.981	0.947
	2	0.435	0.874	0.110	0.890	0.975	0.910
	3	0.445	0.830	0.161	0.794	0.926	0.941
	4	0.467	0.798	0.269	0.787	0.858	0.882
	5	0.352	0.756	0.422	0.746	0.659	0.864
	6	0.346	0.746	0.636	0.755	0.489	0.836
	7	0.352	0.775	0.594	0.793	0.517	0.851
	8	0.366	0.784	0.557	0.809	0.560	0.842
	9	0.394	0.787	0.364	0.841	0.746	0.814
	10	0.387	0.788	0.314	0.794	0.783	0.876
	11	0.378	0.765	0.297	0.763	0.780	0.854
	12	0.466	0.779	0.198	0.834	0.926	0.799

Table 9: Per-layer Task B top- k ranking metrics for the three headline models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. This table is computed on the single-seed v2 split (no M -seed averaging, since the mseed sweep was run only for SIF-conditioned variants; see Table 16 for the multi-seed SIF+ABTT view). **Acc@1** is directory accuracy at rank 1, with an existing/new decomposition. Rows in bold mark the best layer per model, selected by ABTT Acc@1.

C Last-Token vs. Mean Pooling

Decoder-only embedding models often use last-token pooling. We compare last-token pooling with mean pooling under the same ABTT pipeline, train/test split, and ABTT choices. Table 10 reports the layer with the best Task B autonomous-routing accuracy, Task B directory accuracy at rank 1, and cosine gap at the same layer. Mean pooling performs better on assignment accuracy for every model in this experiment, including the decoder-only models.

D Per-Layer Results for Non-T5 Models

The main text focuses on LaTa, PhilTa, and mT5-base. This appendix reports the same per-layer view for LaBSE, Qwen3-0.6B, and KaLM-mini under the same comparison: baseline mean pooling versus ABTT with D tuned per layer on the train split.

Table 10: Mean vs. last-token pooling on the balanced split, restricted to ABTT post-processing. Best layer per (model, pooling) selected by Task A assignment accuracy. Task B metrics (Acc@1, cosine gap) reported at that same layer.

Model	Pool	Layer	Method	Assign Acc	Acc@1	Gap
LaTa	Mean	1	ABTT (D=10)	0.909	0.880	0.565
LaTa	LastTok	12	ABTT (D=10)	0.676	0.425	0.276
PhilTa	Mean	1	ABTT (D=10)	0.911	0.883	0.573
PhilTa	LastTok	9	ABTT (D=10)	0.604	0.385	0.200
LaBSE	Mean	11	ABTT (D=10)	0.908	0.885	0.610
LaBSE	LastTok	12	ABTT (D=10)	0.895	0.871	0.586
Qwen3-0.6B	Mean	5	ABTT (D=10)	0.915	0.894	0.545
Qwen3-0.6B	LastTok	27	ABTT (D=10)	0.907	0.885	0.562
KaLM-mini	Mean	1	ABTT (D=10)	0.921	0.902	0.568
KaLM-mini	LastTok	24	ABTT (D=10)	0.875	0.847	0.483
Qwen3-8B	Mean	1	ABTT (D=10)	0.924	0.897	0.543
Qwen3-8B	LastTok	35	ABTT (D=10)	0.910	0.892	0.642

Table 11: Per-layer Task A pairwise metrics for the remaining three models (LaBSE, Qwen3-0.6B, KaLM-mini), under the same two-method comparison as the main paper. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. Both AUROC and cosine gap are computed over the test $n \times n$ cosine matrix with no threshold and no directory routing: AUROC treats same-directory test pairs as positives, and cosine gap is the difference between mean same-directory and mean different-directory cosine. Rows in bold mark the best layer per model, selected by ABTT AUROC.

Task A: Pairwise Duplicate Detection (appendix models)					
Model	Layer	AUROC (base)	AUROC (ABTT)	Cosine gap (base)	Cosine gap (ABTT)
LaBSE	1	0.807	0.976	0.060	0.519
	2	0.835	0.978	0.038	0.531
	3	0.855	0.977	0.063	0.543
	4	0.864	0.976	0.074	0.525
	5	0.866	0.967	0.066	0.500
	6	0.895	0.974	0.078	0.532
	7	0.889	0.973	0.089	0.523
	8	0.881	0.973	0.080	0.522
	9	0.918	0.976	0.091	0.564
	10	0.936	0.980	0.103	0.584
	11	0.961	0.984	0.061	0.610
	12	0.957	0.979	0.183	0.592
Qwen3-0.6B	1	0.859	0.978	0.033	0.564
	2	0.864	0.972	0.024	0.553
	3	0.864	0.980	0.037	0.555
	4	0.861	0.985	0.036	0.546
	5	0.862	0.980	0.030	0.545
	6	0.868	0.983	0.035	0.545
	7	0.877	0.980	0.033	0.544
	8	0.884	0.972	0.037	0.533
	9	0.882	0.977	0.034	0.532
	10	0.888	0.976	0.031	0.530
	11	0.896	0.970	0.030	0.516
	12	0.903	0.973	0.027	0.531
	13	0.908	0.974	0.026	0.524
	14	0.906	0.974	0.024	0.519
	15	0.902	0.971	0.022	0.514
	16	0.911	0.975	0.026	0.509
	17	0.917	0.978	0.032	0.521
	18	0.934	0.982	0.032	0.537
	19	0.937	0.980	0.035	0.541
	20	0.943	0.982	0.041	0.559
	21	0.953	0.982	0.047	0.562
	22	0.963	0.983	0.044	0.583
	23	0.960	0.983	0.038	0.578
	24	0.962	0.984	0.032	0.576
	25	0.960	0.983	0.027	0.578
	26	0.965	0.985	0.024	0.568
	27	0.966	0.982	0.017	0.559
	28	0.958	0.975	0.098	0.556
KaLM-mini	1	0.888	0.980	0.064	0.568
	2	0.908	0.983	0.044	0.566
	3	0.882	0.980	0.062	0.552
	4	0.872	0.983	0.060	0.553
	5	0.862	0.982	0.062	0.541
	6	0.873	0.982	0.062	0.528
	7	0.889	0.977	0.072	0.524
	8	0.891	0.975	0.072	0.523
	9	0.893	0.972	0.066	0.507
	10	0.893	0.970	0.067	0.494
	11	0.895	0.969	0.061	0.474
	12	0.905	0.968	0.071	0.475
	13	0.901	0.966	0.066	0.467
	14	0.904	0.970	0.063	0.471
	15	0.907	0.972	0.067	0.468
	16	0.917	0.970	0.071	0.488

(continued on next page)

Task A: Pairwise Duplicate Detection (appendix models)					
Model	Layer	AUROC (base)	AUROC (ABTT)	Cosine gap (base)	Cosine gap (ABTT)
	17	0.926	0.973	0.075	0.501
	18	0.944	0.976	0.060	0.519
	19	0.955	0.978	0.061	0.530
	20	0.959	0.979	0.054	0.554
	21	0.961	0.980	0.043	0.554
	22	0.969	0.984	0.050	0.554
	23	0.970	0.985	0.056	0.552
	24	0.954	0.981	0.243	0.520

Table 12: Per-layer Task B autonomous routing accuracy for the remaining three models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. For every test file we compute $s_i = \max_{j \neq i} \cos(e_i, e_j)$ against the rest of the test set and compare to the learned threshold τ (fit on train via F_1 -optimal cut on same-directory vs. different-directory pairs). The decision is binary first (existing if $s_i \geq \tau$, else new); only files predicted existing are then routed to their top-1 neighbour’s directory. We report **existing accuracy** (restricted to files whose true partner is in the test set), **new accuracy** (files that should be flagged as novel), and **overall assignment accuracy**. Rows in bold mark the best layer per model, selected by ABTT overall assignment accuracy.

Task B: Autonomous Routing (appendix models)							
Model	Layer	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)	Overall (base)	Overall (ABTT)
LaBSE	1	0.222	0.879	0.898	0.873	0.477	0.876
	2	0.381	0.875	0.811	0.827	0.543	0.857
	3	0.277	0.886	0.907	0.867	0.514	0.879
	4	0.372	0.897	0.885	0.836	0.565	0.874
	5	0.327	0.845	0.901	0.851	0.543	0.847
	6	0.548	0.860	0.882	0.851	0.674	0.857
	7	0.398	0.849	0.892	0.830	0.584	0.841
	8	0.385	0.869	0.870	0.743	0.568	0.822
	9	0.421	0.871	0.929	0.817	0.612	0.851
	10	0.606	0.892	0.892	0.805	0.713	0.859
	11	0.785	0.923	0.913	0.882	0.833	0.908
	12	0.729	0.914	0.938	0.864	0.808	0.895
Qwen3-0.6B	1	0.264	0.921	0.960	0.885	0.526	0.908
	2	0.437	0.869	0.814	0.916	0.579	0.887
	3	0.250	0.914	0.947	0.907	0.513	0.911
	4	0.411	0.936	0.848	0.842	0.576	0.901
	5	0.181	0.910	0.932	0.923	0.464	0.915
	6	0.422	0.935	0.830	0.796	0.576	0.882
	7	0.250	0.890	0.947	0.876	0.513	0.885
	8	0.409	0.931	0.895	0.799	0.592	0.881
	9	0.460	0.901	0.851	0.851	0.607	0.882
	10	0.533	0.892	0.836	0.864	0.647	0.881
	11	0.293	0.895	0.978	0.814	0.551	0.865
	12	0.407	0.912	0.938	0.817	0.607	0.876
	13	0.432	0.908	0.944	0.836	0.625	0.881
	14	0.456	0.907	0.926	0.811	0.633	0.871
	15	0.469	0.901	0.920	0.845	0.639	0.880
	16	0.639	0.897	0.830	0.808	0.711	0.864
	17	0.480	0.886	0.935	0.851	0.652	0.873
	18	0.553	0.882	0.944	0.839	0.700	0.866
	19	0.514	0.871	0.966	0.858	0.684	0.866
	20	0.654	0.865	0.898	0.892	0.746	0.875
	21	0.647	0.912	0.950	0.898	0.761	0.907
	22	0.731	0.888	0.938	0.916	0.809	0.899
	23	0.622	0.914	0.966	0.895	0.752	0.907
	24	0.748	0.901	0.935	0.935	0.818	0.914
	25	0.574	0.893	0.975	0.923	0.725	0.904
	26	0.613	0.899	0.975	0.895	0.749	0.897
	27	0.793	0.877	0.833	0.923	0.808	0.894

(continued on next page)

Task B: Autonomous Routing (appendix models)							
Model	Layer	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)	Overall (base)	Overall (ABTT)
	28	0.796	0.877	0.867	0.913	0.823	0.890
KaLM-mini	1	0.391	0.921	0.907	0.920	0.585	0.921
	2	0.493	0.916	0.848	0.879	0.627	0.902
	3	0.443	0.923	0.870	0.907	0.604	0.917
	4	0.308	0.931	0.935	0.895	0.544	0.917
	5	0.338	0.905	0.882	0.923	0.543	0.911
	6	0.256	0.895	0.910	0.920	0.502	0.904
	7	0.475	0.864	0.854	0.944	0.618	0.894
	8	0.467	0.862	0.870	0.913	0.619	0.881
	9	0.469	0.873	0.885	0.870	0.626	0.872
	10	0.449	0.875	0.885	0.870	0.613	0.873
	11	0.458	0.862	0.885	0.864	0.619	0.862
	12	0.507	0.862	0.895	0.904	0.653	0.878
	13	0.540	0.860	0.895	0.870	0.674	0.864
	14	0.593	0.836	0.845	0.885	0.688	0.854
	15	0.566	0.864	0.898	0.827	0.691	0.850
	16	0.550	0.849	0.935	0.830	0.695	0.841
	17	0.675	0.864	0.920	0.889	0.767	0.873
	18	0.697	0.875	0.960	0.898	0.796	0.883
	19	0.721	0.873	0.966	0.892	0.814	0.880
	20	0.705	0.791	0.966	0.913	0.803	0.837
	21	0.727	0.811	0.966	0.926	0.817	0.854
	22	0.850	0.879	0.916	0.923	0.875	0.895
	23	0.849	0.884	0.920	0.898	0.875	0.889
	24	0.787	0.892	0.944	0.867	0.846	0.882

Table 13: Per-layer Task B top- k ranking metrics for the remaining three models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. This table is computed on the single-seed v2 split (no M -seed averaging, since the mseed sweep was run only for SIF-conditioned variants; see Table 16 for the multi-seed SIF+ABTT view). **Acc@1** is directory accuracy at rank 1, with an existing/new decomposition. Rows in bold mark the best layer per model, selected by ABTT Acc@1.

Task B: Top-K Ranking (appendix models, single-seed)							
Model	Layer	Acc@1 (base)	Acc@1 (ABTT)	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)
LaBSE	1	0.465	0.851	0.222	0.879	0.898	0.873
	2	0.494	0.830	0.381	0.875	0.811	0.827
	3	0.500	0.848	0.277	0.886	0.907	0.867
	4	0.542	0.841	0.372	0.897	0.885	0.836
	5	0.521	0.811	0.327	0.845	0.901	0.851
	6	0.640	0.831	0.548	0.860	0.882	0.851
	7	0.561	0.810	0.398	0.849	0.892	0.830
	8	0.544	0.779	0.385	0.869	0.870	0.743
	9	0.608	0.823	0.421	0.871	0.929	0.817
	10	0.693	0.832	0.606	0.892	0.892	0.805
	11	0.815	0.885	0.785	0.923	0.913	0.882
	12	0.791	0.868	0.729	0.914	0.938	0.864
Qwen3-0.6B	1	0.521	0.875	0.264	0.921	0.960	0.885
	2	0.558	0.871	0.437	0.869	0.814	0.916
	3	0.508	0.890	0.250	0.914	0.947	0.907
	4	0.559	0.874	0.411	0.936	0.848	0.842
	5	0.450	0.894	0.181	0.910	0.932	0.923
	6	0.551	0.858	0.422	0.935	0.830	0.796
	7	0.507	0.859	0.250	0.890	0.947	0.876
	8	0.580	0.852	0.409	0.931	0.895	0.799
	9	0.587	0.852	0.460	0.901	0.851	0.851
	10	0.622	0.848	0.533	0.892	0.836	0.864
	11	0.548	0.833	0.293	0.895	0.978	0.814
	12	0.603	0.848	0.407	0.912	0.938	0.817
	13	0.618	0.851	0.432	0.908	0.944	0.836

(continued on next page)

Task B: Top-K Ranking (appendix models, single-seed)							
Model	Layer	Acc@1 (base)	Acc@1 (ABTT)	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)
	14	0.621	0.840	0.456	0.907	0.926	0.811
	15	0.627	0.848	0.469	0.901	0.920	0.845
	16	0.682	0.830	0.639	0.897	0.830	0.808
	17	0.641	0.840	0.480	0.886	0.935	0.851
	18	0.691	0.837	0.553	0.882	0.944	0.839
	19	0.676	0.837	0.514	0.871	0.966	0.858
	20	0.733	0.852	0.654	0.865	0.898	0.892
	21	0.752	0.881	0.647	0.912	0.950	0.898
	22	0.795	0.876	0.731	0.888	0.938	0.916
	23	0.742	0.886	0.622	0.914	0.966	0.895
	24	0.808	0.896	0.748	0.901	0.935	0.935
	25	0.718	0.886	0.574	0.893	0.975	0.923
	26	0.740	0.881	0.613	0.899	0.975	0.895
	27	0.787	0.876	0.793	0.877	0.833	0.923
	28	0.803	0.872	0.796	0.877	0.867	0.913
KaLM-mini	1	0.577	0.902	0.391	0.921	0.907	0.920
	2	0.607	0.879	0.493	0.916	0.848	0.879
	3	0.587	0.894	0.443	0.923	0.870	0.907
	4	0.538	0.887	0.308	0.931	0.935	0.895
	5	0.524	0.882	0.338	0.905	0.882	0.923
	6	0.490	0.876	0.256	0.895	0.910	0.920
	7	0.592	0.873	0.475	0.864	0.854	0.944
	8	0.597	0.860	0.467	0.862	0.870	0.913
	9	0.601	0.843	0.469	0.873	0.885	0.870
	10	0.584	0.844	0.449	0.875	0.885	0.870
	11	0.586	0.836	0.458	0.862	0.885	0.864
	12	0.622	0.850	0.507	0.862	0.895	0.904
	13	0.633	0.840	0.540	0.860	0.895	0.870
	14	0.636	0.829	0.593	0.836	0.845	0.885
	15	0.650	0.817	0.566	0.864	0.898	0.827
	16	0.676	0.819	0.550	0.849	0.935	0.830
	17	0.747	0.852	0.675	0.864	0.920	0.889
	18	0.777	0.865	0.697	0.875	0.960	0.898
	19	0.801	0.861	0.721	0.873	0.966	0.892
	20	0.789	0.817	0.705	0.791	0.966	0.913
	21	0.802	0.831	0.727	0.811	0.966	0.926
	22	0.857	0.875	0.850	0.879	0.916	0.923
	23	0.859	0.872	0.849	0.884	0.920	0.898
	24	0.833	0.858	0.787	0.892	0.944	0.867

E SIF-Variant Post-Processing Sweeps

The main text uses ABTT-only as the paper-facing repair. This appendix reports SIF-only, SIF plus fixed- D ABTT, and SIF plus tuned- D ABTT. The tables keep one primary metric per task so the four-method grid stays readable.

Table 14: Per-layer Task A AUROC across the SIF-conditioned post-processing suite for all six models. `sif_only` replaces mean pooling with SIF-weighted pooling; `sif_abtt_fixed` adds ABTT with fixed $D=10$; `sif_abtt_optimal` tunes D per layer on the train split. Gap is omitted to keep the table narrow; the pure-ABTT comparison (not SIF-conditioned) is in Tables 7 and 11. Rows in bold mark the best layer per model, selected by SIF+ABTT (D-tuned) AUROC.

Model	Layer	Task A: SIF-suite AUROC (appendix)			
		AUROC (base)	AUROC (SIF)	AUROC (SIF+ABTT ₁₀)	AUROC (SIF+ABTT)
LaTa	1	0.933	0.944	0.971	0.974
	2	0.565	0.936	0.964	0.970
	3	0.515	0.920	0.966	0.964
	4	0.502	0.890	0.969	0.968
	5	0.498	0.878	0.969	0.969
	6	0.497	0.865	0.964	0.964
	7	0.499	0.858	0.965	0.965
	8	0.504	0.842	0.964	0.964
	9	0.504	0.854	0.964	0.964
	10	0.507	0.859	0.965	0.965
	11	0.510	0.867	0.965	0.965
	12	0.938	0.948	0.963	0.963
PhilTa	1	0.939	0.953	0.973	0.973
	2	0.911	0.941	0.970	0.970
	3	0.601	0.929	0.966	0.966
	4	0.551	0.916	0.970	0.970
	5	0.543	0.904	0.971	0.976
	6	0.541	0.884	0.980	0.980
	7	0.542	0.891	0.980	0.980
	8	0.543	0.895	0.981	0.981
	9	0.539	0.882	0.977	0.977
	10	0.539	0.874	0.977	0.977
	11	0.540	0.872	0.978	0.978
	12	0.911	0.929	0.968	0.968
mT5-base	1	0.823	0.846	0.971	0.971
	2	0.823	0.835	0.964	0.964
	3	0.805	0.809	0.956	0.956
	4	0.800	0.810	0.966	0.966
	5	0.654	0.672	0.972	0.972
	6	0.657	0.679	0.970	0.970
	7	0.657	0.680	0.966	0.966
	8	0.656	0.682	0.967	0.967
	9	0.660	0.685	0.966	0.966
	10	0.659	0.685	0.964	0.964
	11	0.660	0.690	0.964	0.964
	12	0.839	0.834	0.968	0.968
LaBSE	1	0.807	0.819	0.974	0.974
	2	0.835	0.842	0.974	0.974
	3	0.855	0.864	0.968	0.968
	4	0.864	0.871	0.966	0.966
	5	0.866	0.870	0.963	0.963
	6	0.895	0.897	0.964	0.964
	7	0.889	0.893	0.961	0.961
	8	0.881	0.885	0.966	0.966
	9	0.918	0.916	0.971	0.971
	10	0.936	0.936	0.974	0.974
	11	0.961	0.959	0.980	0.980
	12	0.957	0.956	0.978	0.982
Qwen3-0.6B	1	0.859	0.852	0.982	0.980
	2	0.864	0.854	0.972	0.972
	3	0.864	0.859	0.977	0.975
	4	0.861	0.861	0.976	0.977
	5	0.862	0.859	0.971	0.971
	6	0.868	0.862	0.974	0.974
	7	0.877	0.867	0.971	0.971
	8	0.884	0.869	0.969	0.969

(continued on next page)

Task A: SIF-suite AUROC (appendix)						
Model	Layer	AUROC (base)	AUROC (SIF)	AUROC (SIF+ABTT ₁₀)	AUROC (SIF+ABTT)	
	9	0.882	0.867	0.968	0.968	
	10	0.888	0.871	0.966	0.966	
	11	0.896	0.876	0.964	0.964	
	12	0.903	0.882	0.967	0.967	
	13	0.908	0.888	0.967	0.967	
	14	0.906	0.887	0.966	0.966	
	15	0.902	0.884	0.967	0.967	
	16	0.911	0.895	0.965	0.965	
	17	0.917	0.900	0.967	0.974	
	18	0.934	0.915	0.966	0.972	
	19	0.937	0.916	0.966	0.976	
	20	0.943	0.919	0.971	0.971	
	21	0.953	0.931	0.971	0.971	
	22	0.963	0.941	0.973	0.973	
	23	0.960	0.941	0.973	0.973	
	24	0.962	0.942	0.975	0.973	
	25	0.960	0.938	0.978	0.974	
	26	0.965	0.941	0.979	0.979	
	27	0.966	0.935	0.977	0.977	
	28	0.958	0.925	0.967	0.967	
	KaLM-mini	1	0.888	0.881	0.981	0.979
		2	0.908	0.900	0.981	0.981
		3	0.882	0.882	0.977	0.977
		4	0.872	0.873	0.974	0.974
		5	0.862	0.863	0.971	0.971
		6	0.873	0.867	0.966	0.966
		7	0.889	0.879	0.971	0.971
		8	0.891	0.880	0.971	0.975
9		0.893	0.883	0.966	0.966	
10		0.893	0.883	0.963	0.963	
11		0.895	0.886	0.968	0.968	
12		0.905	0.897	0.965	0.965	
13		0.901	0.894	0.963	0.963	
14		0.904	0.898	0.968	0.968	
15		0.907	0.901	0.968	0.968	
16		0.917	0.913	0.966	0.966	
17		0.926	0.923	0.968	0.968	
18		0.944	0.938	0.976	0.976	
19		0.955	0.949	0.978	0.978	
20		0.959	0.955	0.978	0.978	
21		0.961	0.957	0.979	0.982	
22		0.969	0.965	0.982	0.982	
23		0.970	0.963	0.986	0.984	
24		0.954	0.954	0.975	0.975	

Table 15: Per-layer Task B overall assignment accuracy across the SIF-conditioned suite for all six models. Restricted to overall routing accuracy (existing/new decomposition omitted) to keep the table compact; see Tables 8 and 12 for the pure-ABTT pairwise comparison with existing/new broken out. Rows in bold mark the best layer per model, selected by SIF+ABTT (D-tuned) overall assignment accuracy.

Task B: SIF-suite routing (appendix)					
Model	Layer	Overall (base)	Overall (SIF)	Overall (SIF+ABTT ₁₀)	Overall (SIF+ABTT)
LaTa	1	0.738	0.832	0.920	0.899
	2	0.439	0.791	0.904	0.900
	3	0.477	0.720	0.910	0.907
	4	0.513	0.668	0.900	0.902
	5	0.502	0.648	0.906	0.906
	6	0.503	0.668	0.895	0.895
	7	0.501	0.642	0.903	0.903
	8	0.502	0.614	0.892	0.892

(continued on next page)

Task B: SIF-suite routing (appendix)					
Model	Layer	Overall (base)	Overall (SIF)	Overall (SIF+ABTT ₁₀)	Overall (SIF+ABTT)
	9	0.501	0.618	0.889	0.889
	10	0.498	0.655	0.890	0.890
	11	0.490	0.621	0.880	0.880
	12	0.710	0.804	0.902	0.902
PhilTa	1	0.709	0.790	0.914	0.914
	2	0.625	0.738	0.892	0.892
	3	0.455	0.713	0.892	0.892
	4	0.439	0.663	0.886	0.886
	5	0.457	0.626	0.894	0.897
	6	0.474	0.612	0.906	0.906
	7	0.469	0.625	0.906	0.906
	8	0.467	0.614	0.886	0.886
	9	0.476	0.590	0.895	0.895
	10	0.480	0.580	0.888	0.888
	11	0.470	0.565	0.886	0.886
	12	0.549	0.667	0.899	0.899
mT5-base	1	0.453	0.545	0.907	0.907
	2	0.436	0.537	0.903	0.903
	3	0.449	0.568	0.889	0.889
	4	0.491	0.488	0.857	0.857
	5	0.512	0.483	0.826	0.826
	6	0.580	0.558	0.817	0.817
	7	0.565	0.554	0.832	0.832
	8	0.558	0.548	0.834	0.834
	9	0.508	0.495	0.845	0.845
	10	0.491	0.495	0.832	0.832
	11	0.479	0.484	0.816	0.816
	12	0.472	0.541	0.858	0.858
LaBSE	1	0.477	0.521	0.886	0.886
	2	0.543	0.508	0.875	0.875
	3	0.514	0.577	0.860	0.860
	4	0.565	0.587	0.869	0.869
	5	0.543	0.598	0.857	0.857
	6	0.674	0.648	0.872	0.872
	7	0.584	0.634	0.858	0.858
	8	0.568	0.571	0.854	0.854
	9	0.612	0.613	0.859	0.859
	10	0.713	0.698	0.864	0.864
	11	0.833	0.807	0.896	0.896
	12	0.808	0.834	0.890	0.880
Qwen3-0.6B	1	0.526	0.566	0.895	0.909
	2	0.579	0.524	0.894	0.894
	3	0.513	0.537	0.904	0.901
	4	0.576	0.531	0.899	0.904
	5	0.464	0.537	0.908	0.908
	6	0.576	0.526	0.906	0.906
	7	0.513	0.573	0.911	0.911
	8	0.592	0.627	0.911	0.911
	9	0.607	0.576	0.910	0.910
	10	0.647	0.599	0.900	0.900
	11	0.551	0.606	0.893	0.893
	12	0.607	0.663	0.900	0.900
	13	0.625	0.675	0.897	0.897
	14	0.633	0.571	0.893	0.893
	15	0.639	0.600	0.893	0.893
	16	0.711	0.679	0.880	0.880
	17	0.652	0.686	0.900	0.876
	18	0.700	0.731	0.899	0.882
	19	0.684	0.700	0.886	0.850
	20	0.746	0.676	0.882	0.882
	21	0.761	0.739	0.892	0.892
	22	0.809	0.719	0.893	0.893

(continued on next page)

Task B: SIF-suite routing (appendix)					
Model	Layer	Overall (base)	Overall (SIF)	Overall (SIF+ABTT ₁₀)	Overall (SIF+ABTT)
KaLM-mini	23	0.752	0.759	0.893	0.893
	24	0.818	0.706	0.900	0.899
	25	0.725	0.748	0.897	0.893
	26	0.749	0.772	0.903	0.903
	27	0.808	0.760	0.899	0.899
	28	0.823	0.774	0.890	0.890
	1	0.585	0.604	0.906	0.908
	2	0.627	0.594	0.908	0.908
	3	0.604	0.576	0.907	0.907
	4	0.544	0.577	0.901	0.901
	5	0.543	0.533	0.897	0.897
	6	0.502	0.558	0.900	0.900
	7	0.618	0.593	0.901	0.901
	8	0.619	0.607	0.892	0.878
	9	0.626	0.663	0.876	0.876
	10	0.613	0.603	0.866	0.866
	11	0.619	0.627	0.874	0.874
	12	0.653	0.642	0.874	0.874
	13	0.674	0.666	0.864	0.864
	14	0.688	0.648	0.861	0.861
	15	0.691	0.698	0.866	0.866
	16	0.695	0.738	0.859	0.859
	17	0.767	0.763	0.882	0.882
	18	0.796	0.796	0.879	0.879
	19	0.814	0.804	0.878	0.878
	20	0.803	0.821	0.876	0.876
	21	0.817	0.832	0.874	0.876
	22	0.875	0.857	0.876	0.876
	23	0.875	0.862	0.880	0.881
	24	0.846	0.847	0.869	0.869

Table 16: Per-layer Task B top- k ranking accuracy across all six models, averaged over $M = 5$ query/reference reseeds of the v2 train/test split (mean $\pm$ std). `sif_abtt_optimal` applies SIF weighting plus ABTT with D tuned per layer on the train split; the mseed sweep was not run for the pure `abtt_optimal` variant reported in the main paper, so we include this table to give a variance-aware view of the SIF-conditioned best variant. Overall assignment accuracy coincides with Acc@1 in the mseed pipeline, so we report the existing/new decomposition instead. Rows in bold mark the best layer per model, selected by SIF+ABTT Acc@1.

Task B: Top-K Ranking (appendix, 5-seed mean $\pm$ std, SIF+ABTT)							
Model	Layer	Acc@1 (base)	Acc@1 (SIF+ABTT)	Existing (base)	Existing (SIF+ABTT)	New (base)	New (SIF+ABTT)
LaTa	1	0.731 $\pm$ 0.009	0.889 $\pm$ 0.009	0.562 $\pm$ 0.004	0.842 $\pm$ 0.009	0.949 $\pm$ 0.011	0.950 $\pm$ 0.011
	2	0.428 $\pm$ 0.011	0.887 $\pm$ 0.005	0.048 $\pm$ 0.006	0.840 $\pm$ 0.005	0.920 $\pm$ 0.008	0.949 $\pm$ 0.009
	3	0.392 $\pm$ 0.012	0.889 $\pm$ 0.008	0.046 $\pm$ 0.008	0.845 $\pm$ 0.006	0.840 $\pm$ 0.020	0.946 $\pm$ 0.013
	4	0.330 $\pm$ 0.018	0.888 $\pm$ 0.008	0.047 $\pm$ 0.008	0.858 $\pm$ 0.007	0.696 $\pm$ 0.030	0.927 $\pm$ 0.013
	5	0.326 $\pm$ 0.020	0.890 $\pm$ 0.008	0.039 $\pm$ 0.004	0.856 $\pm$ 0.005	0.698 $\pm$ 0.028	0.935 $\pm$ 0.014
	6	0.321 $\pm$ 0.018	0.890 $\pm$ 0.008	0.038 $\pm$ 0.007	0.861 $\pm$ 0.008	0.688 $\pm$ 0.026	0.927 $\pm$ 0.014
	7	0.327 $\pm$ 0.020	0.895 $\pm$ 0.006	0.044 $\pm$ 0.007	0.846 $\pm$ 0.006	0.694 $\pm$ 0.029	0.957 $\pm$ 0.010
	8	0.325 $\pm$ 0.022	0.883 $\pm$ 0.007	0.046 $\pm$ 0.010	0.832 $\pm$ 0.006	0.685 $\pm$ 0.032	0.949 $\pm$ 0.011
	9	0.331 $\pm$ 0.019	0.885 $\pm$ 0.008	0.045 $\pm$ 0.010	0.846 $\pm$ 0.006	0.701 $\pm$ 0.024	0.935 $\pm$ 0.013
	10	0.346 $\pm$ 0.017	0.888 $\pm$ 0.002	0.044 $\pm$ 0.010	0.844 $\pm$ 0.007	0.736 $\pm$ 0.022	0.944 $\pm$ 0.010
	11	0.360 $\pm$ 0.018	0.881 $\pm$ 0.006	0.045 $\pm$ 0.009	0.860 $\pm$ 0.008	0.767 $\pm$ 0.028	0.907 $\pm$ 0.011
	12	0.707 $\pm$ 0.015	0.889 $\pm$ 0.005	0.511 $\pm$ 0.017	0.834 $\pm$ 0.005	0.961 $\pm$ 0.008	0.961 $\pm$ 0.010
PhilTa	1	0.706 $\pm$ 0.016	0.898 $\pm$ 0.007	0.532 $\pm$ 0.018	0.856 $\pm$ 0.009	0.931 $\pm$ 0.012	0.953 $\pm$ 0.011
	2	0.625 $\pm$ 0.018	0.889 $\pm$ 0.009	0.408 $\pm$ 0.016	0.866 $\pm$ 0.009	0.905 $\pm$ 0.017	0.919 $\pm$ 0.017
	3	0.439 $\pm$ 0.013	0.892 $\pm$ 0.006	0.071 $\pm$ 0.014	0.856 $\pm$ 0.008	0.917 $\pm$ 0.013	0.939 $\pm$ 0.009
	4	0.420 $\pm$ 0.016	0.893 $\pm$ 0.009	0.027 $\pm$ 0.007	0.858 $\pm$ 0.006	0.928 $\pm$ 0.015	0.938 $\pm$ 0.015
	5	0.374 $\pm$ 0.015	0.896 $\pm$ 0.010	0.042 $\pm$ 0.011	0.872 $\pm$ 0.007	0.805 $\pm$ 0.018	0.925 $\pm$ 0.014
	6	0.338 $\pm$ 0.015	0.897 $\pm$ 0.008	0.059 $\pm$ 0.010	0.861 $\pm$ 0.007	0.698 $\pm$ 0.025	0.945 $\pm$ 0.014
	7	0.361 $\pm$ 0.016	0.900 $\pm$ 0.009	0.056 $\pm$ 0.010	0.859 $\pm$ 0.005	0.757 $\pm$ 0.021	0.952 $\pm$ 0.016

(continued on next page)

Task B: Top-K Ranking (appendix, 5-seed mean $\pm$ std, SIF+ABTT)							
Model	Layer	Acc@1 (base)	Acc@1 (SIF+ABTT)	Existing (base)	Existing (SIF+ABTT)	New (base)	New (SIF+ABTT)
mT5-base	8	0.390 $\pm$ 0.016	0.891 $\pm$ 0.010	0.058 $\pm$ 0.012	0.858 $\pm$ 0.007	0.820 $\pm$ 0.017	0.933 $\pm$ 0.018
	9	0.372 $\pm$ 0.015	0.894 $\pm$ 0.006	0.066 $\pm$ 0.012	0.861 $\pm$ 0.006	0.768 $\pm$ 0.026	0.936 $\pm$ 0.011
	10	0.377 $\pm$ 0.013	0.892 $\pm$ 0.006	0.066 $\pm$ 0.012	0.855 $\pm$ 0.003	0.781 $\pm$ 0.021	0.941 $\pm$ 0.015
	11	0.394 $\pm$ 0.013	0.889 $\pm$ 0.006	0.058 $\pm$ 0.008	0.851 $\pm$ 0.008	0.829 $\pm$ 0.018	0.937 $\pm$ 0.008
	12	0.563 $\pm$ 0.010	0.896 $\pm$ 0.007	0.279 $\pm$ 0.011	0.863 $\pm$ 0.009	0.932 $\pm$ 0.005	0.938 $\pm$ 0.014
	1	0.486 $\pm$ 0.012	0.898 $\pm$ 0.004	0.100 $\pm$ 0.007	0.856 $\pm$ 0.009	0.987 $\pm$ 0.007	0.953 $\pm$ 0.010
	2	0.475 $\pm$ 0.012	0.892 $\pm$ 0.004	0.082 $\pm$ 0.008	0.851 $\pm$ 0.013	0.984 $\pm$ 0.008	0.946 $\pm$ 0.013
	3	0.482 $\pm$ 0.014	0.872 $\pm$ 0.008	0.123 $\pm$ 0.008	0.809 $\pm$ 0.012	0.947 $\pm$ 0.008	0.954 $\pm$ 0.008
	4	0.498 $\pm$ 0.013	0.852 $\pm$ 0.006	0.205 $\pm$ 0.006	0.772 $\pm$ 0.012	0.879 $\pm$ 0.018	0.956 $\pm$ 0.010
	5	0.370 $\pm$ 0.013	0.825 $\pm$ 0.009	0.155 $\pm$ 0.005	0.724 $\pm$ 0.010	0.649 $\pm$ 0.020	0.956 $\pm$ 0.005
	6	0.359 $\pm$ 0.011	0.809 $\pm$ 0.006	0.263 $\pm$ 0.021	0.765 $\pm$ 0.013	0.483 $\pm$ 0.018	0.865 $\pm$ 0.018
	7	0.368 $\pm$ 0.007	0.827 $\pm$ 0.007	0.256 $\pm$ 0.012	0.758 $\pm$ 0.011	0.512 $\pm$ 0.018	0.916 $\pm$ 0.012
LaBSE	8	0.380 $\pm$ 0.007	0.824 $\pm$ 0.012	0.245 $\pm$ 0.007	0.771 $\pm$ 0.013	0.555 $\pm$ 0.017	0.892 $\pm$ 0.015
	9	0.426 $\pm$ 0.012	0.836 $\pm$ 0.010	0.171 $\pm$ 0.006	0.781 $\pm$ 0.014	0.756 $\pm$ 0.023	0.907 $\pm$ 0.006
	10	0.420 $\pm$ 0.007	0.821 $\pm$ 0.016	0.141 $\pm$ 0.011	0.765 $\pm$ 0.016	0.782 $\pm$ 0.019	0.893 $\pm$ 0.017
	11	0.418 $\pm$ 0.012	0.803 $\pm$ 0.010	0.132 $\pm$ 0.010	0.719 $\pm$ 0.013	0.788 $\pm$ 0.014	0.910 $\pm$ 0.010
	12	0.506 $\pm$ 0.008	0.840 $\pm$ 0.009	0.170 $\pm$ 0.008	0.770 $\pm$ 0.017	0.943 $\pm$ 0.012	0.930 $\pm$ 0.008
	1	0.501 $\pm$ 0.011	0.858 $\pm$ 0.008	0.170 $\pm$ 0.002	0.838 $\pm$ 0.007	0.930 $\pm$ 0.014	0.884 $\pm$ 0.010
	2	0.522 $\pm$ 0.010	0.861 $\pm$ 0.006	0.272 $\pm$ 0.003	0.817 $\pm$ 0.004	0.844 $\pm$ 0.022	0.917 $\pm$ 0.013
	3	0.535 $\pm$ 0.008	0.845 $\pm$ 0.005	0.217 $\pm$ 0.007	0.831 $\pm$ 0.010	0.946 $\pm$ 0.007	0.864 $\pm$ 0.009
	4	0.568 $\pm$ 0.008	0.849 $\pm$ 0.007	0.298 $\pm$ 0.006	0.837 $\pm$ 0.008	0.918 $\pm$ 0.014	0.865 $\pm$ 0.015
	5	0.550 $\pm$ 0.011	0.845 $\pm$ 0.007	0.254 $\pm$ 0.011	0.794 $\pm$ 0.006	0.933 $\pm$ 0.013	0.911 $\pm$ 0.015
	6	0.660 $\pm$ 0.012	0.858 $\pm$ 0.003	0.460 $\pm$ 0.021	0.816 $\pm$ 0.009	0.918 $\pm$ 0.018	0.912 $\pm$ 0.016
	7	0.595 $\pm$ 0.010	0.850 $\pm$ 0.003	0.333 $\pm$ 0.019	0.808 $\pm$ 0.006	0.934 $\pm$ 0.015	0.904 $\pm$ 0.010
Qwen3-0.6B	8	0.571 $\pm$ 0.013	0.848 $\pm$ 0.004	0.313 $\pm$ 0.019	0.802 $\pm$ 0.011	0.905 $\pm$ 0.018	0.906 $\pm$ 0.007
	9	0.630 $\pm$ 0.008	0.855 $\pm$ 0.004	0.380 $\pm$ 0.018	0.828 $\pm$ 0.005	0.954 $\pm$ 0.008	0.890 $\pm$ 0.007
	10	0.711 $\pm$ 0.015	0.856 $\pm$ 0.003	0.550 $\pm$ 0.018	0.835 $\pm$ 0.005	0.920 $\pm$ 0.016	0.884 $\pm$ 0.004
	11	0.820 $\pm$ 0.009	0.904 $\pm$ 0.008	0.732 $\pm$ 0.011	0.873 $\pm$ 0.010	0.935 $\pm$ 0.013	0.944 $\pm$ 0.017
	12	0.792 $\pm$ 0.013	0.883 $\pm$ 0.006	0.668 $\pm$ 0.016	0.849 $\pm$ 0.009	0.952 $\pm$ 0.013	0.928 $\pm$ 0.014
	1	0.549 $\pm$ 0.006	0.892 $\pm$ 0.004	0.221 $\pm$ 0.009	0.865 $\pm$ 0.008	0.972 $\pm$ 0.003	0.927 $\pm$ 0.007
	2	0.570 $\pm$ 0.013	0.882 $\pm$ 0.003	0.362 $\pm$ 0.012	0.828 $\pm$ 0.007	0.839 $\pm$ 0.028	0.951 $\pm$ 0.006
	3	0.540 $\pm$ 0.006	0.889 $\pm$ 0.005	0.214 $\pm$ 0.010	0.827 $\pm$ 0.011	0.962 $\pm$ 0.007	0.969 $\pm$ 0.011
	4	0.581 $\pm$ 0.010	0.891 $\pm$ 0.004	0.354 $\pm$ 0.008	0.844 $\pm$ 0.009	0.874 $\pm$ 0.023	0.952 $\pm$ 0.010
	5	0.486 $\pm$ 0.016	0.894 $\pm$ 0.007	0.133 $\pm$ 0.013	0.839 $\pm$ 0.011	0.943 $\pm$ 0.012	0.965 $\pm$ 0.010
	6	0.566 $\pm$ 0.011	0.897 $\pm$ 0.005	0.349 $\pm$ 0.009	0.857 $\pm$ 0.010	0.848 $\pm$ 0.020	0.949 $\pm$ 0.013
	7	0.540 $\pm$ 0.007	0.902 $\pm$ 0.006	0.211 $\pm$ 0.011	0.862 $\pm$ 0.009	0.967 $\pm$ 0.012	0.955 $\pm$ 0.015
KaLM-mini	8	0.608 $\pm$ 0.010	0.903 $\pm$ 0.006	0.361 $\pm$ 0.016	0.861 $\pm$ 0.011	0.927 $\pm$ 0.015	0.958 $\pm$ 0.009
	9	0.610 $\pm$ 0.011	0.901 $\pm$ 0.005	0.391 $\pm$ 0.015	0.867 $\pm$ 0.009	0.892 $\pm$ 0.019	0.944 $\pm$ 0.014
	10	0.636 $\pm$ 0.012	0.890 $\pm$ 0.008	0.450 $\pm$ 0.009	0.847 $\pm$ 0.014	0.877 $\pm$ 0.020	0.946 $\pm$ 0.010
	11	0.577 $\pm$ 0.006	0.886 $\pm$ 0.005	0.257 $\pm$ 0.013	0.838 $\pm$ 0.016	0.991 $\pm$ 0.005	0.948 $\pm$ 0.015
	12	0.626 $\pm$ 0.009	0.893 $\pm$ 0.004	0.365 $\pm$ 0.019	0.844 $\pm$ 0.013	0.964 $\pm$ 0.013	0.958 $\pm$ 0.010
	13	0.638 $\pm$ 0.010	0.888 $\pm$ 0.006	0.382 $\pm$ 0.016	0.849 $\pm$ 0.013	0.970 $\pm$ 0.015	0.938 $\pm$ 0.009
	14	0.644 $\pm$ 0.010	0.885 $\pm$ 0.007	0.401 $\pm$ 0.019	0.852 $\pm$ 0.012	0.958 $\pm$ 0.015	0.929 $\pm$ 0.008
	15	0.646 $\pm$ 0.013	0.885 $\pm$ 0.005	0.409 $\pm$ 0.020	0.852 $\pm$ 0.012	0.952 $\pm$ 0.016	0.929 $\pm$ 0.017
	16	0.689 $\pm$ 0.015	0.873 $\pm$ 0.006	0.544 $\pm$ 0.017	0.837 $\pm$ 0.010	0.877 $\pm$ 0.016	0.920 $\pm$ 0.020
	17	0.657 $\pm$ 0.010	0.861 $\pm$ 0.008	0.419 $\pm$ 0.020	0.807 $\pm$ 0.014	0.964 $\pm$ 0.012	0.932 $\pm$ 0.017
	18	0.708 $\pm$ 0.011	0.868 $\pm$ 0.004	0.503 $\pm$ 0.019	0.833 $\pm$ 0.016	0.973 $\pm$ 0.005	0.913 $\pm$ 0.020
	19	0.690 $\pm$ 0.011	0.873 $\pm$ 0.005	0.462 $\pm$ 0.018	0.859 $\pm$ 0.008	0.986 $\pm$ 0.006	0.892 $\pm$ 0.015
	20	0.742 $\pm$ 0.019	0.880 $\pm$ 0.006	0.588 $\pm$ 0.023	0.858 $\pm$ 0.009	0.941 $\pm$ 0.017	0.909 $\pm$ 0.016
	21	0.757 $\pm$ 0.016	0.878 $\pm$ 0.006	0.584 $\pm$ 0.021	0.835 $\pm$ 0.009	0.981 $\pm$ 0.010	0.934 $\pm$ 0.013
	22	0.798 $\pm$ 0.014	0.882 $\pm$ 0.006	0.670 $\pm$ 0.021	0.854 $\pm$ 0.004	0.965 $\pm$ 0.006	0.917 $\pm$ 0.017
	23	0.757 $\pm$ 0.013	0.888 $\pm$ 0.007	0.578 $\pm$ 0.016	0.857 $\pm$ 0.005	0.989 $\pm$ 0.004	0.929 $\pm$ 0.016
	24	0.810 $\pm$ 0.014	0.894 $\pm$ 0.005	0.694 $\pm$ 0.022	0.861 $\pm$ 0.005	0.960 $\pm$ 0.003	0.936 $\pm$ 0.015
	25	0.732 $\pm$ 0.012	0.890 $\pm$ 0.009	0.528 $\pm$ 0.016	0.866 $\pm$ 0.005	0.997 $\pm$ 0.003	0.922 $\pm$ 0.017
	26	0.749 $\pm$ 0.013	0.895 $\pm$ 0.007	0.561 $\pm$ 0.017	0.858 $\pm$ 0.002	0.991 $\pm$ 0.003	0.943 $\pm$ 0.017
	27	0.811 $\pm$ 0.004	0.881 $\pm$ 0.003	0.741 $\pm$ 0.008	0.824 $\pm$ 0.007	0.901 $\pm$ 0.011	0.954 $\pm$ 0.011
	28	0.822 $\pm$ 0.010	0.870 $\pm$ 0.007	0.733 $\pm$ 0.014	0.800 $\pm$ 0.012	0.937 $\pm$ 0.006	0.961 $\pm$ 0.009
	1	0.595 $\pm$ 0.011	0.893 $\pm$ 0.004	0.344 $\pm$ 0.002	0.846 $\pm$ 0.007	0.920 $\pm$ 0.021	0.954 $\pm$ 0.009
	2	0.620 $\pm$ 0.017	0.894 $\pm$ 0.005	0.422 $\pm$ 0.014	0.845 $\pm$ 0.005	0.875 $\pm$ 0.021	0.958 $\pm$ 0.013
	3	0.604 $\pm$ 0.010	0.894 $\pm$ 0.003	0.378 $\pm$ 0.009	0.843 $\pm$ 0.003	0.898 $\pm$ 0.017	0.960 $\pm$ 0.011
	4	0.562 $\pm$ 0.015	0.889 $\pm$ 0.004	0.265 $\pm$ 0.009	0.842 $\pm$ 0.010	0.946 $\pm$ 0.016	0.951 $\pm$ 0.011
	5	0.550 $\pm$ 0.012	0.890 $\pm$ 0.005	0.270 $\pm$ 0.007	0.841 $\pm$ 0.006	0.912 $\pm$ 0.018	0.954 $\pm$ 0.012

(continued on next page)

Task B: Top-K Ranking (appendix, 5-seed mean $\pm$ std, SIF+ABTT)							
Model	Layer	Acc@1 (base)	Acc@1 (SIF+ABTT)	Existing (base)	Existing (SIF+ABTT)	New (base)	New (SIF+ABTT)
	6	0.522 $\pm$ 0.015	0.894 $\pm$ 0.004	0.202 $\pm$ 0.011	0.850 $\pm$ 0.006	0.935 $\pm$ 0.013	0.950 $\pm$ 0.015
	7	0.604 $\pm$ 0.017	0.894 $\pm$ 0.003	0.394 $\pm$ 0.016	0.840 $\pm$ 0.011	0.876 $\pm$ 0.020	0.963 $\pm$ 0.012
	8	0.607 $\pm$ 0.013	0.866 $\pm$ 0.003	0.387 $\pm$ 0.010	0.798 $\pm$ 0.010	0.892 $\pm$ 0.016	0.955 $\pm$ 0.010
	9	0.614 $\pm$ 0.010	0.874 $\pm$ 0.005	0.386 $\pm$ 0.009	0.821 $\pm$ 0.011	0.909 $\pm$ 0.013	0.944 $\pm$ 0.011
	10	0.602 $\pm$ 0.012	0.860 $\pm$ 0.007	0.359 $\pm$ 0.005	0.774 $\pm$ 0.012	0.917 $\pm$ 0.017	0.970 $\pm$ 0.012
	11	0.601 $\pm$ 0.012	0.864 $\pm$ 0.008	0.360 $\pm$ 0.009	0.792 $\pm$ 0.012	0.911 $\pm$ 0.020	0.957 $\pm$ 0.008
	12	0.638 $\pm$ 0.009	0.868 $\pm$ 0.006	0.415 $\pm$ 0.008	0.801 $\pm$ 0.009	0.927 $\pm$ 0.016	0.955 $\pm$ 0.005
	13	0.643 $\pm$ 0.010	0.862 $\pm$ 0.006	0.435 $\pm$ 0.007	0.792 $\pm$ 0.005	0.912 $\pm$ 0.016	0.953 $\pm$ 0.016
	14	0.653 $\pm$ 0.009	0.858 $\pm$ 0.007	0.479 $\pm$ 0.002	0.793 $\pm$ 0.009	0.879 $\pm$ 0.016	0.943 $\pm$ 0.009
	15	0.659 $\pm$ 0.007	0.868 $\pm$ 0.005	0.458 $\pm$ 0.006	0.812 $\pm$ 0.009	0.919 $\pm$ 0.010	0.941 $\pm$ 0.006
	16	0.682 $\pm$ 0.005	0.869 $\pm$ 0.008	0.475 $\pm$ 0.006	0.824 $\pm$ 0.008	0.951 $\pm$ 0.010	0.926 $\pm$ 0.009
	17	0.756 $\pm$ 0.007	0.885 $\pm$ 0.009	0.611 $\pm$ 0.018	0.841 $\pm$ 0.006	0.944 $\pm$ 0.010	0.942 $\pm$ 0.015
	18	0.780 $\pm$ 0.010	0.881 $\pm$ 0.007	0.635 $\pm$ 0.016	0.840 $\pm$ 0.002	0.967 $\pm$ 0.007	0.933 $\pm$ 0.016
	19	0.804 $\pm$ 0.013	0.876 $\pm$ 0.007	0.672 $\pm$ 0.022	0.831 $\pm$ 0.007	0.977 $\pm$ 0.003	0.935 $\pm$ 0.008
	20	0.793 $\pm$ 0.011	0.877 $\pm$ 0.004	0.649 $\pm$ 0.020	0.831 $\pm$ 0.003	0.979 $\pm$ 0.004	0.936 $\pm$ 0.008
	21	0.804 $\pm$ 0.011	0.871 $\pm$ 0.007	0.672 $\pm$ 0.018	0.830 $\pm$ 0.005	0.975 $\pm$ 0.006	0.924 $\pm$ 0.012
	22	0.870 $\pm$ 0.008	0.873 $\pm$ 0.004	0.810 $\pm$ 0.007	0.812 $\pm$ 0.005	0.948 $\pm$ 0.008	0.952 $\pm$ 0.009
	23	0.873 $\pm$ 0.007	0.884 $\pm$ 0.004	0.811 $\pm$ 0.008	0.833 $\pm$ 0.004	0.953 $\pm$ 0.007	0.951 $\pm$ 0.007
	24	0.848 $\pm$ 0.011	0.856 $\pm$ 0.005	0.749 $\pm$ 0.019	0.806 $\pm$ 0.009	0.975 $\pm$ 0.007	0.921 $\pm$ 0.010




Figure 4: ρ_{LOO} foreground view for the main candidate-attribution result. Each line connects the baseline and ABTT variants for one model-method cell at the predeclared operational attribution layer. ABTT improves the leave-one-out rank-correlation signal in all six cells; Table 5 shows the ERASER-style disagreements which qualify this narrower faithfulness claim.

F Attribution Metric Sweeps

Figure 4 plots the rank-faithfulness column of Table 5. The rest of this appendix expands the main attribution table across Sufficiency and Comprehensiveness thresholds of 10%, 25%, and 50%, and MinFrac recovery thresholds of 0.70, 0.80, 0.90, and 0.95. The main text uses one global midpoint setting rather than selecting thresholds per model or per method. Table 17 gives the sweep for integrated gradients and retrieval-adapted MaRC, and Table 18 gives it for the supplemental attribution methods.

G Qualitative Attribution Examples

This appendix keeps the example-level views for one PhilTa retrieval pair. Figure 5 gives the integrated-gradient pair matrix before and after ABTT. The attention map in Figure 6 is context only, since ABTT is applied after encoding. The retrieval-adapted MaRC masks in Figure 7 show the learned-mask view for the same pair.

H Cluster Geometry

Figures 8 and 9 give the t-SNE and UMAP panels for LaTa, PhilTa, and mT5-base, and Figures 10 and 11 extend them to LaBSE, Qwen3-0.6B, and KaLM-mini. All four use the same six directories and the same baseline versus ABTT-only comparison.

Model	Method	$\rho_{\text{LOO}} \uparrow$		Suff@10% $\uparrow$		Suff@25% $\uparrow$		Suff@50% $\uparrow$		Comp@10% $\uparrow$		Comp@25% $\uparrow$		Comp@50% $\uparrow$		MinFrac@0.70 $\downarrow$		MinFrac@0.80 $\downarrow$		MinFrac@0.90 $\downarrow$		MinFrac@0.95 $\downarrow$	
		base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt
LaTa	IG	-0.001	0.396	0.967	0.428	0.976	0.681	0.975	0.921	0.320	0.267	0.719	0.581	0.850	0.884	0.040	0.293	0.041	0.379	0.045	0.489	0.047	0.557
	MaRC	0.042	0.401	0.404	0.475	0.534	0.645	0.582	0.863	0.116	0.250	0.361	0.539	0.506	0.845	0.368	0.356	0.383	0.459	0.396	0.568	0.438	0.624
PhilTa	IG	0.182	0.607	0.816	0.848	0.959	0.963	0.988	1.057	0.704	0.436	0.969	0.747	1.010	0.992	0.027	0.090	0.079	0.141	0.162	0.229	0.215	0.300
	MaRC	0.190	0.331	0.786	0.531	0.911	0.754	0.949	0.935	0.657	0.237	0.963	0.521	0.978	0.793	0.058	0.281	0.104	0.358	0.221	0.448	0.327	0.503
mT5-base	IG	0.164	0.597	0.810	0.785	0.892	1.027	0.937	1.192	0.054	0.376	0.110	0.717	0.284	1.071	0.055	0.161	0.116	0.214	0.338	0.291	0.623	0.336
	MaRC	0.250	0.392	0.865	0.597	0.913	0.827	0.951	1.055	0.010	0.279	0.042	0.592	0.118	0.923	0.026	0.277	0.066	0.346	0.212	0.434	0.416	0.488

Table 17: Appendix attribution sweep for the two paper-facing attribution views, IG and retrieval-adapted MARC, across the three main-paper models. Direction arrows are printed in the column headers. Sufficiency and Comprehensiveness are ratio metrics at $\{10\%, 25\%, 50\%\}$; MinFrac is the minimum retained query-token fraction needed to recover $\tau \in \{0.70, 0.80, 0.90, 0.95\}$ of the full cosine (lower is better). Ratio metrics require $|S_v(q_{\text{full}}, c)| \geq 0.05$; ρ_{LOO} excludes $|\Delta| < 10^{-6}$. The completeness report contains no missing metric cells.

Model	Method	$\rho_{\text{LOO}} \uparrow$		Suff@10% $\uparrow$		Suff@25% $\uparrow$		Suff@50% $\uparrow$		Comp@10% $\uparrow$		Comp@25% $\uparrow$		Comp@50% $\uparrow$		MinFrac@0.70 $\downarrow$		MinFrac@0.80 $\downarrow$		MinFrac@0.90 $\downarrow$		MinFrac@0.95 $\downarrow$	
		base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt
LaTa	BERTScore	0.204	0.430	1.035	0.365	1.047	0.626	1.047	0.904	0.916	0.221	1.301	0.436	1.280	0.719	0.018	0.318	0.021	0.394	0.026	0.481	0.036	0.534
	OT	-0.001	0.396	0.967	0.428	0.976	0.681	0.975	0.921	0.320	0.267	0.719	0.581	0.850	0.884	0.040	0.293	0.041	0.379	0.045	0.489	0.047	0.557
	Att-weighted	-0.137	0.106	-0.575	0.142	-0.377	0.295	-0.131	0.588	0.001	0.159	-0.005	0.330	-0.019	0.543	0.746	0.569	0.756	0.653	0.774	0.747	0.796	0.894
	Att-standalone	0.005	0.368	0.115	0.358	0.239	0.640	0.409	0.873	0.043	0.251	0.112	0.536	0.342	0.830	0.516	0.333	0.531	0.422	0.558	0.537	0.584	0.507
	DLA	0.054	0.278	1.017	0.350	1.004	0.607	1.003	0.820	0.937	0.239	1.289	0.505	1.248	0.792	0.018	0.368	0.021	0.471	0.026	0.583	0.036	0.650
	random	0.003	-0.003	0.403	0.222	0.726	0.360	0.906	0.599	0.024	0.072	0.046	0.181	0.105	0.393	0.176	0.561	0.182	0.667	0.192	0.779	0.212	0.842
	inverse	0.001	-0.396	-0.046	-0.014	0.122	-0.001	0.151	0.124	0.000	-0.046	0.001	-0.035	0.025	0.086	0.553	0.837	0.562	0.882	0.602	0.920	0.629	0.942
PhilTa	BERTScore	0.309	0.550	0.861	0.771	0.985	0.942	1.035	1.050	0.687	0.313	1.032	0.580	1.042	0.808	0.026	0.128	0.070	0.188	0.190	0.284	0.280	0.343
	OT	0.182	0.607	0.816	0.848	0.959	0.963	0.988	1.057	0.704	0.436	0.969	0.747	1.010	0.992	0.027	0.090	0.079	0.141	0.162	0.229	0.215	0.300
	Att-weighted	0.053	0.108	0.060	0.341	0.163	0.584	0.219	0.753	0.011	0.157	0.028	0.326	0.059	0.524	0.673	0.397	0.693	0.505	0.742	0.627	0.795	0.698
	Att-standalone	0.046	0.205	0.017	0.351	0.038	0.654	0.064	0.827	0.007	0.163	0.040	0.378	0.069	0.616	0.773	0.323	0.782	0.417	0.796	0.537	0.808	0.616
	DLA	0.203	0.487	0.805	0.863	0.942	0.910	0.991	0.957	0.715	0.438	1.017	0.729	1.028	0.927	0.028	0.087	0.080	0.142	0.194	0.260	0.260	0.342
	random	-0.001	-0.010	0.589	0.288	0.826	0.478	0.953	0.709	0.011	0.050	0.016	0.124	0.066	0.288	0.135	0.437	0.145	0.563	0.204	0.713	0.284	0.798
	inverse	-0.182	-0.607	-0.009	-0.074	-0.025	-0.074	-0.009	0.014	-0.001	-0.049	-0.001	-0.073	0.013	-0.055	0.917	0.920	0.921	0.950	0.932	0.975	0.942	0.987
mT5-base	BERTScore	0.271	0.556	0.615	0.694	0.776	1.005	0.917	1.197	0.014	0.270	0.038	0.579	0.082	0.906	0.148	0.199	0.214	0.255	0.333	0.322	0.452	0.358
	OT	0.164	0.597	0.810	0.785	0.892	1.027	0.937	1.192	0.054	0.376	0.110	0.717	0.284	1.071	0.055	0.161	0.116	0.214	0.338	0.291	0.623	0.336
	Att-weighted	0.173	0.523	0.872	0.890	0.923	1.008	0.942	1.099	0.012	0.364	0.048	0.687	0.249	0.986	0.036	0.117	0.074	0.181	0.239	0.278	0.526	0.346
	Att-standalone	0.110	0.064	0.640	0.180	0.650	0.455	0.862	0.643	0.034	0.184	0.068	0.362	0.135	0.464	0.066	0.471	0.137	0.575	0.412	0.718	0.608	0.786
	DLA	0.129	0.481	0.296	0.922	0.786	0.996	0.949	1.010	0.029	0.408	0.045	0.683	0.083	0.908	0.211	0.097	0.259	0.159	0.347	0.257	0.460	0.337
	random	-0.000	-0.004	0.846	0.296	0.926	0.483	0.969	0.687	0.004	0.059	0.013	0.153	0.033	0.317	0.045	0.460	0.079	0.591	0.178	0.735	0.313	0.813
	inverse	-0.164	-0.597	-0.127	-0.235	0.203	-0.237	0.722	-0.065	0.022	-0.108	0.045	-0.187	0.064	-0.189	0.472	0.896	0.563	0.936	0.740	0.971	0.863	0.985

Table 18: Supplemental attribution sweep for additional methods and sanity baselines. *random* and *inverse* are diagnostic baselines, not paper-facing attribution methods. Direction arrows are printed in the column headers. Sufficiency and Comprehensiveness are ratio metrics at $\{10\%, 25\%, 50\%\}$; MinFrac is the minimum retained query-token fraction needed to recover $\tau \in \{0.70, 0.80, 0.90, 0.95\}$ of the full cosine (lower is better). Ratio metrics require $|S_v(q_{\text{full}}, c)| \geq 0.05$; ρ_{LOO} excludes $|\Delta| < 10^{-6}$. The completeness report contains no missing metric cells.




Figure 5: Derived pairwise heatmap for one PhilTa retrieval case. The left panel shows the baseline representation and the right panel shows the same example after ABTT. Both axes use original decoded tokens. Cell values combine token-token cosine similarity with single-sequence integrated-gradient scores.




Figure 6: Companion self-attention maps for the same PhilTa example as Figure 5. ABTT is applied after encoding, so attention is included only as supporting context.




Figure 7: Retrieval-adapted MaRC masks for the same PhilTa pair as Figure 5. Top row: mask on the query with the candidate fixed. Bottom row: mask on the candidate with the query fixed. Left column: baseline cosine target. Right column: ABTT-corrected cosine target.

2D TSNE of labelled fragments at each model's best layer




Figure 8: t-SNE projections of the 1,705 labelled fragments at each model's selected layer for LaTa, PhilTa, and mT5-base. Columns compare baseline mean pooling with ABTT-only using $D = 10$. Colors mark the six most-populated directories.

2D UMAP of labelled fragments at each model's best layer




Figure 9: UMAP projections with the same configuration as Figure 8. The same cluster-tightening pattern appears under a different projection method.

2D TSNE of labelled fragments at each model's best layer




Figure 10: t-SNE projections of the six most-populated directories at each selected layer for LaBSE, Qwen3-0.6B, and KaLM-mini. Columns compare baseline mean pooling with ABTT-only using $D = 10$.

2D UMAP of labelled fragments at each model's best layer




Figure 11: UMAP projections matching Figure 10. The cluster-tightening pattern appears under a second projection method.