

Semantic Matching in Medieval Latin: A Benchmark, Layerwise Analysis, and Geometric Repair

Anonymous ACL submission

Abstract

Medieval canon law survives through manuscript families shaped by spelling variation, fragmentary copying, and editorial change. Scholars need strict semantic matching to consolidate duplicate passages and link related witnesses. We introduce a Medieval Latin benchmark with 1,705 canon-law fragments grouped into 840 meaning-equivalent families, and we evaluate pairwise duplicate detection plus open-set directory routing. A layerwise analysis shows a geometric failure mode in Latin T5 representations: several intermediate layers collapse into low-rank cones where cosine scores lose class separation. We use All-but-the-Top post-processing, fit only on training embeddings, as a train-free repair after pooling. ABTT reduces top-principal-component dominance, raises effective rank, and improves retrieval separability in the headline layer grid. We then evaluate local attributions for the same retrieval score with integrated gradients and retrieval-adapted MaRC. The strongest attribution result is narrow: ABTT improves leave-one-out rank faithfulness across all LaTa, PhilTa, and mT5-base model-method cells, while ERASER-style rationale metrics show mixed effects.

1 Introduction

Medieval canon law preserves early records of legal reasoning in Western Europe. The corpus does not survive as one stable text. Manuscript copies differ in spelling, wording, fragment boundaries, and editorial history. Scholars often need to decide whether two passages express the same legal content despite those surface changes. This decision supports duplicate consolidation, witness linking, and study of textual transmission.

We study this problem as strict semantic matching. The benchmark contains 1,705 Medieval Latin canon-law fragments arranged into 840 families




Figure 1: Cosine similarity distributions for PhilTa at the last layer and at the diagnostic mid layer, before and after ABTT. At the collapsed layer, ABTT separates equivalent and non-equivalent passages more clearly.

of meaning-equivalent variants. We evaluate two tasks. Task A asks whether two passages express the same meaning. Task B turns the same embeddings into an open-set workflow: a system either routes a new passage to an existing directory or marks the passage as new.

A simple retrieval system embeds each passage and ranks candidates by cosine similarity. This setup leaves a key modeling issue open. A transformer has many internal layers, and a final-layer sentence embedding does not reveal whether useful retrieval signal appears, disappears, or collapses inside the network. We inspect layers because the failure mode is geometric. Several middle layers form low-rank spaces where cosine scores concentrate and no longer separate equivalent from non-equivalent passages.

This diagnosis motivates our use of All-but-the-Top (ABTT) post-processing (Mu and Viswanath, 2018). ABTT is not a new method. Here, ABTT serves as a train-free repair after pooling: fit the dominant principal components on training embeddings, remove those directions from train and test embeddings, then score passages with cosine similarity. On the headline models, ABTT repairs much

of the low-rank collapse and makes intermediate layers usable for retrieval.

The layer analysis also gives readers a more careful selection story. Intrinsic geometry diagnoses collapse and recovery, but geometry does not select the best retrieval layer by itself. For operational retrieval and attribution, we use a train-only rule: choose the earliest layer within 0.5 percentage points of the best training-set directory accuracy at rank 1 under ABTT. This locks LaTa layer 7, PhilTa layer 1, and mT5-base layer 1 before attribution metrics are computed.

We also evaluate local explanations for the retrieval score. Integrated gradients (Sundararajan et al., 2017) gives a standard gradient view of token importance. Retrieval-adapted MaRC (Brinner and Zarrieß, 2023) gives a learned-mask view after replacing MaRC’s classification target with pairwise retrieval cosine. These two views sit at the same level in the paper: both explain a fixed query-candidate score, and both remain local diagnostics rather than global accounts of model behavior.

This paper makes four contributions: (1) we introduce a benchmark for strict semantic matching in Medieval Latin canon law, with pairwise and open-set retrieval tasks; (2) we show through a layerwise geometric analysis where intermediate Latin T5 representations suffer from top-PC dominance, low effective rank, and cosine concentration; (3) we show ABTT repairs much of this collapse and improves retrieval separability without updating model parameters; and (4) we adapt MaRC-style mask optimization to bi-encoder retrieval and compare the method with integrated gradients at predeclared retrieval-selected layers, finding leave-one-out rank faithfulness gains and mixed ERASER-style rationale metrics.

2 Task and Dataset

2.1 Task Definition

We study two related tasks.

Task A: pairwise duplicate detection. Let p_1 and p_2 be two passages and let $\equiv$ denote semantic equivalence. The label is

$$y = \begin{cases} 1 & \text{if } p_1 \equiv p_2, \\ 0 & \text{otherwise.} \end{cases}$$

The model produces $\cos(e_{p_1}, e_{p_2})$ from pooled embeddings. Task A uses no directory structure and no threshold at scoring time. Evaluation uses the test $n \times n$ cosine matrix.

Task B: directory routing. Task B uses the same embeddings in a scholar-facing retrieval workflow. Given a query passage q and the existing labelled directories, the system routes q to a directory or marks q as new. We evaluate two routing variants. In autonomous routing, each query receives the maximum cosine score to any other test file. A threshold τ , fit on training same-directory and different-directory pairs, decides whether the query enters its nearest neighbour’s directory or the new class. In top- k ranking, each directory receives the best document match score, and a pseudo-option `__NEW__` receives score τ . The simulated expert succeeds when the labelled directory appears in the shortlist.

2.2 Dataset Construction

The corpus contains 1,705 Medieval Latin manuscript fragments from canon law. The labelled files belong to 840 directories, where each directory represents one legal text and its manuscript copies. Directory sizes range from one to ten files. The benchmark asks for strict equivalence, so variants inside a directory count as positive matches and files from different directories count as negatives.

2.3 Train/Test Split

We split the corpus into 847 training files and 858 test files with a leak-free protocol (Table 1).¹ The 545 singleton directories are distributed at the file level. Doubleton directories are assigned as whole folders, 54 to train and 54 to test, so both files from a train doubleton contribute a positive training pair. For directories with at least three files, the split varies the train allocation inside each size class while keeping each class near a balanced file split. This avoids a rule where every size-three directory contributes no positive training pairs. The final split contains 565 positive training pairs and 595 positive test pairs.

All fitting uses training data only. This includes ABTT principal components, routing thresholds, ABTT dimensionality choices, and the operational layer rule used for attribution.

3 Methodology

3.1 Models

We evaluate six pre-trained models from three architecture families:

¹Implemented in `src/canon_split_v2.py`; deterministic with seed 42.

Category	Dirs	Train	Test
Singletons (1 file)	545	272	273
Doubletons (2 files)	108	108	108
Multi-file (≥ 3)	187	467	477
Total	840	847	858

Table 1: Train and test split statistics under the v2 varied-allocation protocol. Doubleton directories are assigned as whole folders. The test set contains 535 existing files, with a same-directory partner in test, and 323 new files.

- **LaTa**: a T5-based seq2seq model pre-trained on Latin (Riemenschneider and Frank, 2023). We extract encoder hidden states from 12 transformer blocks.
- **PhilTa**: a philologically oriented Latin T5 variant (Riemenschneider and Frank, 2023), again using 12 encoder blocks.
- **LaBSE**: a language-agnostic BERT sentence encoder covering 109 languages (Feng et al., 2022), with 12 encoder layers.
- **mT5-base**: a multilingual T5 encoder-decoder (Xue et al., 2021), with 12 encoder layers.
- **Qwen3-0.6B**: a multilingual decoder-only embedding model with 28 layers (Qwen Team, 2025).
- **KaLM-mini**: a multilingual decoder-only embedding model with 24 layers (Lu et al., 2025).

3.2 Embedding Extraction and Layer Selection

We extract hidden-state representations from every transformer block, starting at layer 1 and skipping layer 0, the static embedding table. Token representations at each layer are pooled into one passage vector by plain mean pooling or SIF weighting (Arora et al., 2017). For decoder-only models, last-token pooling is another common option; Appendix B compares last-token pooling with mean pooling under ABTT.

The paper-facing pooling baseline removes only tokenizer artifacts with no lexical content, such as whitespace, padding, and tokenizer-prefix-only tokens. Punctuation and numerals remain in the pooled representations. Attribution figures display original decoded tokens, even when a token contributes little to the pooled representation.

Layer geometry and layer selection serve different roles. We use label-free geometry diagnostics, including top-PC dominance, effective rank, and cosine concentration, to identify collapsed layers and test whether ABTT repairs them. For operational retrieval and attribution, we then apply a pre-declared train-only retrieval rule. The rule chooses the earliest layer within 0.5 percentage points of the best training-set directory accuracy at rank 1 under ABTT. The rule uses no test attribution outcome. This rule selects LaTa layer 7, PhilTa layer 1, and mT5-base layer 1 for attribution, with $D = 10$ ABTT components. The strongest diagnostic collapse layers are LaTa layer 8, PhilTa layer 6, and mT5-base layer 5, and we keep those as mechanism checks rather than attribution layers.

3.3 Post-Processing: ABTT

Our main post-processing method is All-but-the-Top (ABTT) (Mu and Viswanath, 2018). For each model and layer, we center the pooled training embeddings, compute the top D principal components by SVD, and remove those components from both train and test embeddings. We search $D \in \{1, 2, 3, 5, 7, 10\}$ and choose D by training-set Task B directory accuracy at rank 1. ABTT itself is borrowed from prior work; our use case is a retrieval-specific repair for anisotropic Latin manuscript representations.

We also compare against two simpler alternatives. SIF weighting changes the pooling weights before passage vectors are formed. PCA whitening (Su et al., 2021) rescales principal directions rather than only removing the top components.

3.4 Local Attribution Methods

We evaluate local explanations for a fixed query-candidate retrieval score. Integrated gradients (Sundararajan et al., 2017) gives a gradient-based attribution of pairwise cosine similarity to input tokens. We compute per-sequence token attributions for each side of the pair. The derived heatmap in Figure 6 combines token-token cosine similarity with the marginal integrated-gradient scores from the two sequences.

Our learned-mask attribution is adapted from MaRC (Brinner and Zarrieß, 2023). MaRC was proposed for classification, where a soft input mask preserves a target class score while sparsity and smoothness regularizers encourage concise rationales. We keep the per-instance mask-optimization idea and replace the classification target with pair-

wise retrieval cosine. Given a query-candidate pair, the mask is optimized on one side while the partner representation stays fixed. Candidate-side attribution is obtained by swapping the two sides.

Let $\mathbf{E}(\mathbf{x}_q) \in \mathbb{R}^{T_q \times d_{\text{emb}}}$ denote the input-embedding lookup, $E_{>0}$ the encoder with the lookup removed, and $\mathbf{z}_c^*$ the candidate representation held fixed during optimization. We parameterize $\mathbf{m} = \sigma(\boldsymbol{\zeta})$ with free logits $\boldsymbol{\zeta} \in \mathbb{R}^{T_q}$ and write $\mathcal{I}$ for content-token positions. For the ABTT variant, define $s_{\text{ABTT}}(\mathbf{a}, \mathbf{b}) = \cos(\text{ABTT}(\mathbf{a}), \mathbf{b})$. The retrieval-adapted MaRC loss is

$$\begin{aligned} \mathcal{L}(\mathbf{m}) = & (s_{\text{ABTT}}(\tilde{\mathbf{z}}_q(\mathbf{m}), \mathbf{z}_c^*) - s_{\text{ABTT}}(\mathbf{z}_q, \mathbf{z}_c^*))^2 \\ & + \lambda \sum_{i \in \mathcal{I}} \frac{m_i}{|\mathcal{I}|} + \gamma \sum_{i, i+1 \in \mathcal{I}} \frac{|m_{i+1} - m_i|}{|\mathcal{I}| - 1}, \end{aligned} \quad (1)$$

where $\tilde{\mathbf{z}}_q(\mathbf{m}) = \text{pool}(E_{>0}(\text{diag}(\mathbf{m}) \mathbf{E}(\mathbf{x}_q)), \mathbf{a}_q)$. The baseline variant uses raw cosine in place of the ABTT-corrected cosine. We run 200 Adam steps per pair, initialize $\zeta_0 = \log 9$, and freeze logits outside $\mathcal{I}$ at $-\infty$.

3.5 Evaluation Metrics

Task A. Both Task A metrics come from the test $n \times n$ pairwise cosine matrix. AUROC treats same-directory test pairs as positives and different-directory pairs as negatives. Cosine gap is the mean same-directory cosine minus the mean different-directory cosine.

Task B. For autonomous routing, we report overall assignment accuracy, existing accuracy, and new accuracy. Overall accuracy requires the existing-vs-new decision to match the ground truth, plus a correct directory when the query enters an existing directory. For top- k ranking, cumulative top- K accuracy is the fraction of test queries whose labelled directory appears among the top K options.

Attribution. For attribution, we adapt ERASER-style sufficiency and comprehensiveness metrics (DeYoung et al., 2020) to retrieval cosine. Sufficiency keeps the top-ranked query tokens and measures how much of the full pair cosine remains. Comprehensiveness removes those tokens and measures the score drop. MinFrac reports the smallest retained query-token fraction needed to recover a fixed share of the full cosine. We also report ρ_{LOO} , the Spearman correlation between attribution magnitude and the leave-one-out drop in the same retrieval score.

4 Results

The results move from diagnosis to use. First, layerwise geometry exposes representation collapse. Next, ABTT repairs much of the collapse and improves retrieval separability. Then Task B supplies the train-only layer rule used for operational attribution. Finally, we compare integrated gradients and retrieval-adapted MaRC at those locked layers.

4.1 Task A: Pairwise Duplicate Detection

Table 2 reports per-layer Task A metrics for LaTa, PhilTa, and mT5-base. The baseline uses mean pooling. The ABTT setting uses ABTT with D tuned per layer on the train split, without SIF weighting. Both AUROC and cosine gap come from the test cosine matrix and use no threshold. The table shows the main pattern behind the paper: baseline mean-pooled representations lose separation in the middle of the Latin T5 encoders, while ABTT recovers much of the lost separation.

4.2 Layerwise Geometry Diagnosis

Figures 2 and 3 show the layerwise retrieval profile. LaTa and PhilTa suffer the largest middle-layer drops under baseline mean pooling. Their AUROC and cosine gap recover toward later layers, so a final-layer-only report would miss the failure mode. mT5-base shows a milder version of the same pattern. LaBSE, Qwen3-0.6B, and KaLM-mini vary less across depth.

The geometry diagnostics explain this behavior. At the diagnostic collapse layers, PC1 variance shares for LaTa, PhilTa, and mT5-base are 0.956, 0.862, and 1.000 before correction. After ABTT-D10, those shares fall to 0.036, 0.039, and 0.044. Effective rank rises from 1.34, 1.83, and 1.00 to 168.40, 155.02, and 151.74. This supports the repair interpretation: ABTT removes dominant global directions which were suppressing cosine separation. ABTT does not add supervision or re-train the encoders.

Figure 1 gives the same diagnosis at the score-distribution level. At PhilTa’s collapsed layer, positive and negative cosine scores overlap heavily before ABTT. After ABTT, the distributions separate more clearly.

4.3 Cluster Geometry

Two-dimensional projections give a visual check on the same geometric change. Figures 4 and 5 show t-SNE (van der Maaten and Hinton, 2008)

Task A: Pairwise Duplicate Detection (main)					
Model	Layer	AUROC (base)	AUROC (ABTT)	Cosine gap (base)	Cosine gap (ABTT)
LaTa	1	0.933	0.969	0.286	0.565
	2	0.565	0.969	0.020	0.553
	3	0.515	0.972	-0.046	0.555
	4	0.502	0.970	-0.065	0.564
	5	0.498	0.971	-0.068	0.572
	6	0.497	0.963	-0.067	0.569
	7	0.499	0.961	-0.070	0.564
	8	0.504	0.964	-0.073	0.558
	9	0.504	0.964	-0.071	0.557
	10	0.507	0.970	-0.070	0.551
	11	0.510	0.973	-0.065	0.550
	12	0.938	0.971	0.238	0.525
PhilTa	1	0.939	0.976	0.231	0.573
	2	0.911	0.971	0.161	0.551
	3	0.601	0.972	0.085	0.549
	4	0.551	0.976	0.064	0.551
	5	0.543	0.978	0.054	0.551
	6	0.541	0.980	0.048	0.551
	7	0.542	0.979	0.047	0.554
	8	0.543	0.983	0.047	0.555
	9	0.539	0.982	0.042	0.540
	10	0.539	0.982	0.040	0.530
	11	0.540	0.982	0.041	0.511
	12	0.911	0.981	0.181	0.487
mT5-base	1	0.823	0.979	0.027	0.550
	2	0.823	0.973	0.026	0.535
	3	0.805	0.968	0.025	0.517
	4	0.800	0.975	0.040	0.503
	5	0.654	0.977	0.218	0.482
	6	0.657	0.973	0.238	0.482
	7	0.657	0.979	0.243	0.482
	8	0.656	0.980	0.248	0.471
	9	0.660	0.981	0.235	0.475
	10	0.659	0.975	0.225	0.459
	11	0.660	0.971	0.213	0.449
	12	0.839	0.970	0.082	0.441

Table 2: Per-layer Task A pairwise duplicate-detection metrics for the three headline models (LaTa, PhilTa, mT5-base). baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. Both AUROC and cosine gap are computed over the test $n \times n$ cosine matrix with no threshold and no directory routing: AUROC treats same-directory test pairs as positives, and cosine gap is the difference between mean same-directory and mean different-directory cosine. Rows in bold mark the best layer per model, selected by ABTT AUROC.

and UMAP (McInnes et al., 2018) projections for the six most-populated directories. Under baseline mean pooling, points from different directories often overlap. Under pure ABTT with $D = 10$, same-directory points form tighter groups, with winnable-fragment silhouette scores rising by 0.09 to 0.75 on the three headline models. We treat these plots as supporting evidence for the geometry story, not as the layer-selection rule.

4.4 Task B: Directory Routing

Task B connects the layerwise analysis to the retrieval workflow. Table 3 reports autonomous routing, where a learned threshold decides between an existing directory and the new class. Table 4 reports top- k ranking, where the system presents directory options to a simulated expert. The main tables show LaTa, PhilTa, and mT5-base; Appendix C reports the remaining models.

The same training signal also selects attribution layers. We choose the earliest layer within 0.5 percentage points of the best training-set directory accuracy at rank 1 under ABTT. This rule selects LaTa layer 7, PhilTa layer 1, and mT5-base layer 1. The choice is locked before attribution metrics are computed, so the attribution results do not choose their own layer.

The cumulative top- K view describes the scholar-facing shortlist. All six models place the labelled directory first at least 89.4% of the time averaged over five seeds. By Top-2, every model exceeds 95%. Gains after Top-3 are small, so two or three options cover nearly all successful cases in this split.

4.5 Attribution at Retrieval-Selected Layers

Table 6 reports the main attribution comparison. The experiment uses 200 positive query-candidate




Figure 2: AUCROC across normalized layer depth for the six models. LaTa and PhilTa show the clearest mid-layer dip before post-processing. ABTT reduces this dip more reliably than SIF or whitening.




Figure 3: Cosine gap across normalized layer depth for the six models. In the Latin T5 models, the baseline gap nearly vanishes in the middle layers. ABTT restores positive separation across more of the depth profile.

pairs per model for LaTa, PhilTa, and mT5-base. Each model has baseline and ABTT variants for both integrated gradients and retrieval-adapted MaRC.

The clearest quantitative result is rank faithfulness. ABTT improves ρ_{LOO} in all six model-method cells: LaTa IG $-0.001 \rightarrow 0.396$, LaTa MaRC $0.042 \rightarrow 0.401$, PhilTa IG $0.182 \rightarrow 0.607$, PhilTa MaRC $0.190 \rightarrow 0.331$, mT5-base IG $0.164 \rightarrow 0.597$, and mT5-base MaRC $0.250 \rightarrow 0.392$. In those cases, tokens with higher attribution align more closely with leave-one-out drops in the same retrieval score.

The ERASER-style metrics qualify this result. At the global headline thresholds, ABTT improves only 6 of 18 Sufficiency, Comprehensiveness, and MinFrac comparisons. MinFrac@0.80 worsens in all six cells. We therefore interpret attribution as support for the ABTT repair story through leave-one-out rank faithfulness, not as a uniform gain across every rationale metric.

4.6 Example-Level Views

The example figures show what the local explanations look like for one PhilTa retrieval case. Figure 6 shows the integrated-gradient pair matrix before and after ABTT. Figure 7 adds companion self-attention maps for context only, since ABTT is applied after encoding. Figure 8 shows the retrieval-adapted MaRC masks for the same pair.

5 Discussion

The results point to representation geometry as a major source of retrieval failure in this corpus. The Latin T5 encoders do not only vary by depth; several middle layers collapse into spaces dominated by one direction. ABTT removes those dominant directions after pooling and recovers much of the separation needed for cosine retrieval.

Layer diagnostics help explain the failure mode, but they are not a standalone layer selector. This distinction matters for practice. A reader should in-

2D TSNE of labelled fragments at each model's best layer




Figure 4: t-SNE projections of the 1,705 labelled fragments at each model's selected layer for LaTa, PhilTa, and mT5-base. Columns compare baseline mean pooling with pure ABTT using $D = 10$. Colors mark the six most-populated directories.

spect geometry to find collapse and recovery, then choose retrieval layers with a train-only task signal. In our attribution experiment, this produces layers different from the strongest diagnostic collapse layers.

The attribution results also require a narrow reading. Baseline attribution explains raw cosine. ABTT attribution explains ABTT-corrected cosine. Since the scalar score changes, cross-variant comparisons are descriptive. The stable result is improved leave-one-out rank faithfulness across the six headline model-method cells. Sufficiency, Comprehensiveness, and MinFrac show threshold-

dependent tradeoffs.

Limitations

The study covers one domain, medieval canon law, and one language, Latin. The 1,705-fragment corpus is large for this field, but small by current NLP standards. We test six models, and models outside this set might behave differently. ABTT also needs enough text to estimate principal components, so smaller archives need separate study.

Task B simulates a human-in-the-loop workflow by counting a query as successful when its labelled directory appears in the shortlist. This measures

2D UMAP of labelled fragments at each model's best layer




Figure 5: UMAP projections with the same configuration as Figure 4. The same cluster-tightening pattern appears under a different projection method.

ranking quality, not the full judgment load for a working scholar. The attribution analysis is also local to fixed query-candidate scores. This analysis does not give a global explanation of model behavior or replace expert reading of retrieved passages. The consistent attribution effect is improved leave-one-out rank faithfulness; other rationale metrics remain threshold-dependent and should not be reported as a broad attribution-quality gain.

References

Sanjeev Arora, Yingyu Liang, and Tengyu Ma. 2017. [A simple but tough-to-beat baseline for sentence em-](#)

[beddings](#). In *Proceedings of the 5th International Conference on Learning Representations (ICLR)*.

Marc Brinner and Sina Zarrieß. 2023. [Model interpretability and rationale extraction by input mask optimization](#). In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13722–13744. Association for Computational Linguistics.

Jay DeYoung, Sarthak Jain, Nazneen Fatema Rajani, Eric Lehman, Caiming Xiong, Richard Socher, and Byron C. Wallace. 2020. [ERASER: A benchmark to evaluate rationalized NLP models](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4443–4458, Online. Association for Computational Linguistics.

Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Ari-

Task B: Autonomous Routing (main, file-level, τ -thresholded)							
Model	Layer	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)	Overall (base)	Overall (ABTT)
LaTa	1	0.630	0.910	0.916	0.907	0.738	0.909
	2	0.153	0.905	0.913	0.892	0.439	0.900
	3	0.256	0.882	0.842	0.950	0.477	0.908
	4	0.389	0.893	0.718	0.920	0.513	0.903
	5	0.374	0.892	0.715	0.913	0.502	0.900
	6	0.378	0.884	0.712	0.913	0.503	0.895
	7	0.372	0.875	0.715	0.916	0.501	0.890
	8	0.379	0.884	0.706	0.885	0.502	0.885
	9	0.366	0.886	0.724	0.873	0.501	0.881
	10	0.344	0.873	0.752	0.882	0.498	0.876
	11	0.310	0.867	0.786	0.901	0.490	0.880
	12	0.576	0.890	0.932	0.833	0.710	0.868
PhilTa	1	0.591	0.918	0.904	0.901	0.709	0.911
	2	0.479	0.895	0.867	0.885	0.625	0.892
	3	0.193	0.897	0.889	0.879	0.455	0.890
	4	0.148	0.886	0.923	0.873	0.439	0.881
	5	0.252	0.879	0.796	0.873	0.457	0.876
	6	0.338	0.905	0.700	0.817	0.474	0.872
	7	0.301	0.901	0.746	0.830	0.469	0.874
	8	0.262	0.875	0.808	0.885	0.467	0.879
	9	0.303	0.890	0.762	0.870	0.476	0.882
	10	0.299	0.908	0.780	0.793	0.480	0.865
	11	0.256	0.869	0.824	0.836	0.470	0.857
	12	0.338	0.880	0.898	0.876	0.549	0.879
mT5-base	1	0.135	0.877	0.981	0.947	0.453	0.903
	2	0.110	0.890	0.975	0.910	0.436	0.897
	3	0.161	0.794	0.926	0.941	0.449	0.850
	4	0.269	0.787	0.858	0.882	0.491	0.823
	5	0.422	0.746	0.659	0.864	0.512	0.790
	6	0.636	0.755	0.489	0.836	0.580	0.786
	7	0.594	0.793	0.517	0.851	0.565	0.815
	8	0.557	0.809	0.560	0.842	0.558	0.822
	9	0.364	0.841	0.746	0.814	0.508	0.831
	10	0.314	0.794	0.783	0.876	0.491	0.825
	11	0.297	0.763	0.780	0.854	0.479	0.797
	12	0.198	0.834	0.926	0.799	0.472	0.821

Table 3: Per-layer Task B autonomous routing accuracy for the three headline models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. For every test file we compute $s_i = \max_{j \neq i} \cos(e_i, e_j)$ against the rest of the test set and compare to the learned threshold τ (fit on train via F_1 -optimal cut on same-directory vs. different-directory pairs). The decision is binary first (existing if $s_i \geq \tau$, else new); only files predicted existing are then routed to their top-1 neighbour’s directory. We report **existing accuracy** (restricted to files whose true partner is in the test set), **new accuracy** (files that should be flagged as novel), and **overall assignment accuracy**. Rows in bold mark the best layer per model, selected by ABTT overall assignment accuracy.




Figure 6: Derived pairwise heatmap for one PhilTa retrieval case. The left panel shows the baseline representation and the right panel shows the same example after ABTT. Both axes use original decoded tokens. Cell values combine token-token cosine similarity with single-sequence integrated-gradient scores.

Task B: Top-K Ranking (main, single-seed v2)							
Model	Layer	Acc@1 (base)	Acc@1 (ABTT)	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)
LaTa	1	0.721	0.880	0.630	0.910	0.916	0.907
	2	0.383	0.869	0.153	0.905	0.913	0.892
	3	0.350	0.882	0.256	0.882	0.842	0.950
	4	0.303	0.881	0.389	0.893	0.718	0.920
	5	0.297	0.875	0.374	0.892	0.715	0.913
	6	0.297	0.874	0.378	0.884	0.712	0.913
	7	0.302	0.868	0.372	0.875	0.715	0.916
	8	0.302	0.861	0.379	0.884	0.706	0.885
	9	0.309	0.858	0.366	0.886	0.724	0.873
	10	0.317	0.852	0.344	0.873	0.752	0.882
	11	0.331	0.857	0.310	0.867	0.786	0.901
	12	0.689	0.841	0.576	0.890	0.932	0.833
PhilTa	1	0.693	0.883	0.591	0.918	0.904	0.901
	2	0.603	0.869	0.479	0.895	0.867	0.885
	3	0.389	0.865	0.193	0.897	0.889	0.879
	4	0.369	0.859	0.148	0.886	0.923	0.873
	5	0.334	0.854	0.252	0.879	0.796	0.873
	6	0.310	0.845	0.338	0.905	0.700	0.817
	7	0.325	0.847	0.301	0.901	0.746	0.830
	8	0.350	0.857	0.262	0.875	0.808	0.885
	9	0.337	0.857	0.303	0.890	0.762	0.870
	10	0.344	0.833	0.299	0.908	0.780	0.793
	11	0.355	0.829	0.256	0.869	0.824	0.836
	12	0.528	0.854	0.338	0.880	0.898	0.876
mT5-base	1	0.451	0.885	0.135	0.877	0.981	0.947
	2	0.435	0.874	0.110	0.890	0.975	0.910
	3	0.445	0.830	0.161	0.794	0.926	0.941
	4	0.467	0.798	0.269	0.787	0.858	0.882
	5	0.352	0.756	0.422	0.746	0.659	0.864
	6	0.346	0.746	0.636	0.755	0.489	0.836
	7	0.352	0.775	0.594	0.793	0.517	0.851
	8	0.366	0.784	0.557	0.809	0.560	0.842
	9	0.394	0.787	0.364	0.841	0.746	0.814
	10	0.387	0.788	0.314	0.794	0.783	0.876
	11	0.378	0.765	0.297	0.763	0.780	0.854
	12	0.466	0.779	0.198	0.834	0.926	0.799

Table 4: Per-layer Task B top- k ranking metrics for the three headline models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. This table is computed on the single-seed v2 split (no M -seed averaging, since the mseed sweep was run only for SIF-conditioned variants; see Table 14 for the multi-seed SIF+ABTT view). **Acc@1** is directory accuracy at rank 1, with an existing/new decomposition. Rows in bold mark the best layer per model, selected by ABTT Acc@1.




Figure 7: Companion self-attention maps for the same PhilTa example. ABTT is applied after encoding, so attention is included only as supporting context.

vazhagan, and Wei Wang. 2022. [Language-agnostic BERT sentence embedding](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages

878–891. Association for Computational Linguistics.

Xinyu Lu, Nurlan Maharramov, and 1 others. 2025. [KaLM-embedding: Superior training data brings a](#)

Model	Top-1	Top-2	Top-3	Top-4	Top-5
LaTa	90.1 $\pm$ 1.1	96.0 $\pm$ 0.4	97.1 $\pm$ 0.1	97.6 $\pm$ 0.4	97.7 $\pm$ 0.4
PhilTa	90.0 $\pm$ 0.9	95.7 $\pm$ 0.4	96.9 $\pm$ 0.4	97.4 $\pm$ 0.3	97.7 $\pm$ 0.2
LaBSE	90.4 $\pm$ 0.8	95.5 $\pm$ 0.4	96.4 $\pm$ 0.3	97.1 $\pm$ 0.5	97.7 $\pm$ 0.4
Qwen3-0.6B	90.3 $\pm$ 0.6	95.5 $\pm$ 0.5	96.8 $\pm$ 0.3	97.2 $\pm$ 0.5	97.4 $\pm$ 0.4
KaLM-mini	89.4 $\pm$ 0.5	95.1 $\pm$ 0.4	96.4 $\pm$ 0.2	96.9 $\pm$ 0.3	97.3 $\pm$ 0.4
mT5-base	89.8 $\pm$ 0.4	95.6 $\pm$ 0.5	96.6 $\pm$ 0.3	96.9 $\pm$ 0.2	97.2 $\pm$ 0.2

Table 5: Cumulative top- K accuracy for Task B across the six models, reported as mean $\pm$ standard deviation over five query/reference reseeds at a fixed v2 split. Each cell gives the percentage of test queries whose labelled directory appears within the top K options. This variance-aware table uses the SIF+ABTT configurations from the multi-seed sweep; pure-ABTT top- k rankings are reported per layer in Tables 4 and 11.

Model	Method	ρ_{LOO} base	ρ_{LOO} ABTT	$\Delta\rho_{\text{LOO}}$	Suff@25% base $\rightarrow$ ABTT	Comp@25% base $\rightarrow$ ABTT	MinFrac@0.80 base $\rightarrow$ ABTT	ERASER wins
LaTa	IG	-0.001	0.396	+0.397	0.976 $\rightarrow$ 0.681	0.719 $\rightarrow$ 0.581	0.041 $\rightarrow$ 0.379	0/3
	MaRC	0.042	0.401	+0.359	0.534 $\rightarrow$ 0.645	0.361 $\rightarrow$ 0.539	0.383 $\rightarrow$ 0.459	2/3
PhilTa	IG	0.182	0.607	+0.425	0.959 $\rightarrow$ 0.963	0.969 $\rightarrow$ 0.747	0.079 $\rightarrow$ 0.141	1/3
	MaRC	0.190	0.331	+0.141	0.911 $\rightarrow$ 0.754	0.963 $\rightarrow$ 0.521	0.104 $\rightarrow$ 0.358	0/3
mT5-base	IG	0.164	0.597	+0.433	0.892 $\rightarrow$ 1.027	0.110 $\rightarrow$ 0.717	0.116 $\rightarrow$ 0.214	2/3
	MaRC	0.250	0.392	+0.142	0.913 $\rightarrow$ 0.827	0.042 $\rightarrow$ 0.592	0.066 $\rightarrow$ 0.346	1/3

Table 6: **Main-text candidate attribution metrics at the predeclared operational layers for LaTa, PhilTa, and mT5-base.** We report 200 positive query-candidate pairs per model and the two attribution methods used in the paper: integrated gradients (IG) and retrieval-adapted MaRC. The primary faithfulness signal is ρ_{LOO} , which improves under ABTT in 6/6 model–method cells. ERASER-style metrics are shown compactly as baseline $\rightarrow$ ABTT values using one global headline threshold choice: Sufficiency and Comprehensiveness at 25% of query tokens, and MinFrac at recovery threshold 0.80 (lower is better for MinFrac; higher is better otherwise). These are the conventional midpoint thresholds from the available sweep grid and are retained as the global main-text setting instead of selecting model- or method-specific thresholds; the appendix exposes the 10/25/50% sufficiency/comprehensiveness sweep and the 0.70/0.80/0.90/0.95 MinFrac sweep. Boldface inside the ERASER cells marks the better variant, making disagreements visible: ABTT wins 6/18 ERASER-style comparisons even though ρ_{LOO} improves consistently. Cross-variant comparisons are descriptive because ABTT changes both the attribution scores and the cosine function being explained.




Figure 8: Retrieval-adapted MaRC masks for the same PhilTa pair as Figure 6. Top row: mask on the query with the candidate fixed. Bottom row: mask on the candidate with the query fixed. Left column: baseline cosine target. Right column: ABTT-corrected cosine target.

stronger embedding model. Accessed: 2025.

arXiv:1802.03426.

Leland McInnes, John Healy, and James Melville. 2018. UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint*

Jiaqi Mu and Pramod Viswanath. 2018. All-but-the-top: Simple and effective postprocessing for word representations. In *Proceedings of the 6th International*

Conference on Learning Representations (ICLR).

Qwen Team. 2025. [Qwen3-embedding: Advancing text embedding and reranking through foundation models](#). Accessed: 2025.

Frederick Riemenschneider and Anette Frank. 2023. [Exploring large language models for classical philology](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15181–15199. Association for Computational Linguistics.

Jianlin Su, Jiarun Cao, Weijie Liu, and Yangyiwen Ou. 2021. [Whitening sentence representations for better semantics and faster retrieval](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Association for Computational Linguistics.

Mukund Sundararajan, Ankur Taly, and Qiqi Yan. 2017. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, pages 3319–3328. PMLR.

Laurens van der Maaten and Geoffrey Hinton. 2008. Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9:2579–2605.

Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. 2021. [mT5: A massively multilingual pre-trained text-to-text transformer](#). In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 483–498. Association for Computational Linguistics.

A Layer Diagnostics and Attribution Layer Selection

This appendix separates diagnostic layers from operational attribution layers. The diagnostic collapse layers identify where geometry is most anisotropic. The operational attribution layers are selected by training-set retrieval accuracy.

Across the main-model layer grid, PC1 dominance and low effective rank are strongly associated with ABTT gains over baseline. These associations explain why ABTT helps, but they do not replace the train-only retrieval rule used for layer selection.

B Last-Token vs. Mean Pooling

Decoder-only embedding models often use last-token pooling. We compare last-token pooling with mean pooling under the same ABTT pipeline,

train/test split, and ABTT choices. Table 8 reports the layer with the best Task B autonomous-routing accuracy, Task B directory accuracy at rank 1, and cosine gap at the same layer. Mean pooling performs better on assignment accuracy for every model in this experiment, including the decoder-only models.

C Per-Layer Results for Remaining Models

The main text focuses on LaTa, PhilTa, and mT5-base. This appendix reports the same per-layer view for LaBSE, Qwen3-0.6B, and KaLM-mini under the same comparison: baseline mean pooling versus ABTT with D tuned per layer on the train split.

Model	Operational layer	Diagnostic layer	PC1 before	PC1 after	Eff. rank before→after
LaTa	7	8	0.956	0.036	1.34→168.40
PhilTa	1	6	0.862	0.039	1.83→155.02
mT5-base	1	5	1.000	0.044	1.00→151.74

Table 7: Diagnostic collapse layers and operational attribution layers for the three headline models. PC1 and effective-rank values compare baseline geometry with ABTT-D10 at the diagnostic collapse layer. Operational layers come from the train-only retrieval rule.

Table 8: Mean vs. last-token pooling on the balanced split, restricted to ABTT post-processing. Best layer per (model, pooling) selected by Task A assignment accuracy. Task B metrics (Acc@1, cosine gap) reported at that same layer.

Model	Pool	Layer	Method	Assign Acc	Acc@1	Gap
LaTa	Mean	1	ABTT (D=10)	0.909	0.880	0.565
LaTa	LastTok	12	ABTT (D=10)	0.676	0.425	0.276
PhilTa	Mean	1	ABTT (D=10)	0.911	0.883	0.573
PhilTa	LastTok	9	ABTT (D=10)	0.604	0.385	0.200
LaBSE	Mean	11	ABTT (D=10)	0.908	0.885	0.610
LaBSE	LastTok	12	ABTT (D=10)	0.895	0.871	0.586
Qwen3-0.6B	Mean	5	ABTT (D=10)	0.915	0.894	0.545
Qwen3-0.6B	LastTok	27	ABTT (D=10)	0.907	0.885	0.562
KaLM-mini	Mean	1	ABTT (D=10)	0.921	0.902	0.568
KaLM-mini	LastTok	24	ABTT (D=10)	0.875	0.847	0.483
Qwen3-8B	Mean	1	ABTT (D=10)	0.924	0.897	0.543
Qwen3-8B	LastTok	35	ABTT (D=10)	0.910	0.892	0.642

Table 9: Per-layer Task A pairwise metrics for the remaining three models (LaBSE, Qwen3-0.6B, KaLM-mini), under the same two-method comparison as the main paper. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. Both AUROC and cosine gap are computed over the test $n \times n$ cosine matrix with no threshold and no directory routing: AUROC treats same-directory test pairs as positives, and cosine gap is the difference between mean same-directory and mean different-directory cosine. Rows in bold mark the best layer per model, selected by ABTT AUROC.

Task A: Pairwise Duplicate Detection (appendix models)					
Model	Layer	AUROC (base)	AUROC (ABTT)	Cosine gap (base)	Cosine gap (ABTT)
LaBSE	1	0.807	0.976	0.060	0.519
	2	0.835	0.978	0.038	0.531
	3	0.855	0.977	0.063	0.543
	4	0.864	0.976	0.074	0.525
	5	0.866	0.967	0.066	0.500
	6	0.895	0.974	0.078	0.532
	7	0.889	0.973	0.089	0.523
	8	0.881	0.973	0.080	0.522
	9	0.918	0.976	0.091	0.564
	10	0.936	0.980	0.103	0.584
	11	0.961	0.984	0.061	0.610
	12	0.957	0.979	0.183	0.592
Qwen3-0.6B	1	0.859	0.978	0.033	0.564
	2	0.864	0.972	0.024	0.553
	3	0.864	0.980	0.037	0.555
	4	0.861	0.985	0.036	0.546
	5	0.862	0.980	0.030	0.545
	6	0.868	0.983	0.035	0.545
	7	0.877	0.980	0.033	0.544
	8	0.884	0.972	0.037	0.533
	9	0.882	0.977	0.034	0.532
	10	0.888	0.976	0.031	0.530
	11	0.896	0.970	0.030	0.516
	12	0.903	0.973	0.027	0.531

(continued on next page)

Task A: Pairwise Duplicate Detection (appendix models)					
Model	Layer	AUROC (base)	AUROC (ABTT)	Cosine gap (base)	Cosine gap (ABTT)
	13	0.908	0.974	0.026	0.524
	14	0.906	0.974	0.024	0.519
	15	0.902	0.971	0.022	0.514
	16	0.911	0.975	0.026	0.509
	17	0.917	0.978	0.032	0.521
	18	0.934	0.982	0.032	0.537
	19	0.937	0.980	0.035	0.541
	20	0.943	0.982	0.041	0.559
	21	0.953	0.982	0.047	0.562
	22	0.963	0.983	0.044	0.583
	23	0.960	0.983	0.038	0.578
	24	0.962	0.984	0.032	0.576
	25	0.960	0.983	0.027	0.578
	26	0.965	0.985	0.024	0.568
	27	0.966	0.982	0.017	0.559
	28	0.958	0.975	0.098	0.556
	KaLM-mini 1	0.888	0.980	0.064	0.568
	2	0.908	0.983	0.044	0.566
	3	0.882	0.980	0.062	0.552
	4	0.872	0.983	0.060	0.553
	5	0.862	0.982	0.062	0.541
	6	0.873	0.982	0.062	0.528
	7	0.889	0.977	0.072	0.524
	8	0.891	0.975	0.072	0.523
	9	0.893	0.972	0.066	0.507
	10	0.893	0.970	0.067	0.494
	11	0.895	0.969	0.061	0.474
	12	0.905	0.968	0.071	0.475
	13	0.901	0.966	0.066	0.467
	14	0.904	0.970	0.063	0.471
	15	0.907	0.972	0.067	0.468
	16	0.917	0.970	0.071	0.488
	17	0.926	0.973	0.075	0.501
	18	0.944	0.976	0.060	0.519
	19	0.955	0.978	0.061	0.530
	20	0.959	0.979	0.054	0.554
	21	0.961	0.980	0.043	0.554
	22	0.969	0.984	0.050	0.554
	23	0.970	0.985	0.056	0.552
	24	0.954	0.981	0.243	0.520

Table 10: Per-layer Task B autonomous routing accuracy for the remaining three models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. For every test file we compute $s_i = \max_{j \neq i} \cos(e_i, e_j)$ against the rest of the test set and compare to the learned threshold τ (fit on train via F_1 -optimal cut on same-directory vs. different-directory pairs). The decision is binary first (existing if $s_i \geq \tau$, else new); only files predicted existing are then routed to their top-1 neighbour’s directory. We report **existing accuracy** (restricted to files whose true partner is in the test set), **new accuracy** (files that should be flagged as novel), and **overall assignment accuracy**. Rows in bold mark the best layer per model, selected by ABTT overall assignment accuracy.

Task B: Autonomous Routing (appendix models)							
Model	Layer	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)	Overall (base)	Overall (ABTT)
LaBSE	1	0.222	0.879	0.898	0.873	0.477	0.876
	2	0.381	0.875	0.811	0.827	0.543	0.857
	3	0.277	0.886	0.907	0.867	0.514	0.879
	4	0.372	0.897	0.885	0.836	0.565	0.874
	5	0.327	0.845	0.901	0.851	0.543	0.847
	6	0.548	0.860	0.882	0.851	0.674	0.857

(continued on next page)

Task B: Autonomous Routing (appendix models)							
Model	Layer	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)	Overall (base)	Overall (ABTT)
	7	0.398	0.849	0.892	0.830	0.584	0.841
	8	0.385	0.869	0.870	0.743	0.568	0.822
	9	0.421	0.871	0.929	0.817	0.612	0.851
	10	0.606	0.892	0.892	0.805	0.713	0.859
	11	0.785	0.923	0.913	0.882	0.833	0.908
	12	0.729	0.914	0.938	0.864	0.808	0.895
Qwen3-0.6B	1	0.264	0.921	0.960	0.885	0.526	0.908
	2	0.437	0.869	0.814	0.916	0.579	0.887
	3	0.250	0.914	0.947	0.907	0.513	0.911
	4	0.411	0.936	0.848	0.842	0.576	0.901
	5	0.181	0.910	0.932	0.923	0.464	0.915
	6	0.422	0.935	0.830	0.796	0.576	0.882
	7	0.250	0.890	0.947	0.876	0.513	0.885
	8	0.409	0.931	0.895	0.799	0.592	0.881
	9	0.460	0.901	0.851	0.851	0.607	0.882
	10	0.533	0.892	0.836	0.864	0.647	0.881
	11	0.293	0.895	0.978	0.814	0.551	0.865
	12	0.407	0.912	0.938	0.817	0.607	0.876
	13	0.432	0.908	0.944	0.836	0.625	0.881
	14	0.456	0.907	0.926	0.811	0.633	0.871
	15	0.469	0.901	0.920	0.845	0.639	0.880
	16	0.639	0.897	0.830	0.808	0.711	0.864
	17	0.480	0.886	0.935	0.851	0.652	0.873
	18	0.553	0.882	0.944	0.839	0.700	0.866
	19	0.514	0.871	0.966	0.858	0.684	0.866
	20	0.654	0.865	0.898	0.892	0.746	0.875
	21	0.647	0.912	0.950	0.898	0.761	0.907
	22	0.731	0.888	0.938	0.916	0.809	0.899
	23	0.622	0.914	0.966	0.895	0.752	0.907
	24	0.748	0.901	0.935	0.935	0.818	0.914
	25	0.574	0.893	0.975	0.923	0.725	0.904
	26	0.613	0.899	0.975	0.895	0.749	0.897
	27	0.793	0.877	0.833	0.923	0.808	0.894
	28	0.796	0.877	0.867	0.913	0.823	0.890
KaLM-mini	1	0.391	0.921	0.907	0.920	0.585	0.921
	2	0.493	0.916	0.848	0.879	0.627	0.902
	3	0.443	0.923	0.870	0.907	0.604	0.917
	4	0.308	0.931	0.935	0.895	0.544	0.917
	5	0.338	0.905	0.882	0.923	0.543	0.911
	6	0.256	0.895	0.910	0.920	0.502	0.904
	7	0.475	0.864	0.854	0.944	0.618	0.894
	8	0.467	0.862	0.870	0.913	0.619	0.881
	9	0.469	0.873	0.885	0.870	0.626	0.872
	10	0.449	0.875	0.885	0.870	0.613	0.873
	11	0.458	0.862	0.885	0.864	0.619	0.862
	12	0.507	0.862	0.895	0.904	0.653	0.878
	13	0.540	0.860	0.895	0.870	0.674	0.864
	14	0.593	0.836	0.845	0.885	0.688	0.854
	15	0.566	0.864	0.898	0.827	0.691	0.850
	16	0.550	0.849	0.935	0.830	0.695	0.841
	17	0.675	0.864	0.920	0.889	0.767	0.873
	18	0.697	0.875	0.960	0.898	0.796	0.883
	19	0.721	0.873	0.966	0.892	0.814	0.880
	20	0.705	0.791	0.966	0.913	0.803	0.837
	21	0.727	0.811	0.966	0.926	0.817	0.854
	22	0.850	0.879	0.916	0.923	0.875	0.895
	23	0.849	0.884	0.920	0.898	0.875	0.889
	24	0.787	0.892	0.944	0.867	0.846	0.882

Table 11: Per-layer Task B top- k ranking metrics for the remaining three models. baseline is mean pooling without correction; abtt_optimal applies ABTT (no SIF weighting) with D tuned per layer on the train split. This table is computed on the single-seed v2 split (no M -seed averaging, since the mseed sweep was run only for SIF-conditioned variants; see Table 14 for the multi-seed SIF+ABTT view). **Acc@1** is directory accuracy at rank 1, with an existing/new decomposition. Rows in bold mark the best layer per model, selected by ABTT Acc@1.

Task B: Top-K Ranking (appendix models, single-seed)							
Model	Layer	Acc@1 (base)	Acc@1 (ABTT)	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)
LaBSE	1	0.465	0.851	0.222	0.879	0.898	0.873
	2	0.494	0.830	0.381	0.875	0.811	0.827
	3	0.500	0.848	0.277	0.886	0.907	0.867
	4	0.542	0.841	0.372	0.897	0.885	0.836
	5	0.521	0.811	0.327	0.845	0.901	0.851
	6	0.640	0.831	0.548	0.860	0.882	0.851
	7	0.561	0.810	0.398	0.849	0.892	0.830
	8	0.544	0.779	0.385	0.869	0.870	0.743
	9	0.608	0.823	0.421	0.871	0.929	0.817
	10	0.693	0.832	0.606	0.892	0.892	0.805
	11	0.815	0.885	0.785	0.923	0.913	0.882
	12	0.791	0.868	0.729	0.914	0.938	0.864
Qwen3-0.6B	1	0.521	0.875	0.264	0.921	0.960	0.885
	2	0.558	0.871	0.437	0.869	0.814	0.916
	3	0.508	0.890	0.250	0.914	0.947	0.907
	4	0.559	0.874	0.411	0.936	0.848	0.842
	5	0.450	0.894	0.181	0.910	0.932	0.923
	6	0.551	0.858	0.422	0.935	0.830	0.796
	7	0.507	0.859	0.250	0.890	0.947	0.876
	8	0.580	0.852	0.409	0.931	0.895	0.799
	9	0.587	0.852	0.460	0.901	0.851	0.851
	10	0.622	0.848	0.533	0.892	0.836	0.864
	11	0.548	0.833	0.293	0.895	0.978	0.814
	12	0.603	0.848	0.407	0.912	0.938	0.817
	13	0.618	0.851	0.432	0.908	0.944	0.836
	14	0.621	0.840	0.456	0.907	0.926	0.811
	15	0.627	0.848	0.469	0.901	0.920	0.845
	16	0.682	0.830	0.639	0.897	0.830	0.808
	17	0.641	0.840	0.480	0.886	0.935	0.851
	18	0.691	0.837	0.553	0.882	0.944	0.839
	19	0.676	0.837	0.514	0.871	0.966	0.858
	20	0.733	0.852	0.654	0.865	0.898	0.892
	21	0.752	0.881	0.647	0.912	0.950	0.898
	22	0.795	0.876	0.731	0.888	0.938	0.916
	23	0.742	0.886	0.622	0.914	0.966	0.895
	24	0.808	0.896	0.748	0.901	0.935	0.935
	25	0.718	0.886	0.574	0.893	0.975	0.923
	26	0.740	0.881	0.613	0.899	0.975	0.895
	27	0.787	0.876	0.793	0.877	0.833	0.923
	28	0.803	0.872	0.796	0.877	0.867	0.913
KaLM-mini	1	0.577	0.902	0.391	0.921	0.907	0.920
	2	0.607	0.879	0.493	0.916	0.848	0.879
	3	0.587	0.894	0.443	0.923	0.870	0.907
	4	0.538	0.887	0.308	0.931	0.935	0.895
	5	0.524	0.882	0.338	0.905	0.882	0.923
	6	0.490	0.876	0.256	0.895	0.910	0.920
	7	0.592	0.873	0.475	0.864	0.854	0.944
	8	0.597	0.860	0.467	0.862	0.870	0.913
	9	0.601	0.843	0.469	0.873	0.885	0.870
	10	0.584	0.844	0.449	0.875	0.885	0.870
	11	0.586	0.836	0.458	0.862	0.885	0.864
	12	0.622	0.850	0.507	0.862	0.895	0.904
	13	0.633	0.840	0.540	0.860	0.895	0.870
	14	0.636	0.829	0.593	0.836	0.845	0.885
	15	0.650	0.817	0.566	0.864	0.898	0.827
	16	0.676	0.819	0.550	0.849	0.935	0.830

(continued on next page)

Task B: Top-K Ranking (appendix models, single-seed)							
Model	Layer	Acc@1 (base)	Acc@1 (ABTT)	Existing (base)	Existing (ABTT)	New (base)	New (ABTT)
	17	0.747	0.852	0.675	0.864	0.920	0.889
	18	0.777	0.865	0.697	0.875	0.960	0.898
	19	0.801	0.861	0.721	0.873	0.966	0.892
	20	0.789	0.817	0.705	0.791	0.966	0.913
	21	0.802	0.831	0.727	0.811	0.966	0.926
	22	0.857	0.875	0.850	0.879	0.916	0.923
	23	0.859	0.872	0.849	0.884	0.920	0.898
	24	0.833	0.858	0.787	0.892	0.944	0.867

D SIF-Variant Post-Processing Sweeps

The main text uses pure ABTT as the paper-facing repair. This appendix reports SIF-only, SIF plus fixed- D ABTT, and SIF plus tuned- D ABTT. The tables keep one primary metric per task so the four-method grid stays readable.

Table 12: Per-layer Task A AUROC across the SIF-conditioned post-processing suite for all six models. `sif_only` replaces mean pooling with SIF-weighted pooling; `sif_abtt_fixed` adds ABTT with fixed $D=10$; `sif_abtt_optimal` tunes D per layer on the train split. Gap is omitted to keep the table narrow; the pure-ABTT comparison (not SIF-conditioned) is in Tables 2 and 9. Rows in bold mark the best layer per model, selected by SIF+ABTT (D-tuned) AUROC.

Model	Layer	Task A: SIF-suite AUROC (appendix)			
		AUROC (base)	AUROC (SIF)	AUROC (SIF+ABTT ₁₀)	AUROC (SIF+ABTT)
LaTa	1	0.933	0.944	0.971	0.974
	2	0.565	0.936	0.964	0.970
	3	0.515	0.920	0.966	0.964
	4	0.502	0.890	0.969	0.968
	5	0.498	0.878	0.969	0.969
	6	0.497	0.865	0.964	0.964
	7	0.499	0.858	0.965	0.965
	8	0.504	0.842	0.964	0.964
	9	0.504	0.854	0.964	0.964
	10	0.507	0.859	0.965	0.965
	11	0.510	0.867	0.965	0.965
	12	0.938	0.948	0.963	0.963
PhilTa	1	0.939	0.953	0.973	0.973
	2	0.911	0.941	0.970	0.970
	3	0.601	0.929	0.966	0.966
	4	0.551	0.916	0.970	0.970
	5	0.543	0.904	0.971	0.976
	6	0.541	0.884	0.980	0.980
	7	0.542	0.891	0.980	0.980
	8	0.543	0.895	0.981	0.981
	9	0.539	0.882	0.977	0.977
	10	0.539	0.874	0.977	0.977
	11	0.540	0.872	0.978	0.978
	12	0.911	0.929	0.968	0.968
mT5-base	1	0.823	0.846	0.971	0.971
	2	0.823	0.835	0.964	0.964
	3	0.805	0.809	0.956	0.956
	4	0.800	0.810	0.966	0.966
	5	0.654	0.672	0.972	0.972
	6	0.657	0.679	0.970	0.970
	7	0.657	0.680	0.966	0.966
	8	0.656	0.682	0.967	0.967
	9	0.660	0.685	0.966	0.966
	10	0.659	0.685	0.964	0.964
	11	0.660	0.690	0.964	0.964
	12	0.839	0.834	0.968	0.968
LaBSE	1	0.807	0.819	0.974	0.974
	2	0.835	0.842	0.974	0.974
	3	0.855	0.864	0.968	0.968
	4	0.864	0.871	0.966	0.966
	5	0.866	0.870	0.963	0.963
	6	0.895	0.897	0.964	0.964
	7	0.889	0.893	0.961	0.961
	8	0.881	0.885	0.966	0.966
	9	0.918	0.916	0.971	0.971
	10	0.936	0.936	0.974	0.974
	11	0.961	0.959	0.980	0.980
	12	0.957	0.956	0.978	0.982
Qwen3-0.6B	1	0.859	0.852	0.982	0.980
	2	0.864	0.854	0.972	0.972
	3	0.864	0.859	0.977	0.975
	4	0.861	0.861	0.976	0.977
	5	0.862	0.859	0.971	0.971
	6	0.868	0.862	0.974	0.974
	7	0.877	0.867	0.971	0.971

(continued on next page)

Task A: SIF-suite AUROC (appendix)					
Model	Layer	AUROC (base)	AUROC (SIF)	AUROC (SIF+ABTT ₁₀)	AUROC (SIF+ABTT)
	8	0.884	0.869	0.969	0.969
	9	0.882	0.867	0.968	0.968
	10	0.888	0.871	0.966	0.966
	11	0.896	0.876	0.964	0.964
	12	0.903	0.882	0.967	0.967
	13	0.908	0.888	0.967	0.967
	14	0.906	0.887	0.966	0.966
	15	0.902	0.884	0.967	0.967
	16	0.911	0.895	0.965	0.965
	17	0.917	0.900	0.967	0.974
	18	0.934	0.915	0.966	0.972
	19	0.937	0.916	0.966	0.976
	20	0.943	0.919	0.971	0.971
	21	0.953	0.931	0.971	0.971
	22	0.963	0.941	0.973	0.973
	23	0.960	0.941	0.973	0.973
	24	0.962	0.942	0.975	0.973
	25	0.960	0.938	0.978	0.974
	26	0.965	0.941	0.979	0.979
	27	0.966	0.935	0.977	0.977
	28	0.958	0.925	0.967	0.967
KaLM-mini	1	0.888	0.881	0.981	0.979
	2	0.908	0.900	0.981	0.981
	3	0.882	0.882	0.977	0.977
	4	0.872	0.873	0.974	0.974
	5	0.862	0.863	0.971	0.971
	6	0.873	0.867	0.966	0.966
	7	0.889	0.879	0.971	0.971
	8	0.891	0.880	0.971	0.975
	9	0.893	0.883	0.966	0.966
	10	0.893	0.883	0.963	0.963
	11	0.895	0.886	0.968	0.968
	12	0.905	0.897	0.965	0.965
	13	0.901	0.894	0.963	0.963
	14	0.904	0.898	0.968	0.968
	15	0.907	0.901	0.968	0.968
	16	0.917	0.913	0.966	0.966
	17	0.926	0.923	0.968	0.968
	18	0.944	0.938	0.976	0.976
	19	0.955	0.949	0.978	0.978
	20	0.959	0.955	0.978	0.978
	21	0.961	0.957	0.979	0.982
	22	0.969	0.965	0.982	0.982
	23	0.970	0.963	0.986	0.984
	24	0.954	0.954	0.975	0.975

Table 13: Per-layer Task B overall assignment accuracy across the SIF-conditioned suite for all six models. Restricted to overall routing accuracy (existing/new decomposition omitted) to keep the table compact; see Tables 3 and 10 for the pure-ABTT pairwise comparison with existing/new broken out. Rows in bold mark the best layer per model, selected by SIF+ABTT (D-tuned) overall assignment accuracy.

Task B: SIF-suite routing (appendix)					
Model	Layer	Overall (base)	Overall (SIF)	Overall (SIF+ABTT ₁₀)	Overall (SIF+ABTT)
LaTa	1	0.738	0.832	0.920	0.899
	2	0.439	0.791	0.904	0.900
	3	0.477	0.720	0.910	0.907
	4	0.513	0.668	0.900	0.902
	5	0.502	0.648	0.906	0.906
	6	0.503	0.668	0.895	0.895

(continued on next page)

Task B: SIF-suite routing (appendix)					
Model	Layer	Overall (base)	Overall (SIF)	Overall (SIF+ABTT ₁₀)	Overall (SIF+ABTT)
	7	0.501	0.642	0.903	0.903
	8	0.502	0.614	0.892	0.892
	9	0.501	0.618	0.889	0.889
	10	0.498	0.655	0.890	0.890
	11	0.490	0.621	0.880	0.880
	12	0.710	0.804	0.902	0.902
PhilTa	1	0.709	0.790	0.914	0.914
	2	0.625	0.738	0.892	0.892
	3	0.455	0.713	0.892	0.892
	4	0.439	0.663	0.886	0.886
	5	0.457	0.626	0.894	0.897
	6	0.474	0.612	0.906	0.906
	7	0.469	0.625	0.906	0.906
	8	0.467	0.614	0.886	0.886
	9	0.476	0.590	0.895	0.895
	10	0.480	0.580	0.888	0.888
	11	0.470	0.565	0.886	0.886
	12	0.549	0.667	0.899	0.899
mT5-base	1	0.453	0.545	0.907	0.907
	2	0.436	0.537	0.903	0.903
	3	0.449	0.568	0.889	0.889
	4	0.491	0.488	0.857	0.857
	5	0.512	0.483	0.826	0.826
	6	0.580	0.558	0.817	0.817
	7	0.565	0.554	0.832	0.832
	8	0.558	0.548	0.834	0.834
	9	0.508	0.495	0.845	0.845
	10	0.491	0.495	0.832	0.832
	11	0.479	0.484	0.816	0.816
	12	0.472	0.541	0.858	0.858
LaBSE	1	0.477	0.521	0.886	0.886
	2	0.543	0.508	0.875	0.875
	3	0.514	0.577	0.860	0.860
	4	0.565	0.587	0.869	0.869
	5	0.543	0.598	0.857	0.857
	6	0.674	0.648	0.872	0.872
	7	0.584	0.634	0.858	0.858
	8	0.568	0.571	0.854	0.854
	9	0.612	0.613	0.859	0.859
	10	0.713	0.698	0.864	0.864
	11	0.833	0.807	0.896	0.896
	12	0.808	0.834	0.890	0.880
Qwen3-0.6B	1	0.526	0.566	0.895	0.909
	2	0.579	0.524	0.894	0.894
	3	0.513	0.537	0.904	0.901
	4	0.576	0.531	0.899	0.904
	5	0.464	0.537	0.908	0.908
	6	0.576	0.526	0.906	0.906
	7	0.513	0.573	0.911	0.911
	8	0.592	0.627	0.911	0.911
	9	0.607	0.576	0.910	0.910
	10	0.647	0.599	0.900	0.900
	11	0.551	0.606	0.893	0.893
	12	0.607	0.663	0.900	0.900
	13	0.625	0.675	0.897	0.897
	14	0.633	0.571	0.893	0.893
	15	0.639	0.600	0.893	0.893
	16	0.711	0.679	0.880	0.880
	17	0.652	0.686	0.900	0.876
	18	0.700	0.731	0.899	0.882
	19	0.684	0.700	0.886	0.850
	20	0.746	0.676	0.882	0.882

(continued on next page)

Task B: SIF-suite routing (appendix)					
Model	Layer	Overall (base)	Overall (SIF)	Overall (SIF+ABTT ₁₀)	Overall (SIF+ABTT)
	21	0.761	0.739	0.892	0.892
	22	0.809	0.719	0.893	0.893
	23	0.752	0.759	0.893	0.893
	24	0.818	0.706	0.900	0.899
	25	0.725	0.748	0.897	0.893
	26	0.749	0.772	0.903	0.903
	27	0.808	0.760	0.899	0.899
	28	0.823	0.774	0.890	0.890
KaLM-mini	1	0.585	0.604	0.906	0.908
	2	0.627	0.594	0.908	0.908
	3	0.604	0.576	0.907	0.907
	4	0.544	0.577	0.901	0.901
	5	0.543	0.533	0.897	0.897
	6	0.502	0.558	0.900	0.900
	7	0.618	0.593	0.901	0.901
	8	0.619	0.607	0.892	0.878
	9	0.626	0.663	0.876	0.876
	10	0.613	0.603	0.866	0.866
	11	0.619	0.627	0.874	0.874
	12	0.653	0.642	0.874	0.874
	13	0.674	0.666	0.864	0.864
	14	0.688	0.648	0.861	0.861
	15	0.691	0.698	0.866	0.866
	16	0.695	0.738	0.859	0.859
	17	0.767	0.763	0.882	0.882
	18	0.796	0.796	0.879	0.879
	19	0.814	0.804	0.878	0.878
	20	0.803	0.821	0.876	0.876
	21	0.817	0.832	0.874	0.876
	22	0.875	0.857	0.876	0.876
	23	0.875	0.862	0.880	0.881
	24	0.846	0.847	0.869	0.869

Table 14: Per-layer Task B top- k ranking accuracy across all six models, averaged over $M = 5$ query/reference reseeds of the v2 train/test split (mean $\pm$ std). `sif_abtt_optimal` applies SIF weighting plus ABTT with D tuned per layer on the train split; the mseed sweep was not run for the pure `abtt_optimal` variant reported in the main paper, so we include this table to give a variance-aware view of the SIF-conditioned best variant. Overall assignment accuracy coincides with Acc@1 in the mseed pipeline, so we report the existing/new decomposition instead. Rows in bold mark the best layer per model, selected by SIF+ABTT Acc@1.

Task B: Top-K Ranking (appendix, 5-seed mean $\pm$ std, SIF+ABTT)							
Model	Layer	Acc@1 (base)	Acc@1 (SIF+ABTT)	Existing (base)	Existing (SIF+ABTT)	New (base)	New (SIF+ABTT)
LaTa	1	0.731 $\pm$ 0.009	0.889 $\pm$ 0.009	0.562 $\pm$ 0.004	0.842 $\pm$ 0.009	0.949 $\pm$ 0.011	0.950 $\pm$ 0.011
	2	0.428 $\pm$ 0.011	0.887 $\pm$ 0.005	0.048 $\pm$ 0.006	0.840 $\pm$ 0.005	0.920 $\pm$ 0.008	0.949 $\pm$ 0.009
	3	0.392 $\pm$ 0.012	0.889 $\pm$ 0.008	0.046 $\pm$ 0.008	0.845 $\pm$ 0.006	0.840 $\pm$ 0.020	0.946 $\pm$ 0.013
	4	0.330 $\pm$ 0.018	0.888 $\pm$ 0.008	0.047 $\pm$ 0.008	0.858 $\pm$ 0.007	0.696 $\pm$ 0.030	0.927 $\pm$ 0.013
	5	0.326 $\pm$ 0.020	0.890 $\pm$ 0.008	0.039 $\pm$ 0.004	0.856 $\pm$ 0.005	0.698 $\pm$ 0.028	0.935 $\pm$ 0.014
	6	0.321 $\pm$ 0.018	0.890 $\pm$ 0.008	0.038 $\pm$ 0.007	0.861 $\pm$ 0.008	0.688 $\pm$ 0.026	0.927 $\pm$ 0.014
	7	0.327 $\pm$ 0.020	0.895 $\pm$ 0.006	0.044 $\pm$ 0.007	0.846 $\pm$ 0.006	0.694 $\pm$ 0.029	0.957 $\pm$ 0.010
	8	0.325 $\pm$ 0.022	0.883 $\pm$ 0.007	0.046 $\pm$ 0.010	0.832 $\pm$ 0.006	0.685 $\pm$ 0.032	0.949 $\pm$ 0.011
	9	0.331 $\pm$ 0.019	0.885 $\pm$ 0.008	0.045 $\pm$ 0.010	0.846 $\pm$ 0.006	0.701 $\pm$ 0.024	0.935 $\pm$ 0.013
	10	0.346 $\pm$ 0.017	0.888 $\pm$ 0.002	0.044 $\pm$ 0.010	0.844 $\pm$ 0.007	0.736 $\pm$ 0.022	0.944 $\pm$ 0.010
	11	0.360 $\pm$ 0.018	0.881 $\pm$ 0.006	0.045 $\pm$ 0.009	0.860 $\pm$ 0.008	0.767 $\pm$ 0.028	0.907 $\pm$ 0.011
	12	0.707 $\pm$ 0.015	0.889 $\pm$ 0.005	0.511 $\pm$ 0.017	0.834 $\pm$ 0.005	0.961 $\pm$ 0.008	0.961 $\pm$ 0.010
PhilTa	1	0.706 $\pm$ 0.016	0.898 $\pm$ 0.007	0.532 $\pm$ 0.018	0.856 $\pm$ 0.009	0.931 $\pm$ 0.012	0.953 $\pm$ 0.011
	2	0.625 $\pm$ 0.018	0.889 $\pm$ 0.009	0.408 $\pm$ 0.016	0.866 $\pm$ 0.009	0.905 $\pm$ 0.017	0.919 $\pm$ 0.017
	3	0.439 $\pm$ 0.013	0.892 $\pm$ 0.006	0.071 $\pm$ 0.014	0.856 $\pm$ 0.008	0.917 $\pm$ 0.013	0.939 $\pm$ 0.009
	4	0.420 $\pm$ 0.016	0.893 $\pm$ 0.009	0.027 $\pm$ 0.007	0.858 $\pm$ 0.006	0.928 $\pm$ 0.015	0.938 $\pm$ 0.015

(continued on next page)

Task B: Top-K Ranking (appendix, 5-seed mean $\pm$ std, SIF+ABTT)							
Model	Layer	Acc@1 (base)	Acc@1 (SIF+ABTT)	Existing (base)	Existing (SIF+ABTT)	New (base)	New (SIF+ABTT)
	5	0.374 $\pm$ 0.015	0.896 $\pm$ 0.010	0.042 $\pm$ 0.011	0.872 $\pm$ 0.007	0.805 $\pm$ 0.018	0.925 $\pm$ 0.014
	6	0.338 $\pm$ 0.015	0.897 $\pm$ 0.008	0.059 $\pm$ 0.010	0.861 $\pm$ 0.007	0.698 $\pm$ 0.025	0.945 $\pm$ 0.014
	7	0.361 $\pm$ 0.016	0.900 $\pm$ 0.009	0.056 $\pm$ 0.010	0.859 $\pm$ 0.005	0.757 $\pm$ 0.021	0.952 $\pm$ 0.016
	8	0.390 $\pm$ 0.016	0.891 $\pm$ 0.010	0.058 $\pm$ 0.012	0.858 $\pm$ 0.007	0.820 $\pm$ 0.017	0.933 $\pm$ 0.018
	9	0.372 $\pm$ 0.015	0.894 $\pm$ 0.006	0.066 $\pm$ 0.012	0.861 $\pm$ 0.006	0.768 $\pm$ 0.026	0.936 $\pm$ 0.011
	10	0.377 $\pm$ 0.013	0.892 $\pm$ 0.006	0.066 $\pm$ 0.012	0.855 $\pm$ 0.003	0.781 $\pm$ 0.021	0.941 $\pm$ 0.015
	11	0.394 $\pm$ 0.013	0.889 $\pm$ 0.006	0.058 $\pm$ 0.008	0.851 $\pm$ 0.008	0.829 $\pm$ 0.018	0.937 $\pm$ 0.008
	12	0.563 $\pm$ 0.010	0.896 $\pm$ 0.007	0.279 $\pm$ 0.011	0.863 $\pm$ 0.009	0.932 $\pm$ 0.005	0.938 $\pm$ 0.014
	1	0.486 $\pm$ 0.012	0.898 $\pm$ 0.004	0.100 $\pm$ 0.007	0.856 $\pm$ 0.009	0.987 $\pm$ 0.007	0.953 $\pm$ 0.010
	2	0.475 $\pm$ 0.012	0.892 $\pm$ 0.004	0.082 $\pm$ 0.008	0.851 $\pm$ 0.013	0.984 $\pm$ 0.008	0.946 $\pm$ 0.013
	3	0.482 $\pm$ 0.014	0.872 $\pm$ 0.008	0.123 $\pm$ 0.008	0.809 $\pm$ 0.012	0.947 $\pm$ 0.008	0.954 $\pm$ 0.008
	4	0.498 $\pm$ 0.013	0.852 $\pm$ 0.006	0.205 $\pm$ 0.006	0.772 $\pm$ 0.012	0.879 $\pm$ 0.018	0.956 $\pm$ 0.010
mT5-base	5	0.370 $\pm$ 0.013	0.825 $\pm$ 0.009	0.155 $\pm$ 0.005	0.724 $\pm$ 0.010	0.649 $\pm$ 0.020	0.956 $\pm$ 0.005
	6	0.359 $\pm$ 0.011	0.809 $\pm$ 0.006	0.263 $\pm$ 0.021	0.765 $\pm$ 0.013	0.483 $\pm$ 0.018	0.865 $\pm$ 0.018
	7	0.368 $\pm$ 0.007	0.827 $\pm$ 0.007	0.256 $\pm$ 0.012	0.758 $\pm$ 0.011	0.512 $\pm$ 0.018	0.916 $\pm$ 0.012
	8	0.380 $\pm$ 0.007	0.824 $\pm$ 0.012	0.245 $\pm$ 0.007	0.771 $\pm$ 0.013	0.555 $\pm$ 0.017	0.892 $\pm$ 0.015
	9	0.426 $\pm$ 0.012	0.836 $\pm$ 0.010	0.171 $\pm$ 0.006	0.781 $\pm$ 0.014	0.756 $\pm$ 0.023	0.907 $\pm$ 0.006
	10	0.420 $\pm$ 0.007	0.821 $\pm$ 0.016	0.141 $\pm$ 0.011	0.765 $\pm$ 0.016	0.782 $\pm$ 0.019	0.893 $\pm$ 0.017
	11	0.418 $\pm$ 0.012	0.803 $\pm$ 0.010	0.132 $\pm$ 0.010	0.719 $\pm$ 0.013	0.788 $\pm$ 0.014	0.910 $\pm$ 0.010
	12	0.506 $\pm$ 0.008	0.840 $\pm$ 0.009	0.170 $\pm$ 0.008	0.770 $\pm$ 0.017	0.943 $\pm$ 0.012	0.930 $\pm$ 0.008
	1	0.501 $\pm$ 0.011	0.858 $\pm$ 0.008	0.170 $\pm$ 0.002	0.838 $\pm$ 0.007	0.930 $\pm$ 0.014	0.884 $\pm$ 0.010
	2	0.522 $\pm$ 0.010	0.861 $\pm$ 0.006	0.272 $\pm$ 0.003	0.817 $\pm$ 0.004	0.844 $\pm$ 0.022	0.917 $\pm$ 0.013
	3	0.535 $\pm$ 0.008	0.845 $\pm$ 0.005	0.217 $\pm$ 0.007	0.831 $\pm$ 0.010	0.946 $\pm$ 0.007	0.864 $\pm$ 0.009
	4	0.568 $\pm$ 0.008	0.849 $\pm$ 0.007	0.298 $\pm$ 0.006	0.837 $\pm$ 0.008	0.918 $\pm$ 0.014	0.865 $\pm$ 0.015
LaBSE	5	0.550 $\pm$ 0.011	0.845 $\pm$ 0.007	0.254 $\pm$ 0.011	0.794 $\pm$ 0.006	0.933 $\pm$ 0.013	0.911 $\pm$ 0.015
	6	0.660 $\pm$ 0.012	0.858 $\pm$ 0.003	0.460 $\pm$ 0.021	0.816 $\pm$ 0.009	0.918 $\pm$ 0.018	0.912 $\pm$ 0.016
	7	0.595 $\pm$ 0.010	0.850 $\pm$ 0.003	0.333 $\pm$ 0.019	0.808 $\pm$ 0.006	0.934 $\pm$ 0.015	0.904 $\pm$ 0.010
	8	0.571 $\pm$ 0.013	0.848 $\pm$ 0.004	0.313 $\pm$ 0.019	0.802 $\pm$ 0.011	0.905 $\pm$ 0.018	0.906 $\pm$ 0.007
	9	0.630 $\pm$ 0.008	0.855 $\pm$ 0.004	0.380 $\pm$ 0.018	0.828 $\pm$ 0.005	0.954 $\pm$ 0.008	0.890 $\pm$ 0.007
	10	0.711 $\pm$ 0.015	0.856 $\pm$ 0.003	0.550 $\pm$ 0.018	0.835 $\pm$ 0.005	0.920 $\pm$ 0.016	0.884 $\pm$ 0.004
	11	0.820 $\pm$ 0.009	0.904 $\pm$ 0.008	0.732 $\pm$ 0.011	0.873 $\pm$ 0.010	0.935 $\pm$ 0.013	0.944 $\pm$ 0.017
	12	0.792 $\pm$ 0.013	0.883 $\pm$ 0.006	0.668 $\pm$ 0.016	0.849 $\pm$ 0.009	0.952 $\pm$ 0.013	0.928 $\pm$ 0.014
	1	0.549 $\pm$ 0.006	0.892 $\pm$ 0.004	0.221 $\pm$ 0.009	0.865 $\pm$ 0.008	0.972 $\pm$ 0.003	0.927 $\pm$ 0.007
	2	0.570 $\pm$ 0.013	0.882 $\pm$ 0.003	0.362 $\pm$ 0.012	0.828 $\pm$ 0.007	0.839 $\pm$ 0.028	0.951 $\pm$ 0.006
	3	0.540 $\pm$ 0.006	0.889 $\pm$ 0.005	0.214 $\pm$ 0.010	0.827 $\pm$ 0.011	0.962 $\pm$ 0.007	0.969 $\pm$ 0.011
	4	0.581 $\pm$ 0.010	0.891 $\pm$ 0.004	0.354 $\pm$ 0.008	0.844 $\pm$ 0.009	0.874 $\pm$ 0.023	0.952 $\pm$ 0.010
Qwen3-0.6B	5	0.486 $\pm$ 0.016	0.894 $\pm$ 0.007	0.133 $\pm$ 0.013	0.839 $\pm$ 0.011	0.943 $\pm$ 0.012	0.965 $\pm$ 0.010
	6	0.566 $\pm$ 0.011	0.897 $\pm$ 0.005	0.349 $\pm$ 0.009	0.857 $\pm$ 0.010	0.848 $\pm$ 0.020	0.949 $\pm$ 0.013
	7	0.540 $\pm$ 0.007	0.902 $\pm$ 0.006	0.211 $\pm$ 0.011	0.862 $\pm$ 0.009	0.967 $\pm$ 0.012	0.955 $\pm$ 0.015
	8	0.608 $\pm$ 0.010	0.903 $\pm$ 0.006	0.361 $\pm$ 0.016	0.861 $\pm$ 0.011	0.927 $\pm$ 0.015	0.958 $\pm$ 0.009
	9	0.610 $\pm$ 0.011	0.901 $\pm$ 0.005	0.391 $\pm$ 0.015	0.867 $\pm$ 0.009	0.892 $\pm$ 0.019	0.944 $\pm$ 0.014
	10	0.636 $\pm$ 0.012	0.890 $\pm$ 0.008	0.450 $\pm$ 0.009	0.847 $\pm$ 0.014	0.877 $\pm$ 0.020	0.946 $\pm$ 0.010
	11	0.577 $\pm$ 0.006	0.886 $\pm$ 0.005	0.257 $\pm$ 0.013	0.838 $\pm$ 0.016	0.991 $\pm$ 0.005	0.948 $\pm$ 0.015
	12	0.626 $\pm$ 0.009	0.893 $\pm$ 0.004	0.365 $\pm$ 0.019	0.844 $\pm$ 0.013	0.964 $\pm$ 0.013	0.958 $\pm$ 0.010
	13	0.638 $\pm$ 0.010	0.888 $\pm$ 0.006	0.382 $\pm$ 0.016	0.849 $\pm$ 0.013	0.970 $\pm$ 0.015	0.938 $\pm$ 0.009
	14	0.644 $\pm$ 0.010	0.885 $\pm$ 0.007	0.401 $\pm$ 0.019	0.852 $\pm$ 0.012	0.958 $\pm$ 0.015	0.929 $\pm$ 0.008
	15	0.646 $\pm$ 0.013	0.885 $\pm$ 0.005	0.409 $\pm$ 0.020	0.852 $\pm$ 0.012	0.952 $\pm$ 0.016	0.929 $\pm$ 0.017
	16	0.689 $\pm$ 0.015	0.873 $\pm$ 0.006	0.544 $\pm$ 0.017	0.837 $\pm$ 0.010	0.877 $\pm$ 0.016	0.920 $\pm$ 0.020
	17	0.657 $\pm$ 0.010	0.861 $\pm$ 0.008	0.419 $\pm$ 0.020	0.807 $\pm$ 0.014	0.964 $\pm$ 0.012	0.932 $\pm$ 0.017
	18	0.708 $\pm$ 0.011	0.868 $\pm$ 0.004	0.503 $\pm$ 0.019	0.833 $\pm$ 0.016	0.973 $\pm$ 0.005	0.913 $\pm$ 0.020
	19	0.690 $\pm$ 0.011	0.873 $\pm$ 0.005	0.462 $\pm$ 0.018	0.859 $\pm$ 0.008	0.986 $\pm$ 0.006	0.892 $\pm$ 0.015
	20	0.742 $\pm$ 0.019	0.880 $\pm$ 0.006	0.588 $\pm$ 0.023	0.858 $\pm$ 0.009	0.941 $\pm$ 0.017	0.909 $\pm$ 0.016
	21	0.757 $\pm$ 0.016	0.878 $\pm$ 0.006	0.584 $\pm$ 0.021	0.835 $\pm$ 0.009	0.981 $\pm$ 0.010	0.934 $\pm$ 0.013
	22	0.798 $\pm$ 0.014	0.882 $\pm$ 0.006	0.670 $\pm$ 0.021	0.854 $\pm$ 0.004	0.965 $\pm$ 0.006	0.917 $\pm$ 0.017
	23	0.757 $\pm$ 0.013	0.888 $\pm$ 0.007	0.578 $\pm$ 0.016	0.857 $\pm$ 0.005	0.989 $\pm$ 0.004	0.929 $\pm$ 0.016
	24	0.810 $\pm$ 0.014	0.894 $\pm$ 0.005	0.694 $\pm$ 0.022	0.861 $\pm$ 0.005	0.960 $\pm$ 0.003	0.936 $\pm$ 0.015
	25	0.732 $\pm$ 0.012	0.890 $\pm$ 0.009	0.528 $\pm$ 0.016	0.866 $\pm$ 0.005	0.997 $\pm$ 0.003	0.922 $\pm$ 0.017
	26	0.749 $\pm$ 0.013	0.895 $\pm$ 0.007	0.561 $\pm$ 0.017	0.858 $\pm$ 0.002	0.991 $\pm$ 0.003	0.943 $\pm$ 0.017
	27	0.811 $\pm$ 0.004	0.881 $\pm$ 0.003	0.741 $\pm$ 0.008	0.824 $\pm$ 0.007	0.901 $\pm$ 0.011	0.954 $\pm$ 0.011
	28	0.822 $\pm$ 0.010	0.870 $\pm$ 0.007	0.733 $\pm$ 0.014	0.800 $\pm$ 0.012	0.937 $\pm$ 0.006	0.961 $\pm$ 0.009
KaLM-mini	1	0.595 $\pm$ 0.011	0.893 $\pm$ 0.004	0.344 $\pm$ 0.002	0.846 $\pm$ 0.007	0.920 $\pm$ 0.021	0.954 $\pm$ 0.009
	2	0.620 $\pm$ 0.017	0.894 $\pm$ 0.005	0.422 $\pm$ 0.014	0.845 $\pm$ 0.005	0.875 $\pm$ 0.021	0.958 $\pm$ 0.013

(continued on next page)

Task B: Top-K Ranking (appendix, 5-seed mean $\pm$ std, SIF+ABTT)							
Model	Layer	Acc@1 (base)	Acc@1 (SIF+ABTT)	Existing (base)	Existing (SIF+ABTT)	New (base)	New (SIF+ABTT)
	3	0.604 $\pm$ 0.010	0.894 $\pm$ 0.003	0.378 $\pm$ 0.009	0.843 $\pm$ 0.003	0.898 $\pm$ 0.017	0.960 $\pm$ 0.011
	4	0.562 $\pm$ 0.015	0.889 $\pm$ 0.004	0.265 $\pm$ 0.009	0.842 $\pm$ 0.010	0.946 $\pm$ 0.016	0.951 $\pm$ 0.011
	5	0.550 $\pm$ 0.012	0.890 $\pm$ 0.005	0.270 $\pm$ 0.007	0.841 $\pm$ 0.006	0.912 $\pm$ 0.018	0.954 $\pm$ 0.012
	6	0.522 $\pm$ 0.015	0.894 $\pm$ 0.004	0.202 $\pm$ 0.011	0.850 $\pm$ 0.006	0.935 $\pm$ 0.013	0.950 $\pm$ 0.015
	7	0.604 $\pm$ 0.017	0.894 $\pm$ 0.003	0.394 $\pm$ 0.016	0.840 $\pm$ 0.011	0.876 $\pm$ 0.020	0.963 $\pm$ 0.012
	8	0.607 $\pm$ 0.013	0.866 $\pm$ 0.003	0.387 $\pm$ 0.010	0.798 $\pm$ 0.010	0.892 $\pm$ 0.016	0.955 $\pm$ 0.010
	9	0.614 $\pm$ 0.010	0.874 $\pm$ 0.005	0.386 $\pm$ 0.009	0.821 $\pm$ 0.011	0.909 $\pm$ 0.013	0.944 $\pm$ 0.011
	10	0.602 $\pm$ 0.012	0.860 $\pm$ 0.007	0.359 $\pm$ 0.005	0.774 $\pm$ 0.012	0.917 $\pm$ 0.017	0.970 $\pm$ 0.012
	11	0.601 $\pm$ 0.012	0.864 $\pm$ 0.008	0.360 $\pm$ 0.009	0.792 $\pm$ 0.012	0.911 $\pm$ 0.020	0.957 $\pm$ 0.008
	12	0.638 $\pm$ 0.009	0.868 $\pm$ 0.006	0.415 $\pm$ 0.008	0.801 $\pm$ 0.009	0.927 $\pm$ 0.016	0.955 $\pm$ 0.005
	13	0.643 $\pm$ 0.010	0.862 $\pm$ 0.006	0.435 $\pm$ 0.007	0.792 $\pm$ 0.005	0.912 $\pm$ 0.016	0.953 $\pm$ 0.016
	14	0.653 $\pm$ 0.009	0.858 $\pm$ 0.007	0.479 $\pm$ 0.002	0.793 $\pm$ 0.009	0.879 $\pm$ 0.016	0.943 $\pm$ 0.009
	15	0.659 $\pm$ 0.007	0.868 $\pm$ 0.005	0.458 $\pm$ 0.006	0.812 $\pm$ 0.009	0.919 $\pm$ 0.010	0.941 $\pm$ 0.006
	16	0.682 $\pm$ 0.005	0.869 $\pm$ 0.008	0.475 $\pm$ 0.006	0.824 $\pm$ 0.008	0.951 $\pm$ 0.010	0.926 $\pm$ 0.009
	17	0.756 $\pm$ 0.007	0.885 $\pm$ 0.009	0.611 $\pm$ 0.018	0.841 $\pm$ 0.006	0.944 $\pm$ 0.010	0.942 $\pm$ 0.015
	18	0.780 $\pm$ 0.010	0.881 $\pm$ 0.007	0.635 $\pm$ 0.016	0.840 $\pm$ 0.002	0.967 $\pm$ 0.007	0.933 $\pm$ 0.016
	19	0.804 $\pm$ 0.013	0.876 $\pm$ 0.007	0.672 $\pm$ 0.022	0.831 $\pm$ 0.007	0.977 $\pm$ 0.003	0.935 $\pm$ 0.008
	20	0.793 $\pm$ 0.011	0.877 $\pm$ 0.004	0.649 $\pm$ 0.020	0.831 $\pm$ 0.003	0.979 $\pm$ 0.004	0.936 $\pm$ 0.008
	21	0.804 $\pm$ 0.011	0.871 $\pm$ 0.007	0.672 $\pm$ 0.018	0.830 $\pm$ 0.005	0.975 $\pm$ 0.006	0.924 $\pm$ 0.012
	22	0.870 $\pm$ 0.008	0.873 $\pm$ 0.004	0.810 $\pm$ 0.007	0.812 $\pm$ 0.005	0.948 $\pm$ 0.008	0.952 $\pm$ 0.009
	23	0.873 $\pm$ 0.007	0.884 $\pm$ 0.004	0.811 $\pm$ 0.008	0.833 $\pm$ 0.004	0.953 $\pm$ 0.007	0.951 $\pm$ 0.007
	24	0.848 $\pm$ 0.011	0.856 $\pm$ 0.005	0.749 $\pm$ 0.019	0.806 $\pm$ 0.009	0.975 $\pm$ 0.007	0.921 $\pm$ 0.010

E Attribution Metric Sweeps

This appendix expands the main attribution table across Sufficiency and Comprehensiveness thresholds of 10%, 25%, and 50%, and MinFrac recovery thresholds of 0.70, 0.80, 0.90, and 0.95. The main text uses one global midpoint setting rather than selecting thresholds per model or per method.

F Cluster Geometry: Additional Models

Figures 9 and 10 extend the cluster-geometry panels to LaBSE, Qwen3-0.6B, and KaLM-mini. They use the same six directories and the same baseline versus pure ABTT comparison.

Model	Method	$\rho_{\text{LOO}} \uparrow$		Suff@10% $\uparrow$		Suff@25% $\uparrow$		Suff@50% $\uparrow$		Comp@10% $\uparrow$		Comp@25% $\uparrow$		Comp@50% $\uparrow$		MinFrac@0.70 $\downarrow$		MinFrac@0.80 $\downarrow$		MinFrac@0.90 $\downarrow$		MinFrac@0.95 $\downarrow$	
		base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt
LaTa	IG	-0.001	0.396	0.967	0.428	0.976	0.681	0.975	0.921	0.320	0.267	0.719	0.581	0.850	0.884	0.040	0.293	0.041	0.379	0.045	0.489	0.047	0.557
	MaRC	0.042	0.401	0.404	0.475	0.534	0.645	0.582	0.863	0.116	0.250	0.361	0.539	0.506	0.845	0.368	0.356	0.383	0.459	0.396	0.568	0.438	0.624
PhilTa	IG	0.182	0.607	0.816	0.848	0.959	0.963	0.988	1.057	0.704	0.436	0.969	0.747	1.010	0.992	0.027	0.090	0.079	0.141	0.162	0.229	0.215	0.300
	MaRC	0.190	0.331	0.786	0.531	0.911	0.754	0.949	0.935	0.657	0.237	0.963	0.521	0.978	0.793	0.058	0.281	0.104	0.358	0.221	0.448	0.327	0.503
mT5-base	IG	0.164	0.597	0.810	0.785	0.892	1.027	0.937	1.192	0.054	0.376	0.110	0.717	0.284	1.071	0.055	0.161	0.116	0.214	0.338	0.291	0.623	0.336
	MaRC	0.250	0.392	0.865	0.597	0.913	0.827	0.951	1.055	0.010	0.279	0.042	0.592	0.118	0.923	0.026	0.277	0.066	0.346	0.212	0.434	0.416	0.488

Table 15: Appendix attribution sweep for the two paper-facing attribution views, IG and retrieval-adapted MaRC, across the three main-paper models. Direction arrows are printed in the column headers. Sufficiency and Comprehensiveness are ratio metrics at $\{10\%, 25\%, 50\%\}$; MinFrac is the minimum retained query-token fraction needed to recover $\tau \in \{0.70, 0.80, 0.90, 0.95\}$ of the full cosine (lower is better). Ratio metrics require $|S_v(q_{\text{full}}, c)| \geq 0.05$; ρ_{LOO} excludes $|\Delta| < 10^{-6}$. The completeness report contains no missing metric cells.

Model	Method	$\rho_{\text{LOO}} \uparrow$		Suff@10% $\uparrow$		Suff@25% $\uparrow$		Suff@50% $\uparrow$		Comp@10% $\uparrow$		Comp@25% $\uparrow$		Comp@50% $\uparrow$		MinFrac@0.70 $\downarrow$		MinFrac@0.80 $\downarrow$		MinFrac@0.90 $\downarrow$		MinFrac@0.95 $\downarrow$	
		base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt	base	abtt
LaTa	BERTScore	0.204	0.430	1.035	0.365	1.047	0.626	1.047	0.904	0.916	0.221	1.301	0.436	1.280	0.719	0.018	0.318	0.021	0.394	0.026	0.481	0.036	0.534
	OT	-0.001	0.396	0.967	0.428	0.976	0.681	0.975	0.921	0.320	0.267	0.719	0.581	0.850	0.884	0.040	0.293	0.041	0.379	0.045	0.489	0.047	0.557
	Att-weighted	-0.137	0.106	-0.575	0.142	-0.377	0.295	-0.131	0.588	0.001	0.159	-0.005	0.330	-0.019	0.543	0.746	0.569	0.756	0.653	0.774	0.747	0.796	0.804
	Att-standalone	0.005	0.368	0.115	0.358	0.239	0.640	0.409	0.873	0.043	0.251	0.112	0.536	0.342	0.830	0.516	0.333	0.531	0.422	0.558	0.537	0.584	0.597
	DLA	0.054	0.278	1.017	0.350	1.004	0.607	1.003	0.820	0.937	0.239	1.289	0.505	1.248	0.792	0.018	0.368	0.021	0.471	0.026	0.583	0.036	0.650
	random	0.003	-0.003	0.403	0.222	0.726	0.360	0.906	0.599	0.024	0.072	0.046	0.181	0.105	0.393	0.176	0.561	0.182	0.667	0.192	0.779	0.212	0.842
	inverse	0.001	-0.396	-0.046	-0.014	0.122	-0.001	0.151	0.124	0.000	-0.046	0.001	-0.035	0.025	0.086	0.553	0.837	0.562	0.882	0.602	0.920	0.629	0.942
PhilTa	BERTScore	0.309	0.550	0.861	0.771	0.985	0.942	1.035	1.050	0.687	0.313	1.032	0.580	1.042	0.808	0.026	0.128	0.070	0.188	0.190	0.284	0.280	0.343
	OT	0.182	0.607	0.816	0.848	0.959	0.963	0.988	1.057	0.704	0.436	0.969	0.747	1.010	0.992	0.027	0.090	0.079	0.141	0.162	0.229	0.215	0.300
	Att-weighted	0.053	0.108	0.060	0.341	0.163	0.584	0.219	0.753	0.011	0.157	0.028	0.326	0.059	0.524	0.673	0.397	0.693	0.505	0.742	0.627	0.795	0.698
	Att-standalone	0.046	0.205	0.017	0.351	0.038	0.654	0.064	0.827	0.007	0.163	0.040	0.378	0.069	0.616	0.773	0.323	0.782	0.417	0.796	0.537	0.808	0.612
	DLA	0.203	0.487	0.805	0.863	0.942	0.910	0.991	0.957	0.715	0.438	1.017	0.729	1.028	0.927	0.028	0.087	0.080	0.142	0.194	0.260	0.260	0.346
	random	-0.001	-0.010	0.589	0.288	0.826	0.478	0.953	0.709	0.011	0.050	0.016	0.124	0.066	0.288	0.135	0.437	0.145	0.563	0.204	0.713	0.284	0.798
	inverse	-0.182	-0.607	-0.009	-0.074	-0.025	-0.074	-0.009	0.014	-0.001	-0.049	-0.001	-0.073	0.013	-0.055	0.917	0.920	0.921	0.950	0.932	0.975	0.942	0.987
mT5-base	BERTScore	0.271	0.556	0.615	0.694	0.776	1.005	0.917	1.197	0.014	0.270	0.038	0.579	0.082	0.906	0.148	0.199	0.214	0.255	0.333	0.322	0.452	0.358
	OT	0.164	0.597	0.810	0.785	0.892	1.027	0.937	1.192	0.054	0.376	0.110	0.717	0.284	1.071	0.055	0.161	0.116	0.214	0.338	0.291	0.623	0.336
	Att-weighted	0.173	0.523	0.872	0.890	0.923	1.008	0.942	1.099	0.012	0.364	0.048	0.687	0.249	0.986	0.036	0.117	0.074	0.181	0.239	0.278	0.526	0.346
	Att-standalone	0.110	0.064	0.640	0.180	0.650	0.455	0.862	0.643	0.034	0.184	0.068	0.362	0.135	0.464	0.066	0.471	0.137	0.575	0.412	0.718	0.608	0.786
	DLA	0.129	0.481	0.296	0.922	0.786	0.996	0.949	1.010	0.029	0.408	0.045	0.683	0.083	0.908	0.211	0.097	0.259	0.159	0.347	0.257	0.460	0.337
	random	-0.000	-0.004	0.846	0.296	0.926	0.483	0.969	0.687	0.004	0.059	0.013	0.153	0.033	0.317	0.045	0.460	0.079	0.591	0.178	0.735	0.313	0.813
	inverse	-0.164	-0.597	-0.127	-0.235	0.203	-0.237	0.722	-0.065	0.022	-0.108	0.045	-0.187	0.064	-0.189	0.472	0.896	0.563	0.936	0.740	0.971	0.863	0.985

Table 16: Supplemental attribution sweep for additional methods and sanity baselines. *random* and *inverse* are diagnostic baselines, not paper-facing attribution methods. Direction arrows are printed in the column headers. Sufficiency and Comprehensiveness are ratio metrics at $\{10\%, 25\%, 50\%\}$; MinFrac is the minimum retained query-token fraction needed to recover $\tau \in \{0.70, 0.80, 0.90, 0.95\}$ of the full cosine (lower is better). Ratio metrics require $|S_v(q_{\text{full}}, c)| \geq 0.05$; ρ_{LOO} excludes $|\Delta| < 10^{-6}$. The completeness report contains no missing metric cells.

2D TSNE of labelled fragments at each model's best layer




Figure 9: t-SNE projections of the six most-populated directories at each selected layer for LaBSE, Qwen3-0.6B, and KaLM-mini. Columns compare baseline mean pooling with pure ABTT using $D = 10$.

2D UMAP of labelled fragments at each model's best layer




Figure 10: UMAP projections matching Figure 9. The cluster-tightening pattern appears under a second projection method.