Attribution-quality metrics: MaRC and IG, baseline vs ABTT

200 POSITIVE test pairs per model (same-folder pairs only; harder/more meaningful slice than the 200-random set); LaTa + PhilTa.

Generated 2026-04-26.

Method	Suff@25% ↑		Comp@25% ↑		Cmpct@0.8 ↓		ρ_{LOO} ↑	
	base	abtt	base	abtt	base	abtt	base	abtt
LaTa								
IG	0.521	**0.735**	0.036	**0.613**	**0.359**	0.363	-0.055	**0.460**
MaRC	0.039	**0.740**	-0.019	**0.564**	0.905	**0.393**	-0.124	**0.455**
PhilTa								
IG	**0.807**	0.512	0.176	**0.416**	**0.132**	0.501	0.058	**0.245**
MaRC	-0.350	**0.424**	-0.002	**0.430**	0.648	**0.618**	-0.080	**0.295**

Within each (method, metric) cell, the better of {baseline, ABTT} is **bolded**. Group means over 200 positive (same-folder) test pairs per model. **Cosine note:** ABTT `full_cos_mean` is now *below* baseline (0.564 vs 0.934 for LaTa; 0.574 vs 0.954 for PhilTa), reversing the direction-of-cleaning bug from the 200-random run (where ABTT was 0.985 > baseline 0.926); per the cosine-inflation investigation, values remain elevated above the D=10 refit demo range (~ 0.30), so this is a partial fix.

Per-cell ABTT-wins count (out of 4 metrics):

- LaTa × IG: **3/4** (suff + comp + rho; cmpct essentially tied at 0.359 vs 0.363)
- LaTa × MaRC: **4/4** (clean sweep; cmpct collapses from 0.905 to 0.393)
- PhilTa × IG: **2/4** (comp + rho; baseline already strong on suff and cmpct)
- PhilTa × MaRC: **4/4** (clean sweep)

Reading guide. Suff@25% (↑): similarity retained with only the top 25% attributed tokens. **Comp@25%** (↑): similarity drop when those tokens are removed. **Cmpct@0.8** (↓): token fraction to recover 80% of full-pair cosine. $\boldsymbol{\rho}_{\text{LOO}}$ (↑): rank correlation between attribution magnitude and true leave-one-out $\Delta\cos$.

Bottom line at 200 *positive* pairs/model.

- The "≥ 3/4 metrics for both MaRC and IG on at least one T5 model" threshold is now **cleared on 3 of 4 cells** (LaTa×IG 3/4, LaTa×MaRC 4/4, PhilTa×MaRC 4/4). On the 200-random set it cleared on 0/4. The positive-pair slice tells a substantially stronger story.
- ρ_{LOO} now improves on **all 4** T5 cells (range +0.19 to +0.58). The "ABTT universally lifts ρ_{LOO}" claim from the original 20-pair snapshot *is* supported on this slice; the regression seen at PhilTa×MaRC on the random set ($-0.036 \rightarrow -0.120$) does not appear here ($-0.080 \rightarrow +0.295$).
- The holdout is PhilTa×IG: ABTT loses Suff@25% (0.807 → 0.512) and Cmpct@0.8 (0.132 → 0.501). The PhilTa baseline already retains 81% of the cosine with only the top quartile of tokens, and recovers 80% of full-pair cosine using only 13% of tokens — a hard ceiling for ABTT to clear.
- Caveat unchanged: cross-variant deltas conflate two effects (ABTT changes both the attribution scores *and* the cosine being explained, and the cosine is now meaningfully lower under ABTT — 0.56–0.57 vs 0.93–0.95), so these are descriptive shifts on different fidelity targets rather than apples-to-apples.